


DOCKER AND PYTHON


Making them play nicely and securely for Data Science and Machine Learning

TANIA ALLARD, PHD

Sr. Developer Advocate @Microsoft.



ixek | <https://bit.ly/europython-ml-docker>



@ixek



@trallard



trallard.dev

THESE SLIDES

<https://bit.ly/europython-ml-docker>

WHAT YOU'LL LEARN TODAY

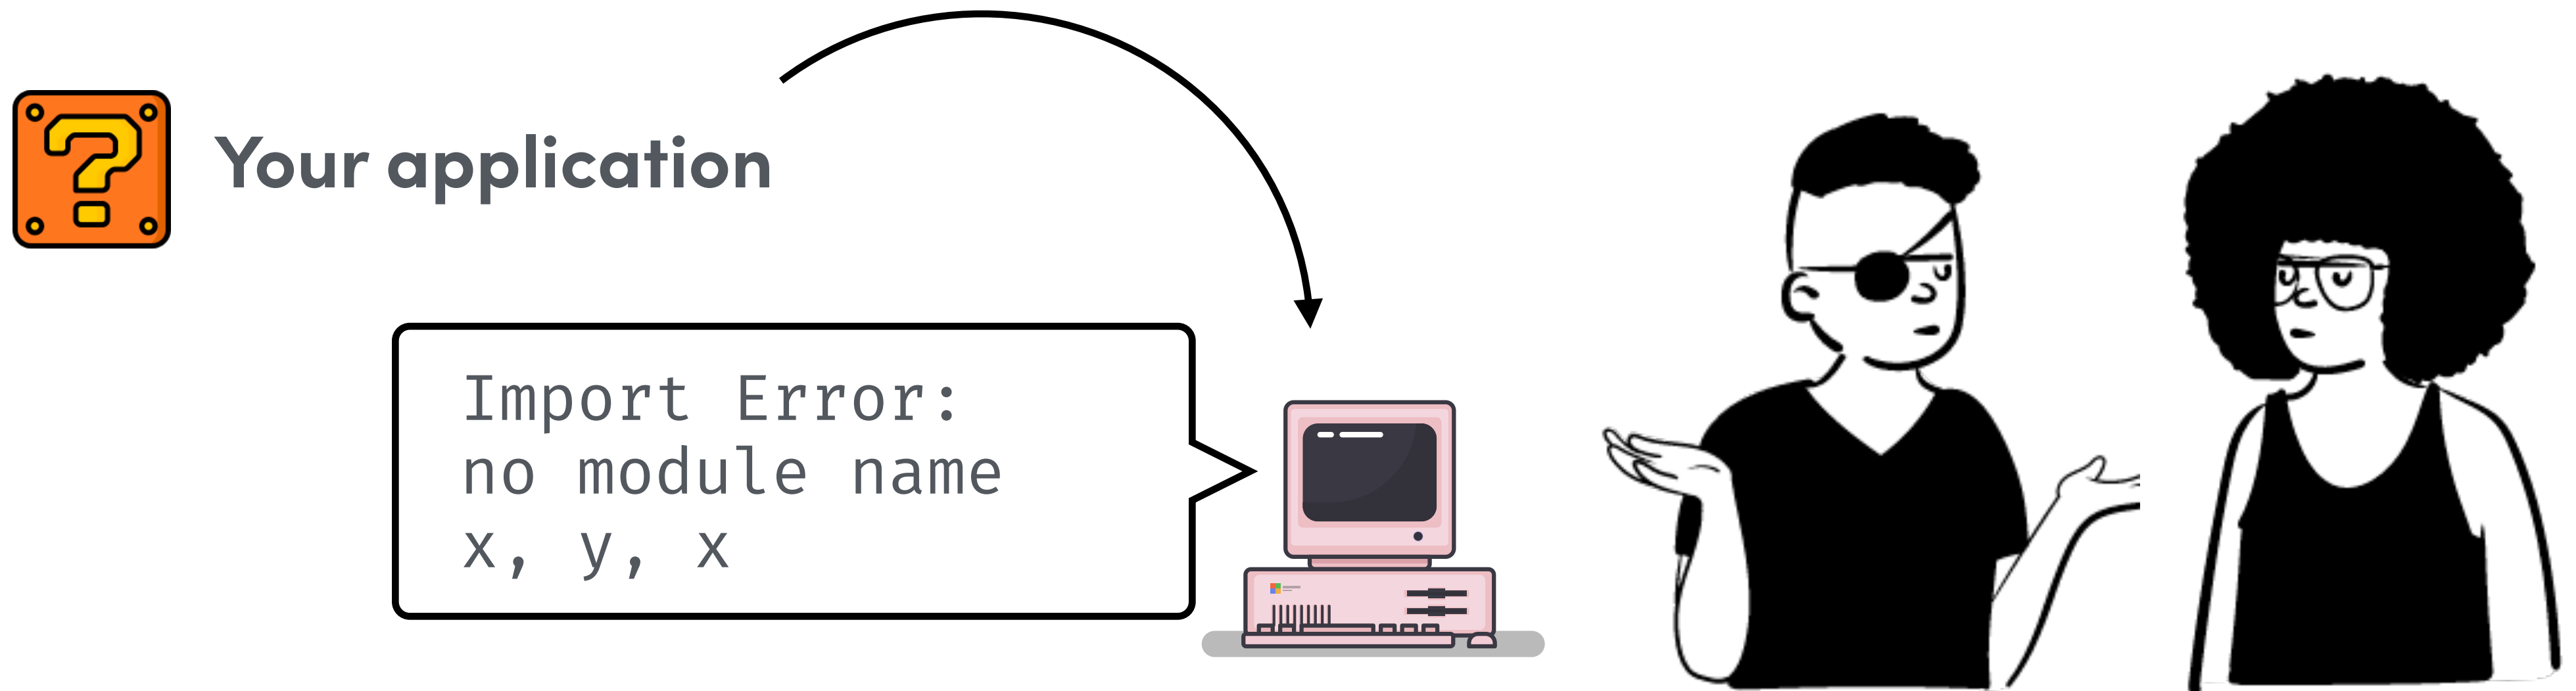
- Why using Docker?
- Docker for Data Science and Machine Learning
- Security and performance
- Do not reinvent the wheel, automate
- Tips and trick to use Docker






WHY DOCKER?

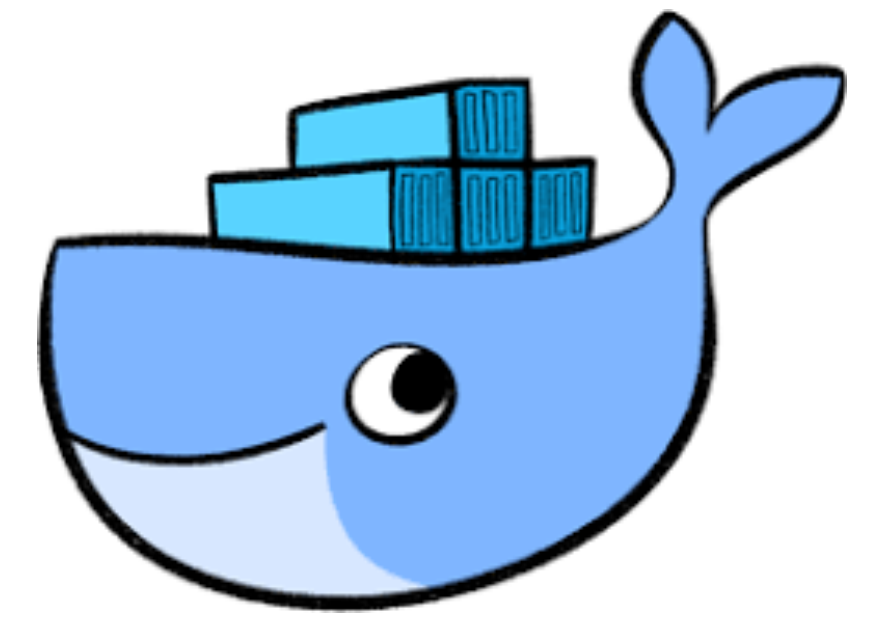
DEV LIFE WITHOUT DOCKER OR CONTAINERS



How are your users or colleagues meant to know what dependencies they need?




WHAT IS DOCKER?



A tool that helps you to create, deploy and run your applications or projects **by using containers.**

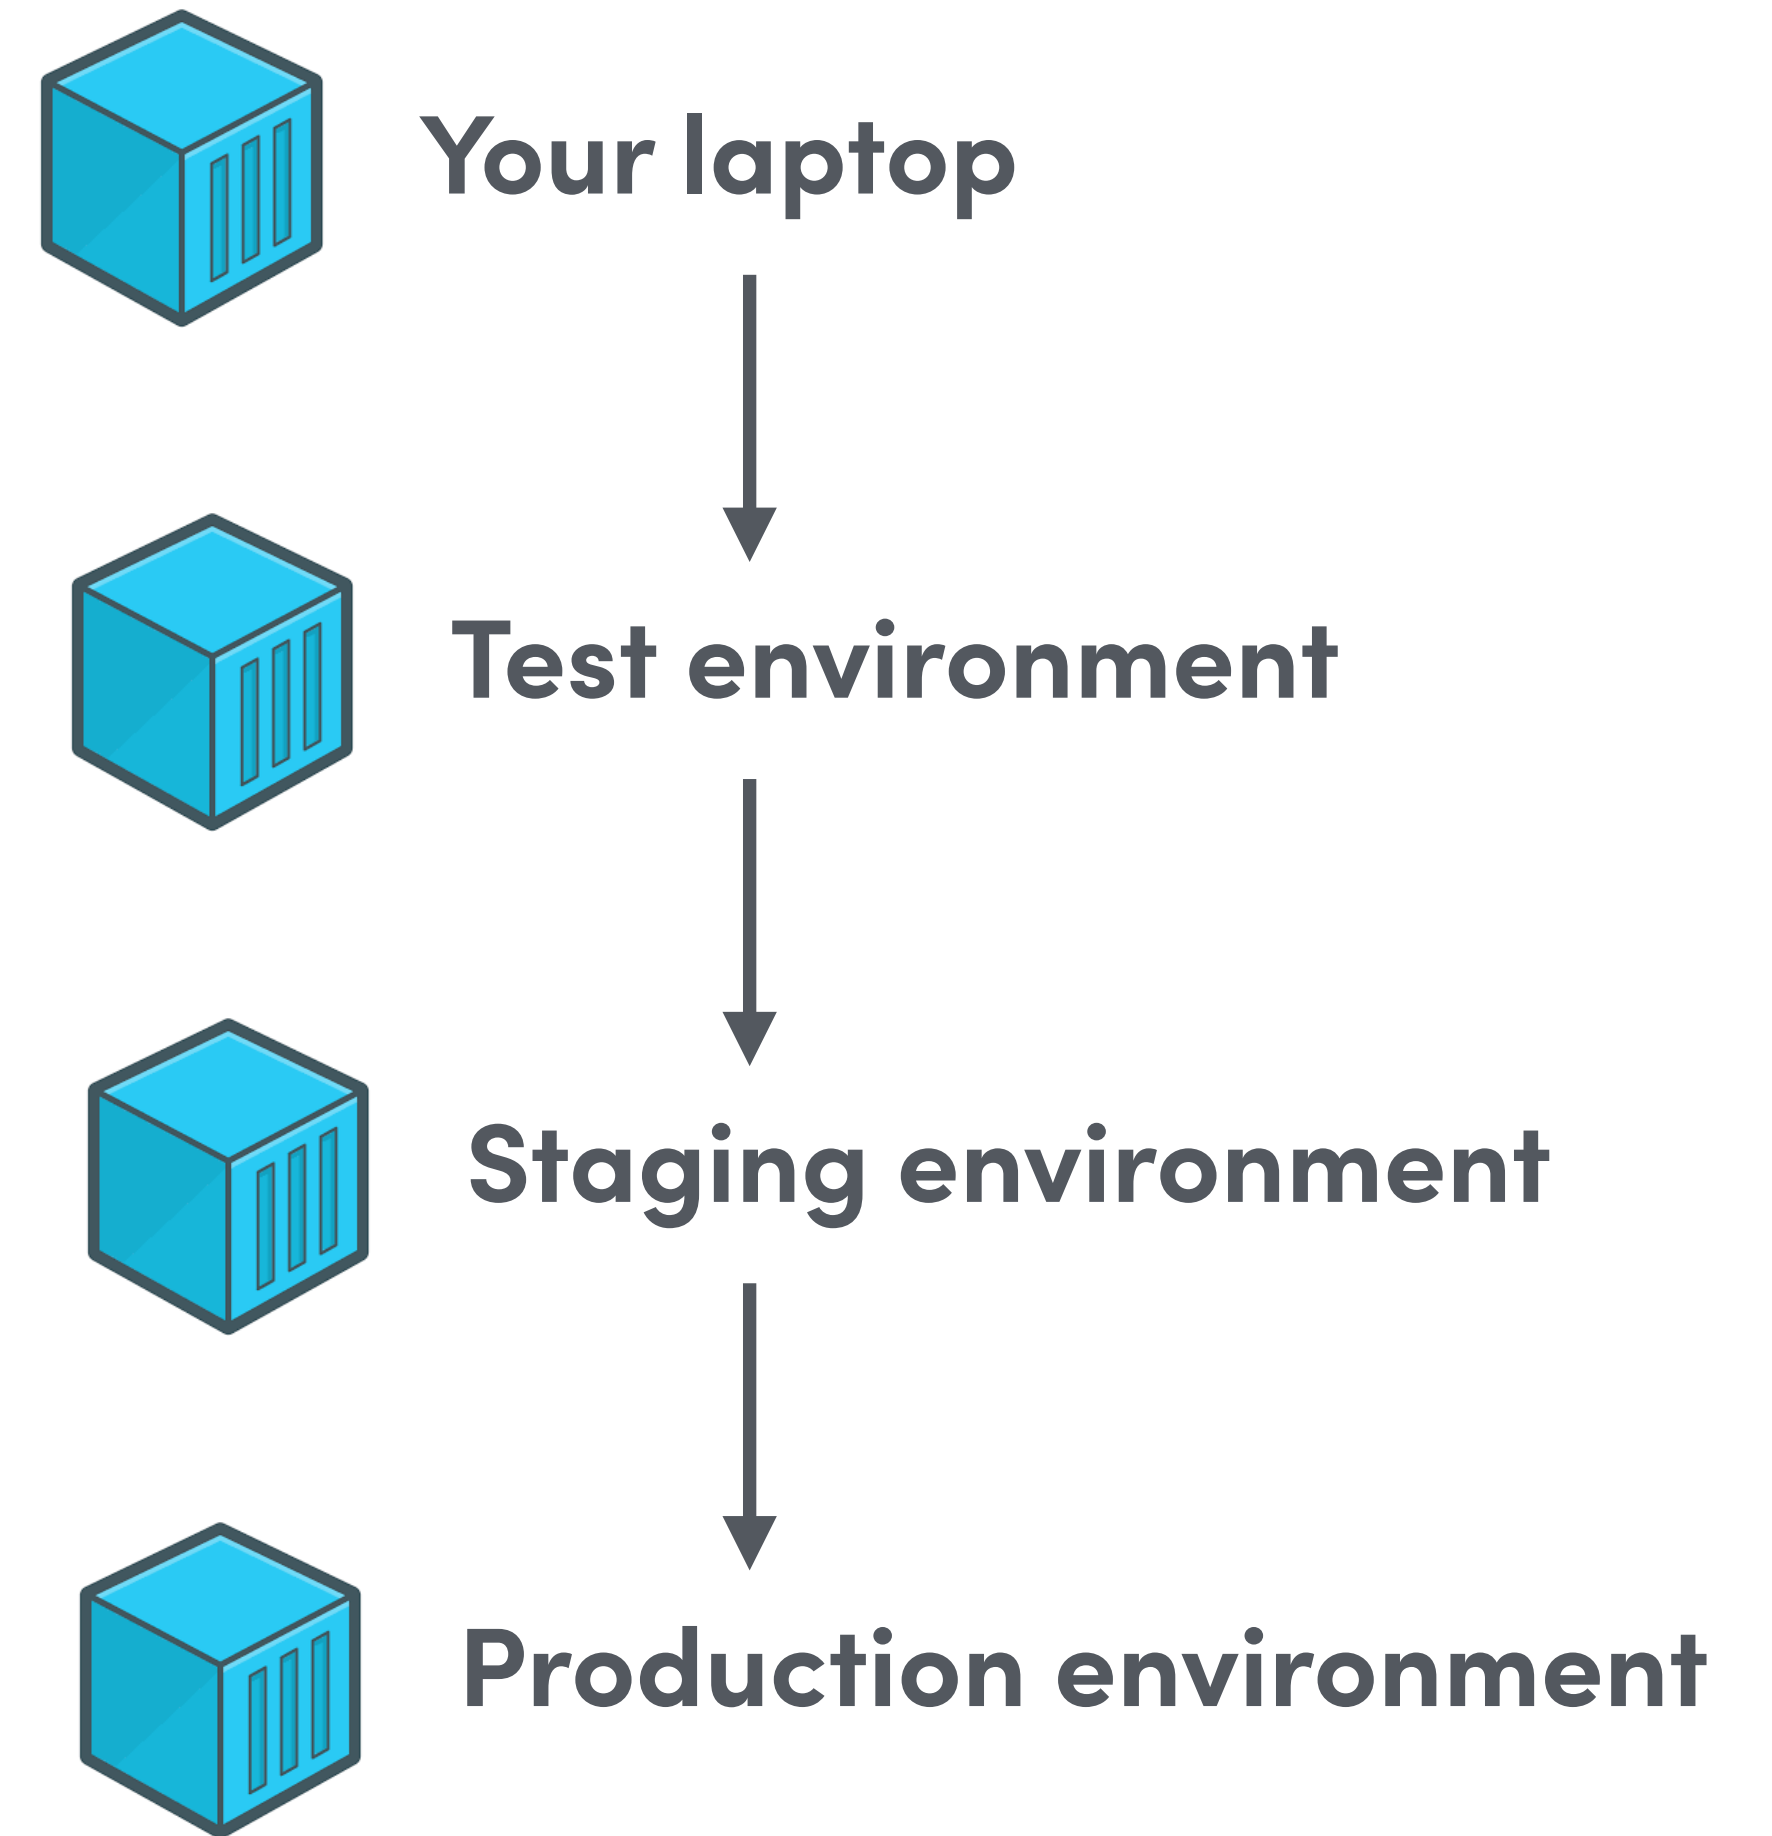


This is a container


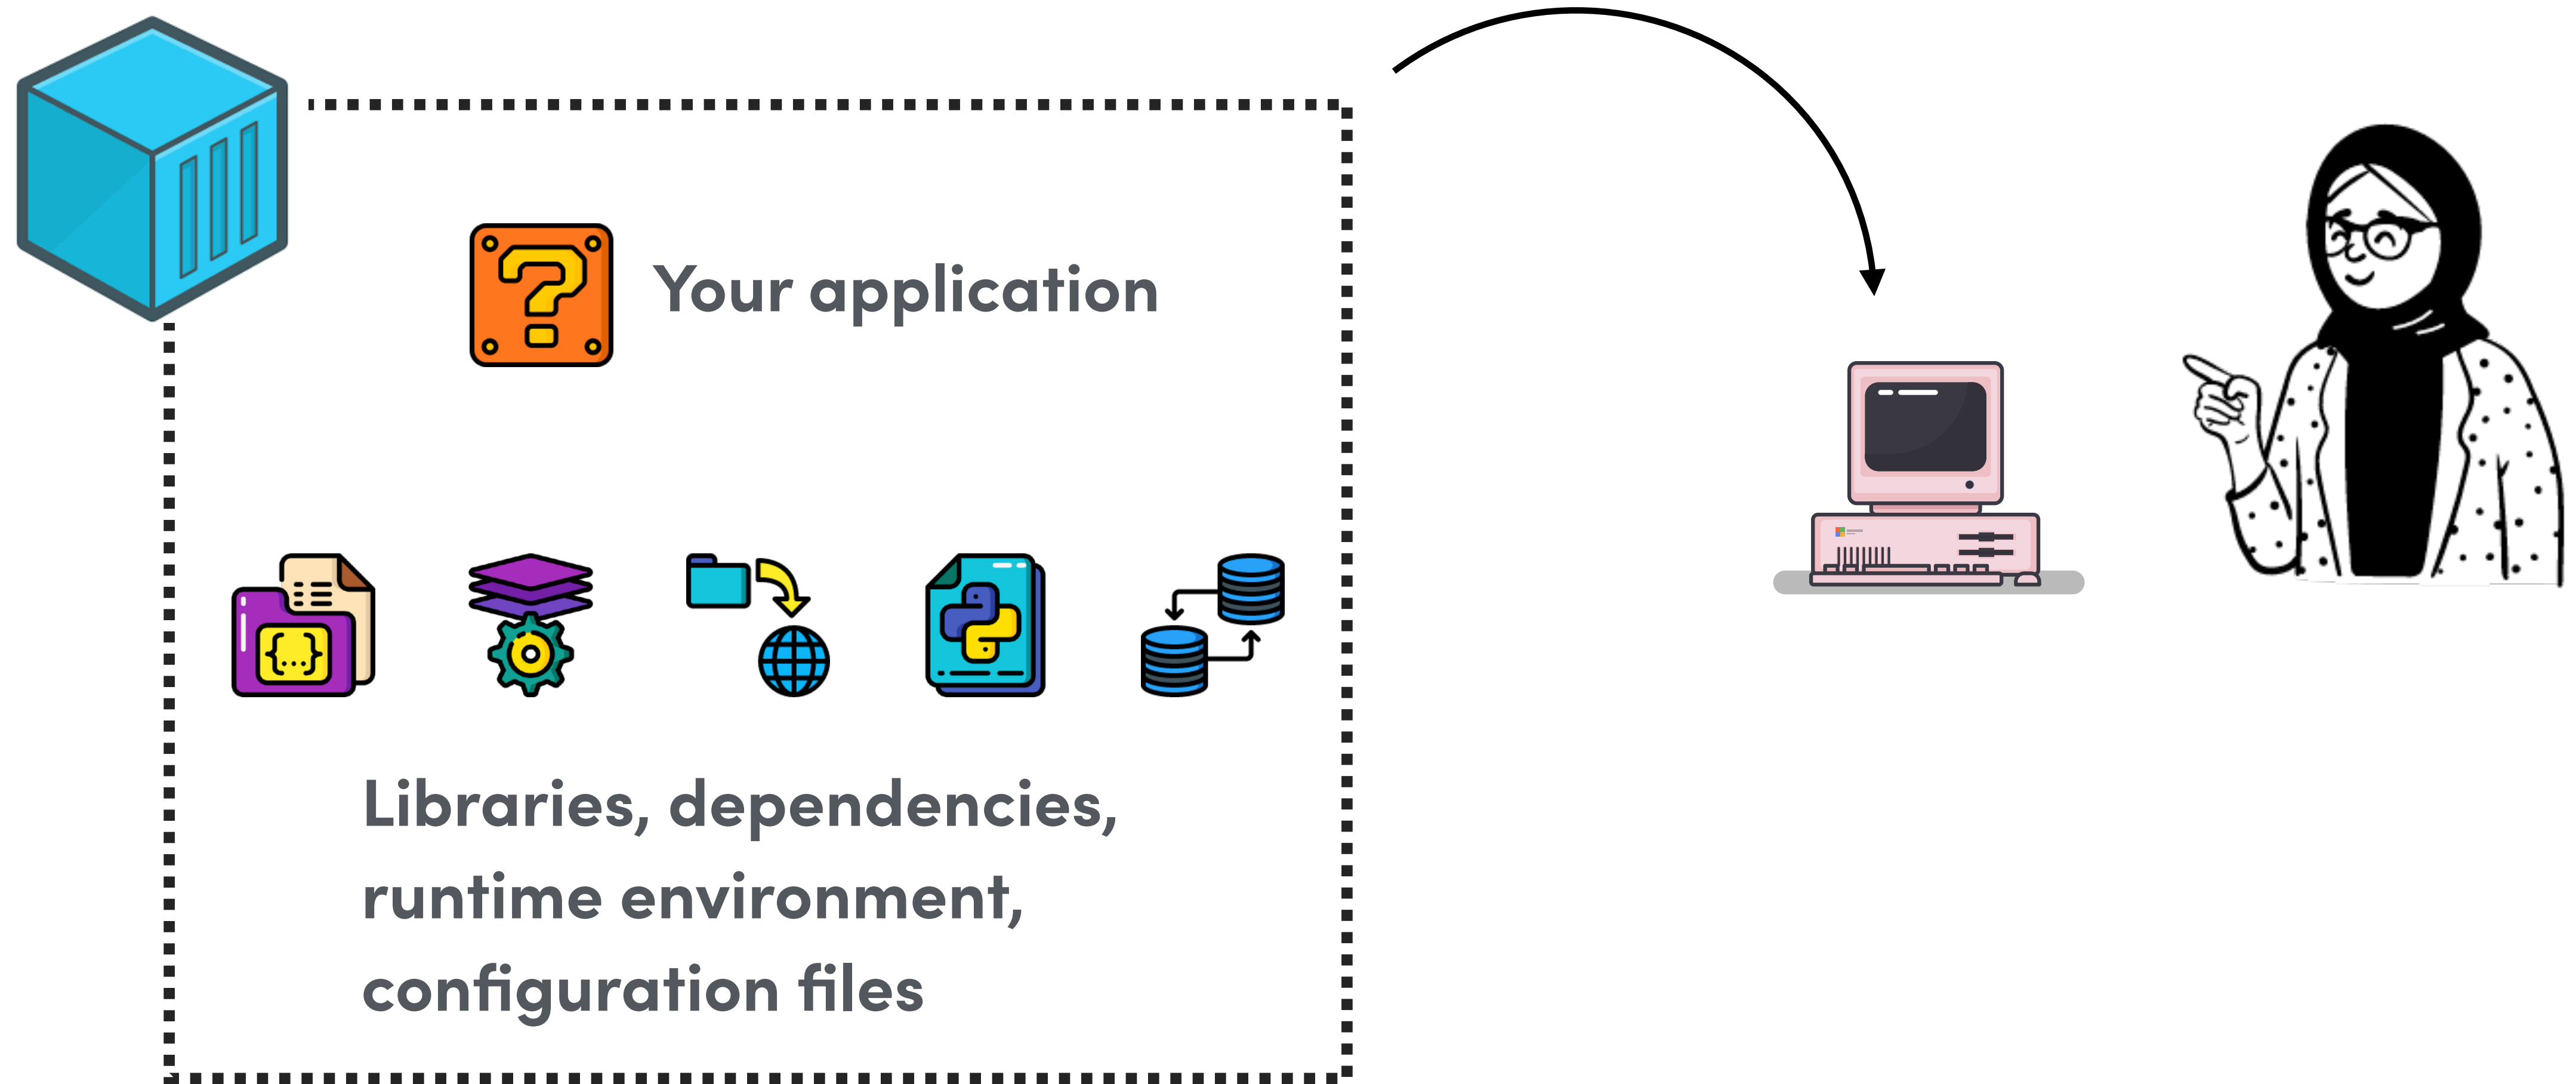


HOW DO CONTAINERS HELP ME?

They provide a solution to the problem of how to get software to run reliably when **moved from one computing environment to another**




DEV LIFE WITH CONTAINERS




THAT SOUNDS A LOT LIKE A VIRTUAL MACHINE

Each app is
containerised




At the app level:
Each runs as an isolated process



THAT SOUNDS A LOT LIKE A VIRTUAL MACHINE

CONTAINERS



Full OS + app +
binaries +
libraries



VIRTUAL MACHINE

At the hardware level







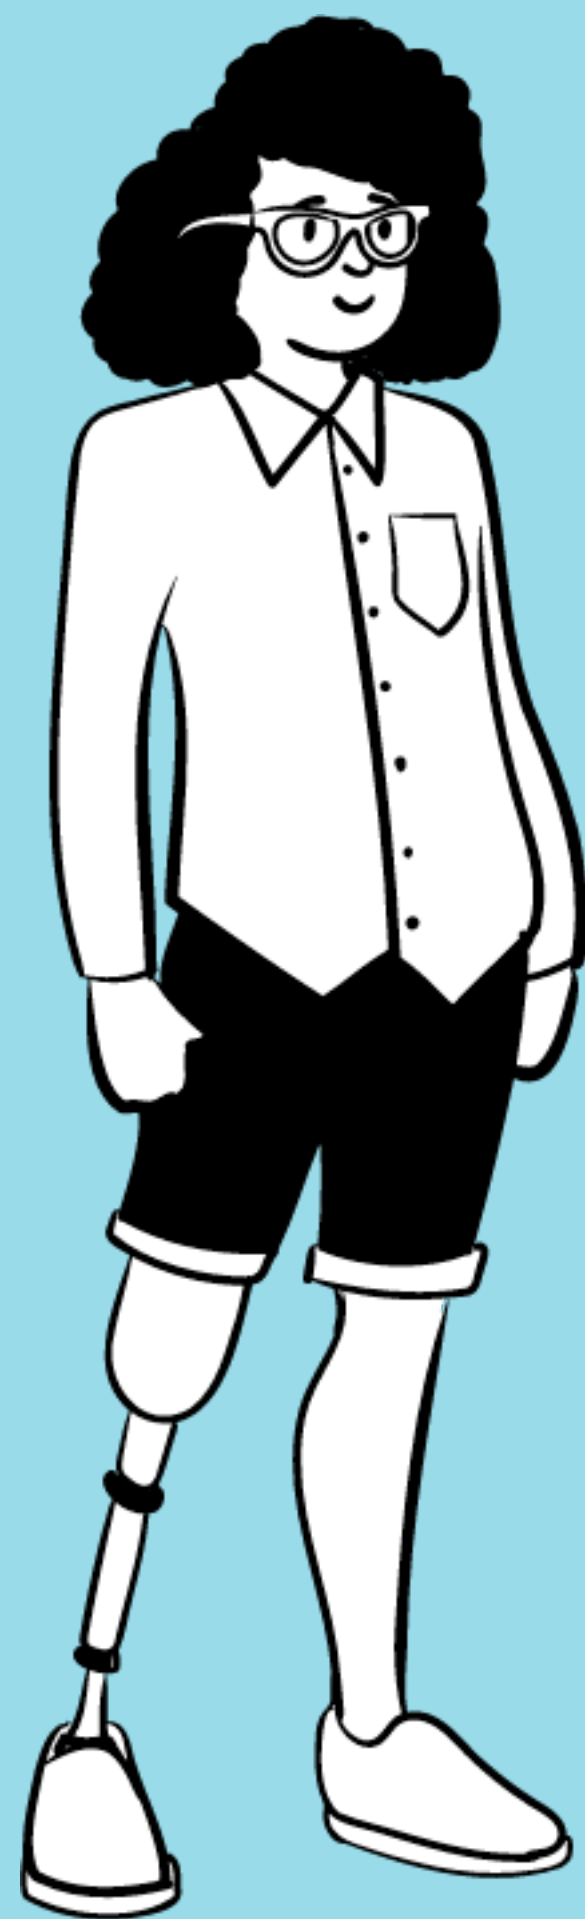
IMAGE VS CONTAINER

- **Image**: archive with all the data needed to run the app
- When you run an image it creates a **container**



COMMON PAIN POINTS IN DS AND ML

- Complex setups / dependencies
- Reliance on data / databases
- Fast evolving projects (iterative R&D process)
- Docker is complex and can take a lot of time to upskill
- Are containers secure enough for my data / model /algorithm?



DOCKER FOR DATA SCIENCE AND MACHINE LEARNING

HOW IS IT DIFFERENT FROM WEB APPS FOR EXAMPLE?



Drawing a line to the scope of Python packaging

■ Packaging



Tzu-ping Chung uranusjr

Feb '19

Another topic in the Big Picture thread I found interesting 🤔

twitter.com




Sylvain Corlay (SylvainCorlay)

@dstufft @pwang @WillingCarol @brettsky @vorpalsmith @zooba @uranusjr
@acanthamoeba @ncoghlan_dev @kushaldas Although there is a continuum of things between "we need BLAS", and "we need R". Where is the limit?

It feels to me that with language-specific packaging tools, you soon find a cliff somewhere between those two usecases.

6:46 PM - 11 Feb 2019

<https://twitter.com/dstufft/status/1095164069802397696>



ixek | <https://bit.ly/europython-ml-docker>


HOW IS IT DIFFERENT FROM WEB APPS FOR EXAMPLE?

- Not every deliverable is an app
- Not every deliverable is a model either
- Heavily relies on data
- Mixture of wheels and compiled packages
- Security access levels – for data and software
- Mixture of stakeholders: data scientists, software engineers, ML engineers




BUILDING DOCKER IMAGES

Dockerfiles are used to create Docker images by providing a set of instructions to install software, configure your image or copy files



```
docker-ds - Dockerfile
1  # word of caution - this is a bad example
2  FROM python:3
3
4  COPY yourscript.py /
5
6  RUN pip install flask
7
8  CMD [ "python", "./yourscript.py" ]
```




DISSECTING DOCKER IMAGES

Base image

Main instructions

Entry command


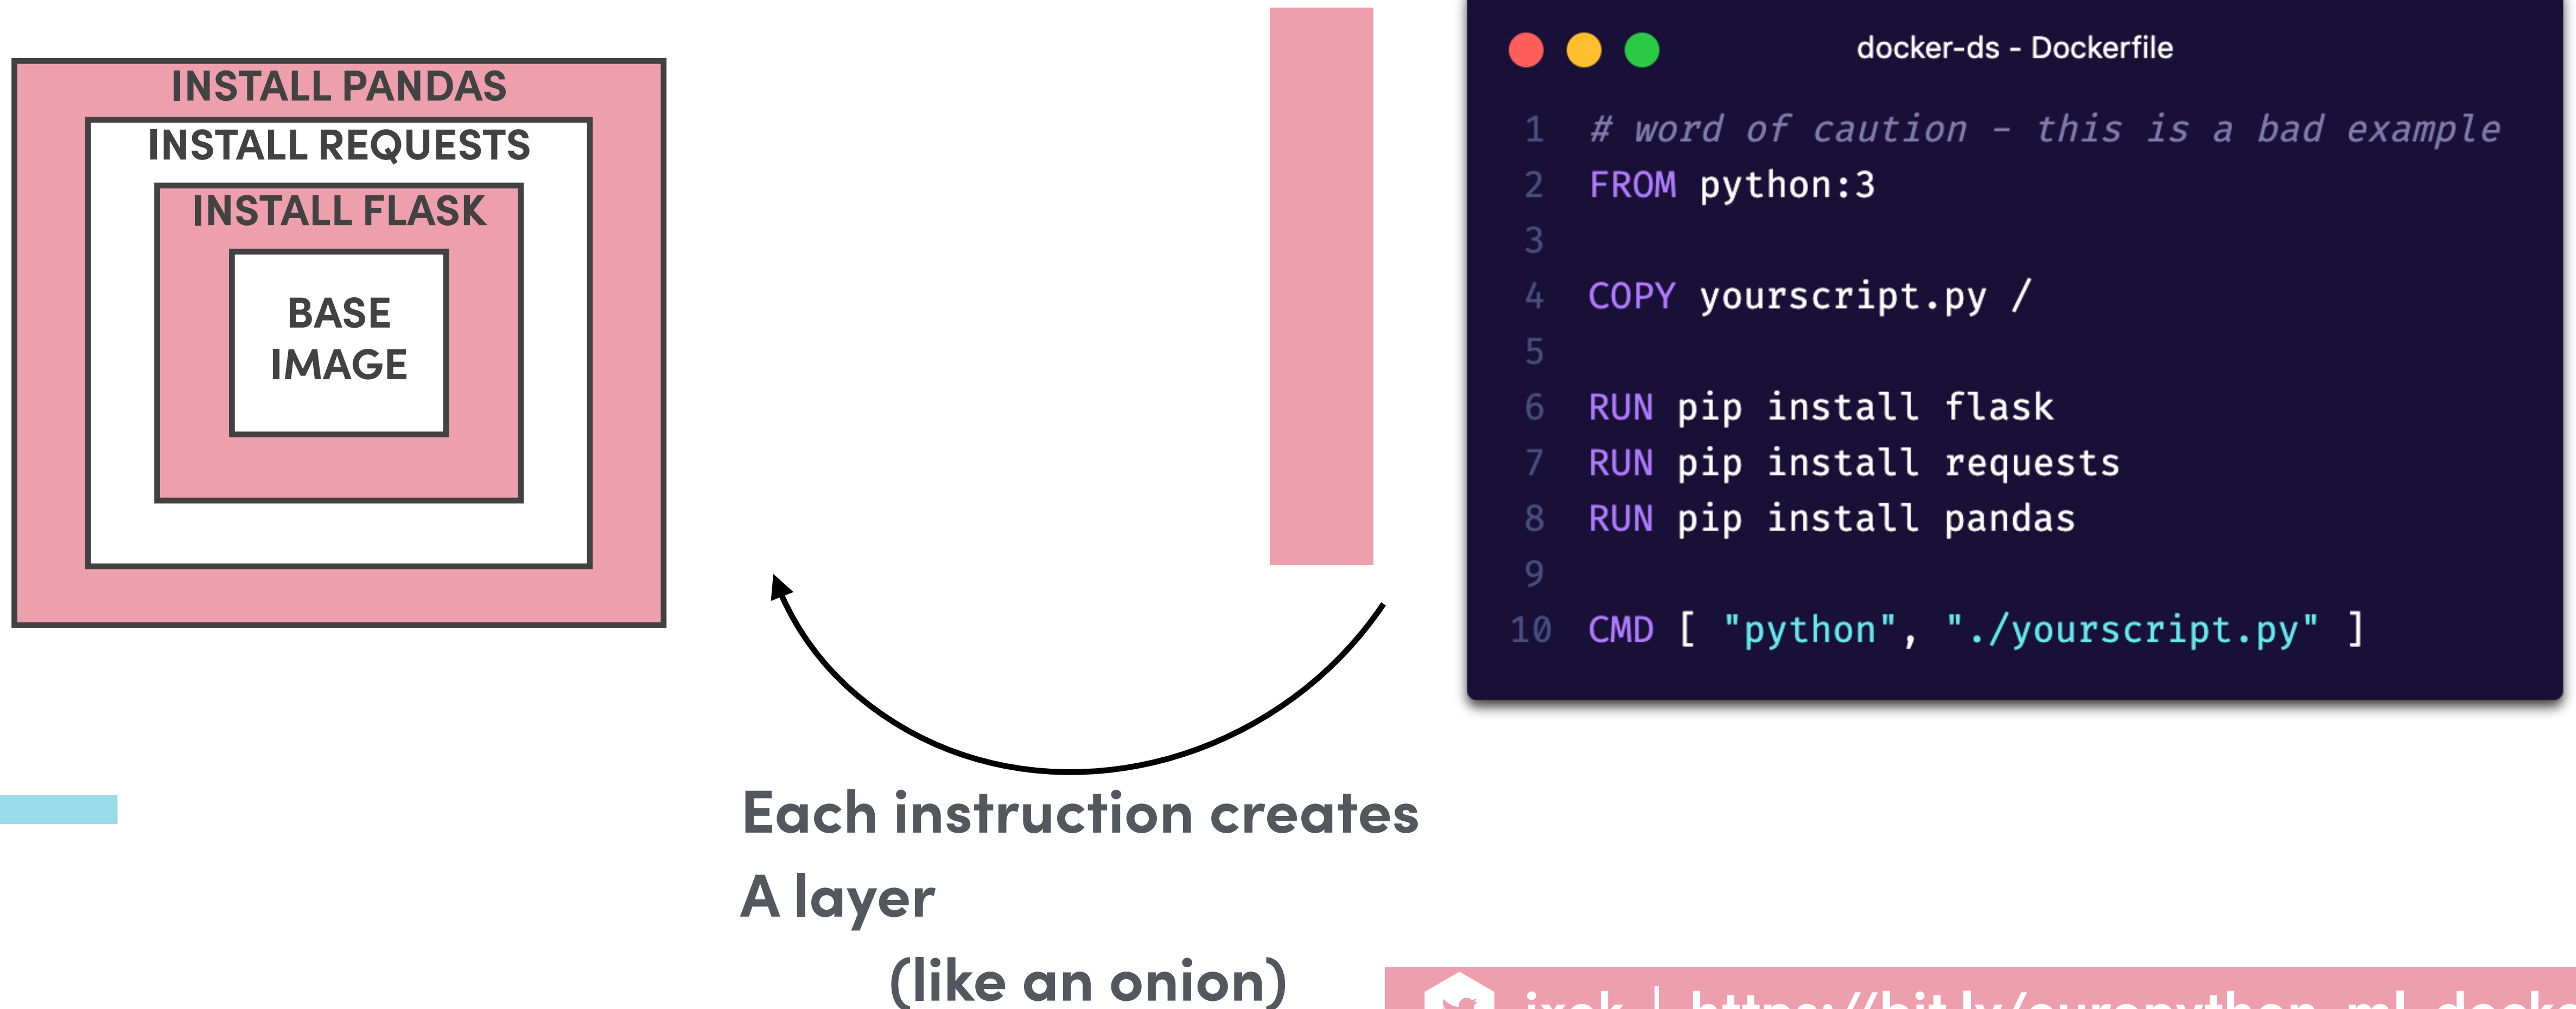


```
docker-ds - Dockerfile

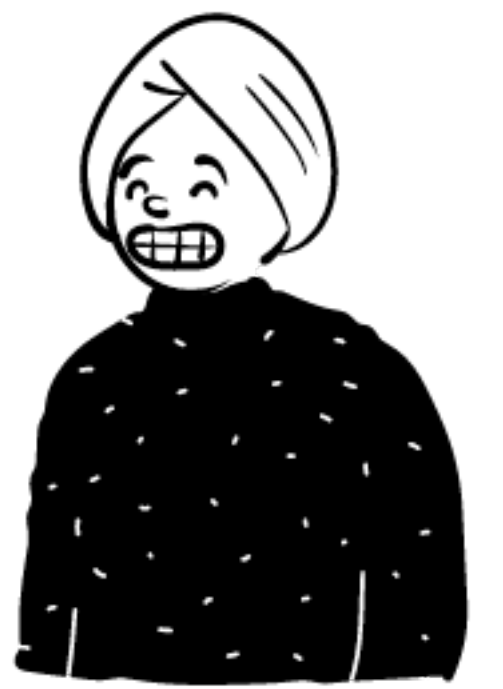
1  # word of caution - this is a bad example
2  FROM python:3
3
4  COPY yourscript.py /
5
6  RUN pip install flask
7
8  CMD [ "python", "./yourscript.py" ]
```



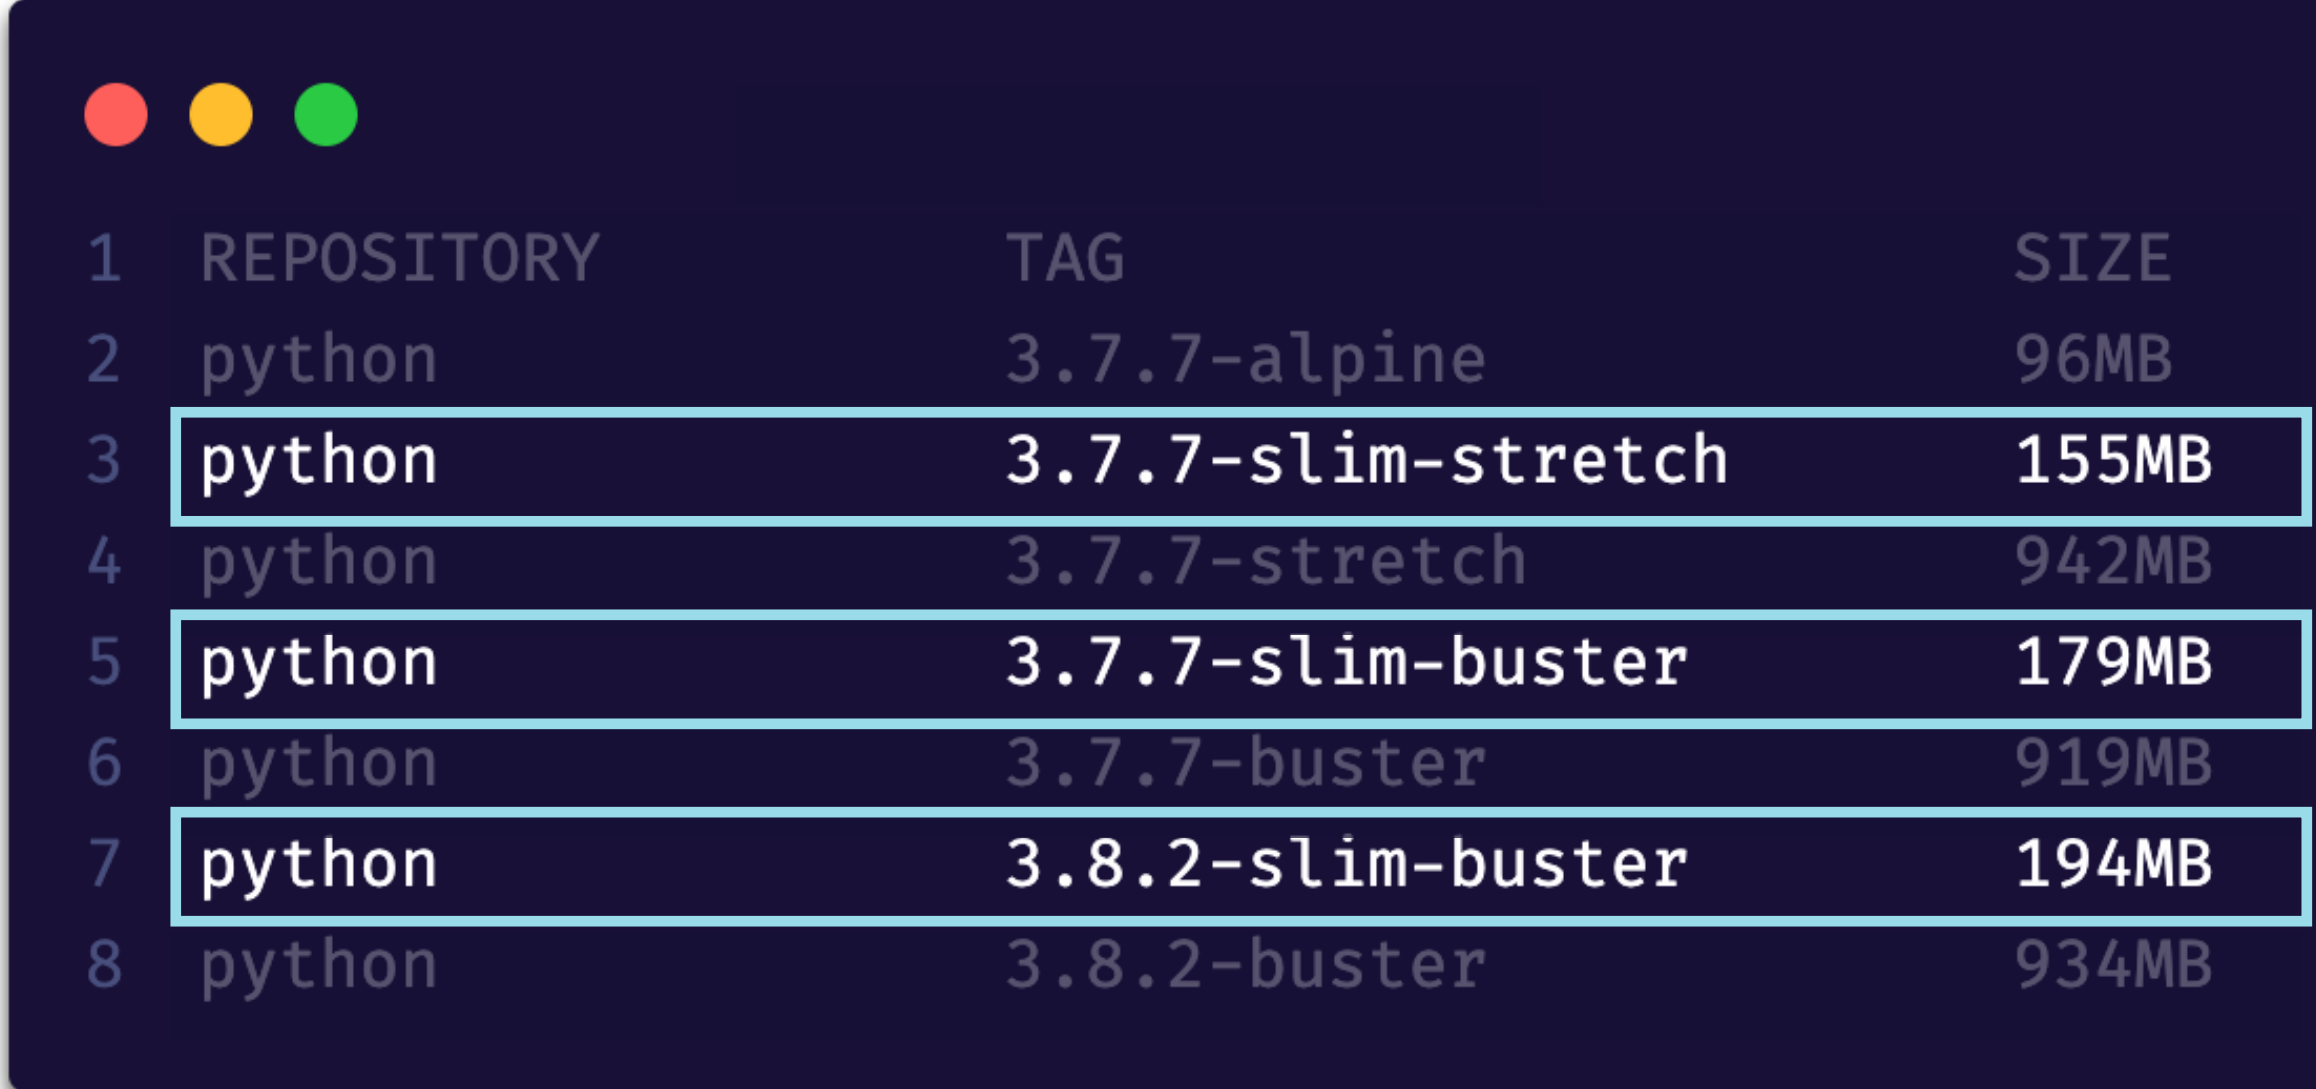
DISSECTING DOCKER IMAGES



CHOOSING THE BEST BASE IMAGE



If building from scratch use the
official Python images



	REPOSITORY	TAG	SIZE
1	python	3.7.7-alpine	96MB
2	python	3.7.7-slim-stretch	155MB
3	python	3.7.7-stretch	942MB
4	python	3.7.7-slim-buster	179MB
5	python	3.7.7-buster	919MB
6	python	3.8.2-slim-buster	194MB
7	python	3.8.2-buster	934MB
8			

https://hub.docker.com/_/python

<https://github.com/docker-library/docs/tree/master/python>




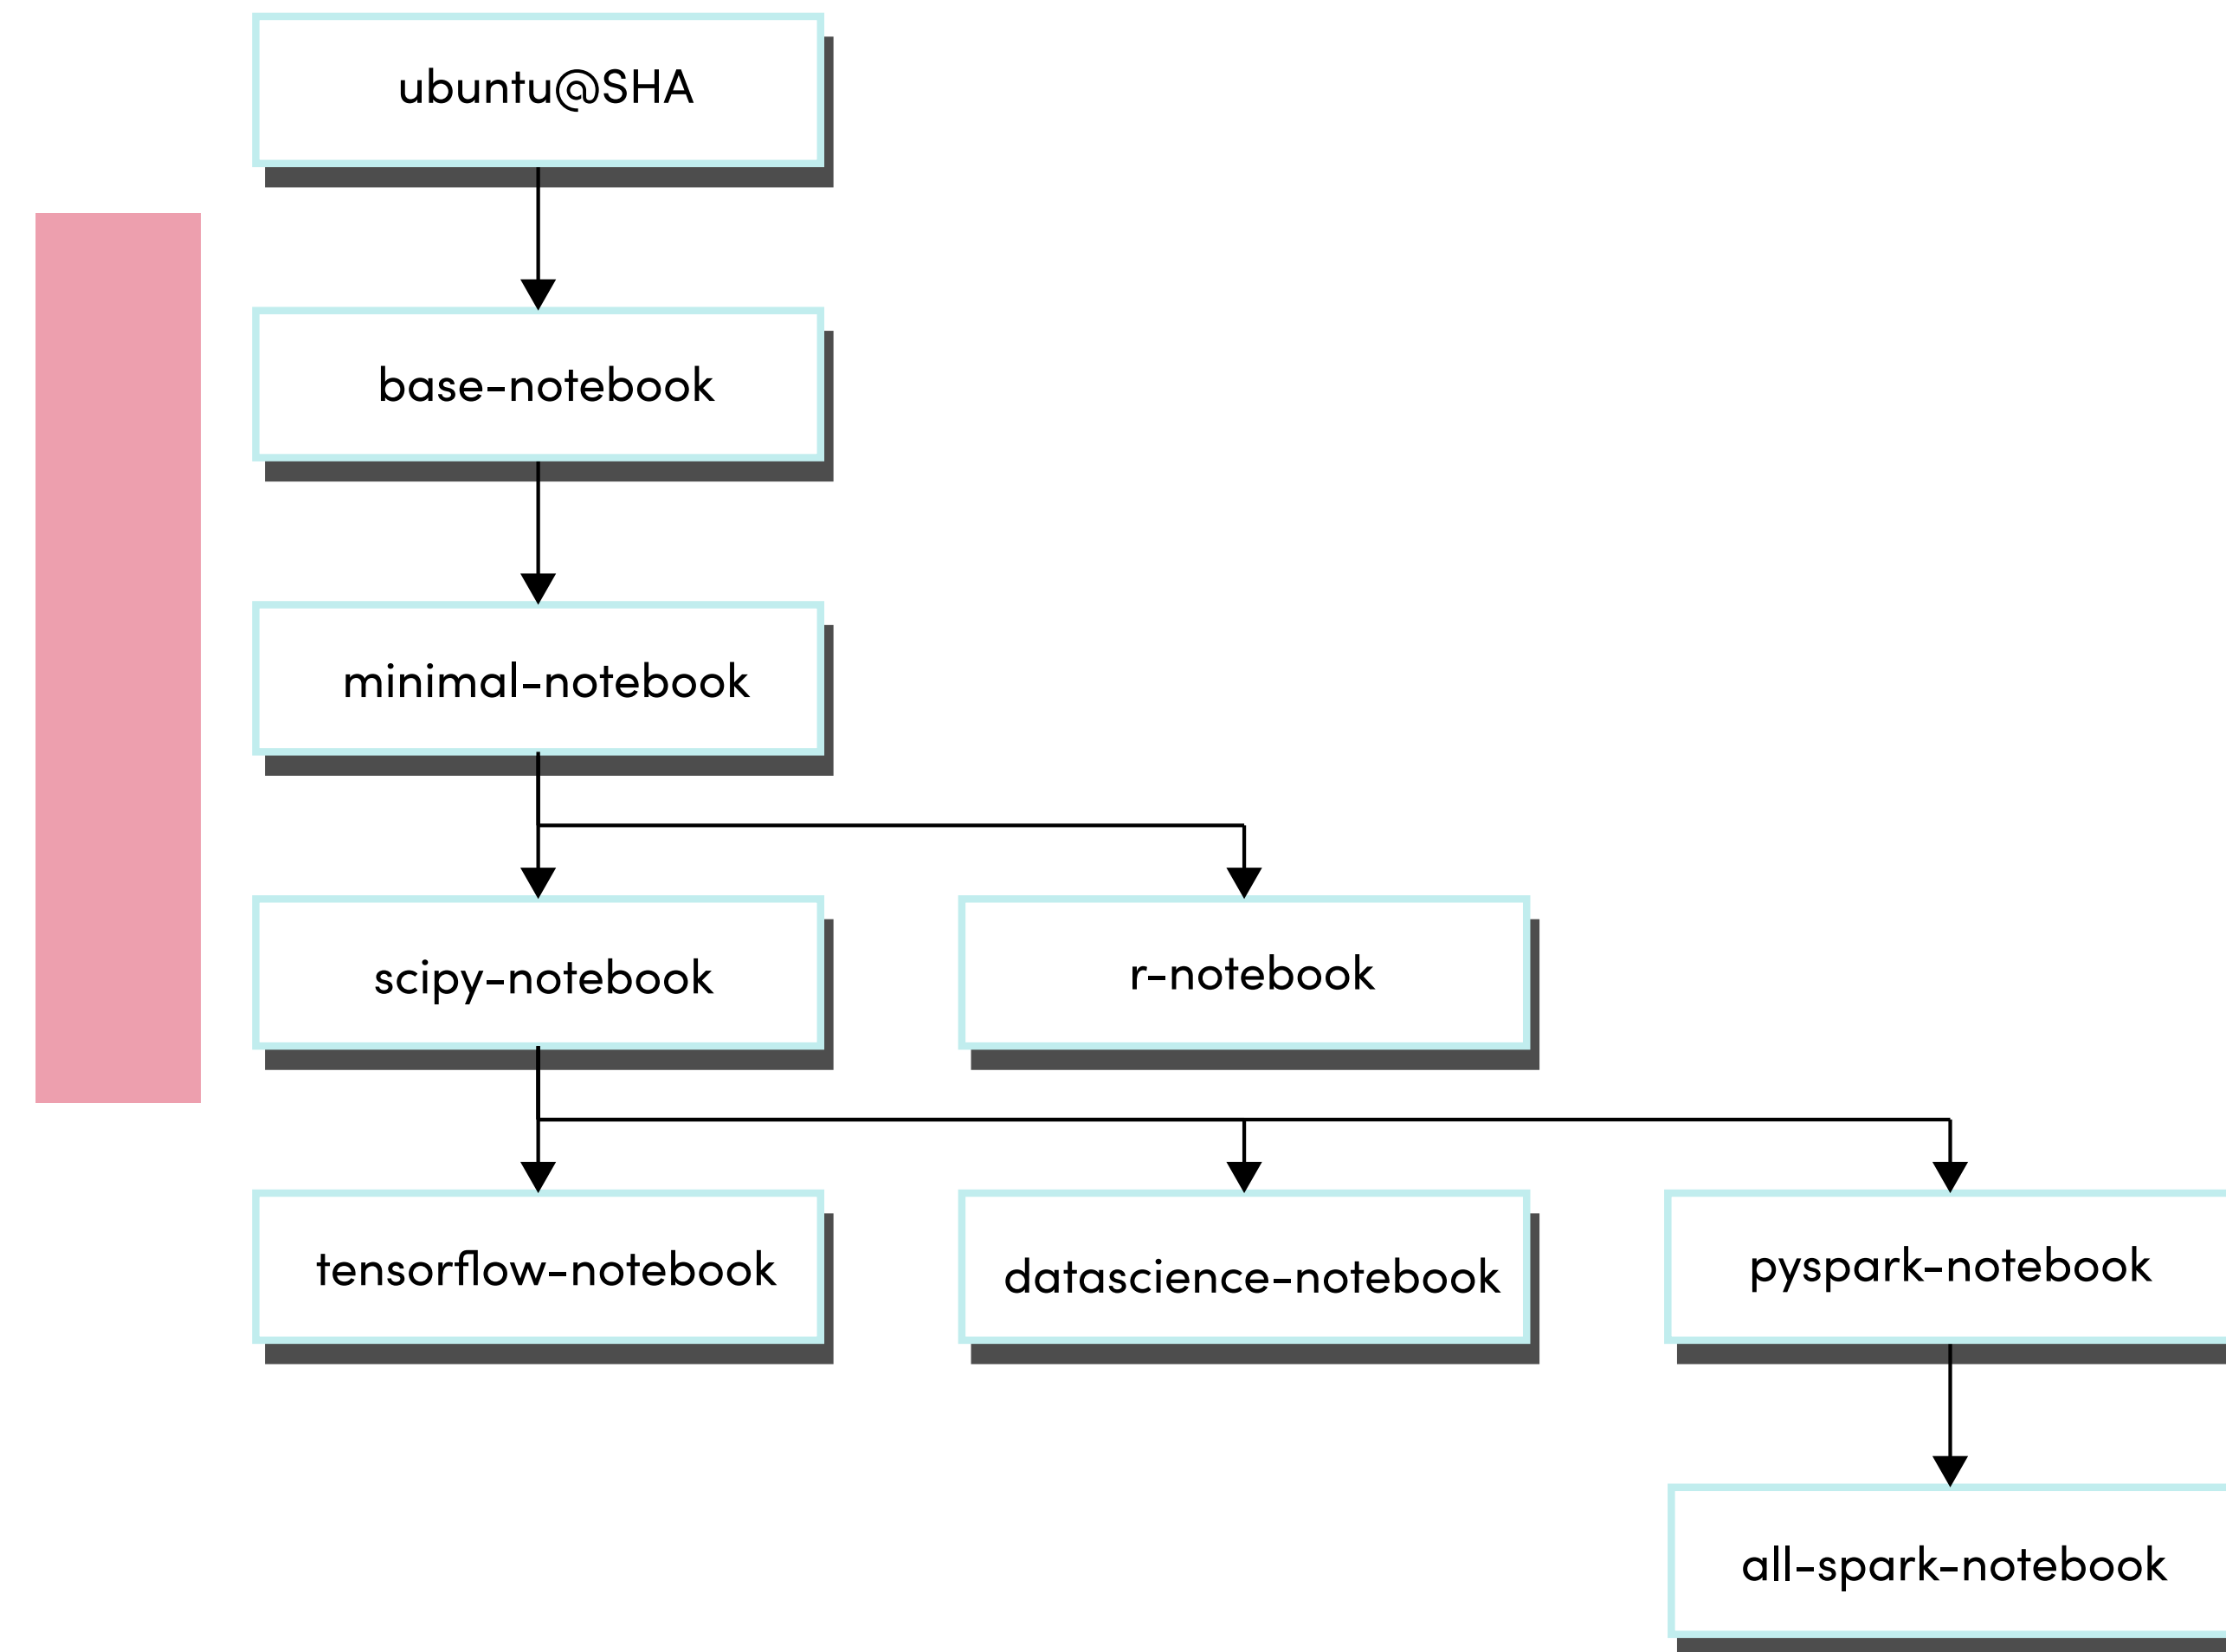
ixek | <https://bit.ly/euopython-ml-docker>

THE JUPYTER DOCKER STACK

Need Conda, notebooks and scientific Python ecosystem?

Try **Jupyter Docker stacks**

<https://jupyter-docker-stacks.readthedocs.io/>



BEST PRACTICES




- Always know what you are expecting
- Provide context with LABELS
- Split complex RUN statements and sort them
- Prefer COPY to add files

```
docker-ds - Dockerfile

1  # Always use a concrete tag (avoid LATEST)
2  FROM jupyter/base-notebook:6.0.3
3
4  # Add metadata
5  LABEL maintainer="Tania Allard"
6  LABEL securitytxt="https://www.example.com/.well-known/security.txt"
7
8  # Use pinned versions always
9  RUN conda install --quiet --yes \
10     'pandas=1.0.3' \
11     'dask=2.14.*' \
12     && \
13     # do not forget to clean - reduce image
14     conda clean --all -f -y
15
16 # separate instructions per scope
17 RUN mkdir data-sci-demo
18
19 COPY ./your-project data-sci-demo/
20
21
```

https://docs.docker.com/develop/develop-images/dockerfile_best-practices/



ixek | <https://bit.ly/europython-ml-docker>

SPEED UP YOUR BUILD

- Leverage build cache
- Install only necessary packages


docker-ds - requirements.txt

```
1 pandas=1.0.3
2 dask=2.14.*
```

```
docker-ds - Dockerfile

1  # Always use a concrete tag (avoid LATEST)
2  FROM jupyter/base-notebook:6.0.3
3
4  # Add metadata
5  LABEL maintainer="Tania Allard"
6  LABEL securitytxt="https://www.example.com/.well-known/security.txt"
7
8  # Leveraging build cache
9  COPY ./requirements.txt /tmp/
10
11 RUN conda install --quiet --yes --file /tmp/requirements.txt &&\
12     # do not forget to clean - reduce image
13     conda clean --all -f -y
14
15 # Separate instructions per scope
16 RUN mkdir data-sci-demo
17
18 COPY ./your-project data-sci-demo/
```

https://docs.docker.com/develop/develop-images/dockerfile_best-practices/



ixek | <https://bit.ly/europython-ml-docker>

SPEED UP YOUR BUILD AND PROOF

- Leverage build cache
- Install only necessary packages
- Explicitly ignore files


docker-ds - .dockerignore

```
1 # Documentation
2 Readme.md
3
4 # Never add data
5 ./yourproject/data/
6
7 # Secrets
8 appsettings.json
9 .env
10 supersecretkeys.json
```

docker-ds - Dockerfile

```
1 # Always use a concrete tag (avoid LATEST)
2 FROM jupyter/base-notebook:6.0.3
3
4 # Add metadata
5 LABEL maintainer="Tania Allard"
6 LABEL securitytxt="https://www.example.com/.well-known/security.txt"
7
8 # Leveraging build cache
9 COPY ./requirements.txt /tmp/
10
11 RUN conda install --quiet --yes --file /tmp/requirements.txt &&\
12     # do not forget to clean - reduce image
13     conda clean --all -f -y
14
15 # Separate instructions per scope
16 RUN mkdir data-sci-demo
17
18 COPY ./your-project data-sci-demo/
```

https://docs.docker.com/develop/develop-images/dockerfile_best-practices/



ixek | <https://bit.ly/europython-ml-docker>


MOUNT VOLUMES TO ACCESS DATA

- You can use bind mounts to directories (unless you are using a database)
- Avoid issues by creating a non-root user

```
docker-ds - workflow.yml
1  # mount directory
2  docker run --volume /home/user/yourproject:/yourproject mycontainer
3  # mount directory as read-only
4  docker run --volume /home/user/yourproject:/yourproject:ro mycontainer
5  # mount multiple directories, one with write access relative to current path
   (Linux)

6  docker run --volume /home/user/article-x-supplement/data:/data:ro \
7  --volume $(pwd)/outputs:/output-data:rw mycontainer
```

https://docs.docker.com/develop/develop-images/dockerfile_best-practices/



ixek | <https://bit.ly/europython-ml-docker>




SECURITY AND PERFORMANCE


MINIMISE PRIVILEGE - FAVOUR LESS PRIVILEGED USER

Lock down your container:

- Run as non-root user (Docker runs as root by default)
- Minimise capabilities



```
docker-ds - Dockerfile  
  
1 FROM python:3.8.2-slim-buster  
2 RUN useradd --create-home jovyan  
3 WORKDIR /home/jovyan  
4 USER jovyan
```




DON'T LEAK SENSITIVE INFORMATION

Remember Docker images are like onions. If you copy keys in an intermediate layer **they are cached**.



Keep them out of your **Dockerfile**.



USE MULTI STAGE BUILDS

- Fetch and manage secrets in an intermediate layer
- Not all your dependencies will have been packed as wheels so you might need a compiler – build a compile and a runtime image
- Smaller images overall

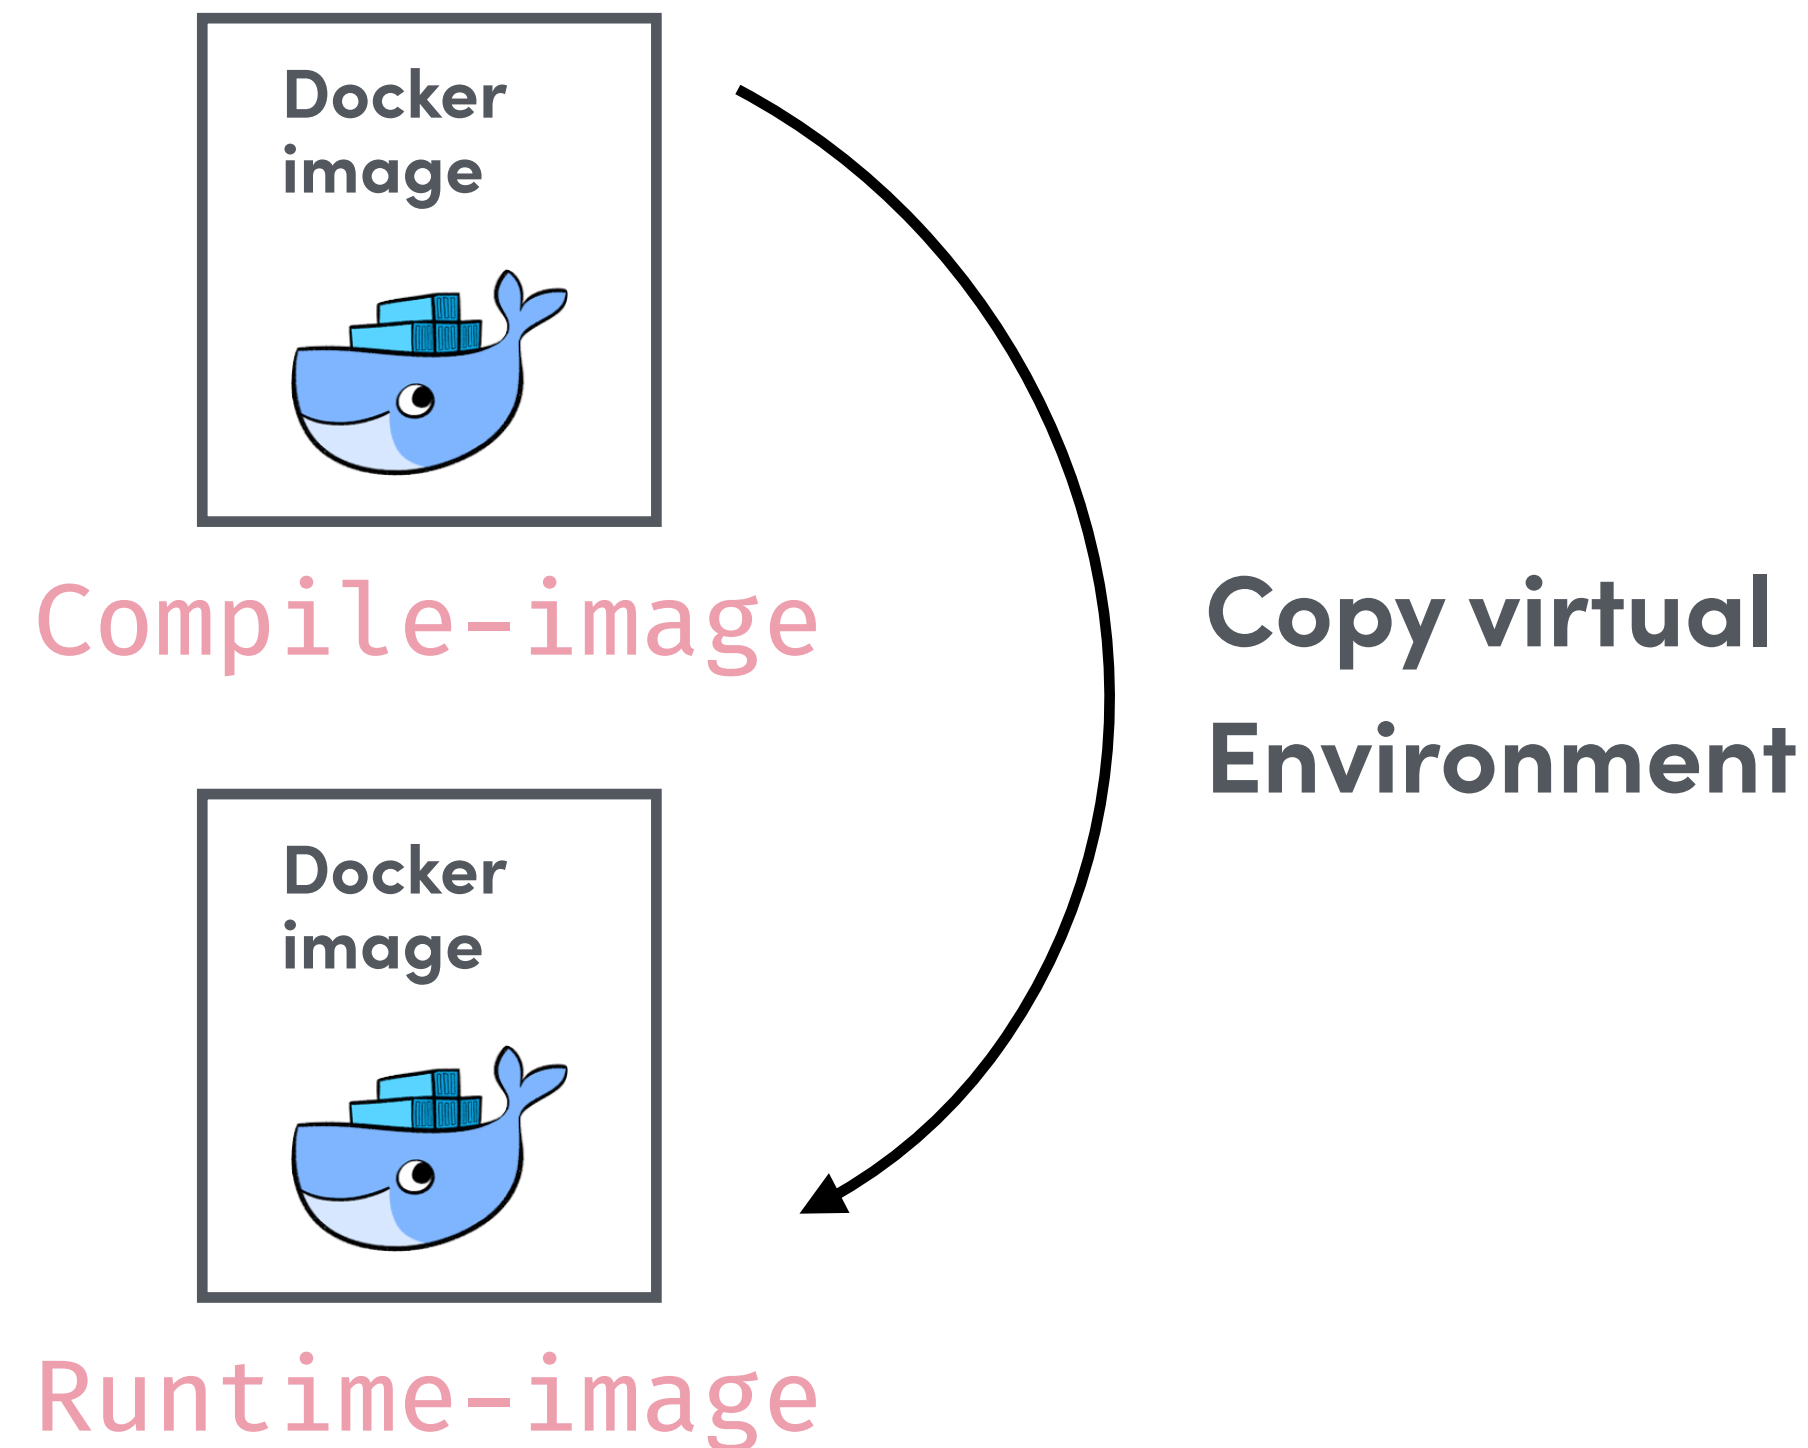


```
docker-ds - Dockerfile

1  # Always use a concrete tag (avoid LATEST)
2  FROM python:3.8.2-slim-buster as compile-image
3
4  # Add metadata
5  LABEL maintainer="Tania Allard"
6  LABEL securitytxt=
   "https://www.example.com/.well-known/security.txt"
7
8  RUN apt-get update
9  RUN
   apt-get install -y --no-install-recommends gcc build-essen
   tial gcc gfortran
10
11 RUN python -m venv /opt/venv
12
13 # Ensure we use the virtualenv
14 ENV PATH="/opt/venv/bin:$PATH"
15
16 COPY requirements.txt /tmp/
17
18 RUN CFLAGS=
   "-g0 -Wl,--strip-all -I/usr/include:/usr/local/include -L/u
   sr/lib:/usr/local/lib"
   \
19     pip install \
20     --no-cache-dir \
21     --compile \
22     --global-option=build_ext \
23     --global-option="-j 4" \
24     -r /tmp/requirements.txt
25
26 # -----
27 # This is the second image that copies the compiled librar
   y
28
29 FROM python:3.8.2-slim-buster as runtime-image
30
31 COPY --from=compile-image /opt/venv /opt/venv
32 # Ensure we use the virtualenv
33 ENV PATH="/opt/venv/bin:$PATH"
```


USE MULTI STAGE BUILDS

```
$ docker build --pull --rm -f "Dockerfile" \
-t trallard:data-scratch-1.0 "."
```




```
docker-ds - Dockerfile

1  # Always use a concrete tag (avoid LATEST)
2  FROM python:3.8.2-slim-buster as compile-image
3
4  # Add metadata
5  LABEL maintainer="Tania Allard"
6  LABEL securitytxt=
   "https://www.example.com/.well-known/security.txt"
7
8  RUN apt-get update
9  RUN
   apt-get install -y --no-install-recommends gcc build-essen
   tial gcc gfortran
10
11 RUN python -m venv /opt/venv
12
13 # Ensure we use the virtualenv
14 ENV PATH="/opt/venv/bin:$PATH"
15
16 COPY requirements.txt /tmp/
17
18 RUN CFLAGS=
   "-g0 -Wl,--strip-all -I/usr/include:/usr/local/include -L/u
   sr/lib:/usr/local/lib"
   \
19     pip install \
20     --no-cache-dir \
21     --compile \
22     --global-option=build_ext \
23     --global-option="-j 4" \
24     -r /tmp/requirements.txt
25
26 # -----
27 # This is the second image that copies the compiled librar
   y
28
29 FROM python:3.8.2-slim-buster as runtime-image
30
31 COPY --from=compile-image /opt/venv /opt/venv
32 # Ensure we use the virtualenv
33 ENV PATH="/opt/venv/bin:$PATH"
```

USE MULTI STAGE BUILDS

FINAL IMAGE



Runtime-image

trallard:data-scratch-1.0

```
docker-ds - Dockerfile

1  # Always use a concrete tag (avoid LATEST)
2  FROM python:3.8.2-slim-buster as compile-image
3
4  # Add metadata
5  LABEL maintainer="Tania Allard"
6  LABEL securitytxt=
   "https://www.example.com/.well-known/security.txt"
7
8  RUN apt-get update
9  RUN
   apt-get install -y --no-install-recommends gcc build-essen
   tial gcc gfortran
10
11  RUN python -m venv /opt/venv
12
13  # Ensure we use the virtualenv
14  ENV PATH="/opt/venv/bin:$PATH"
15
16  COPY requirements.txt /tmp/
17
18  RUN CFLAGS=
   "-g0 -Wl,--strip-all -I/usr/include:/usr/local/include -L/u
   sr/lib:/usr/local/lib"
   \
19     pip install \
20     --no-cache-dir \
21     --compile \
22     --global-option=build_ext \
23     --global-option="-j 4" \
24     -r /tmp/requirements.txt
25
26  # -----
27  # This is the second image that copies the compiled librar
   y
28
29  FROM python:3.8.2-slim-buster as runtime-image
30
31  COPY --from=compile-image /opt/venv /opt/venv
32  # Ensure we use the virtualenv
33  ENV PATH="/opt/venv/bin:$PATH"
```



AUTOMATE

PROJECT TEMPLATES



Need a standard project template?

Use cookie cutter data science

Or cookie cutter docker science

docker-ds - workflow.yml	
1 — Makefile	← Makefile contains many targets such as create docker container or get input files.
2	
3 — config	← This directory contains configuration files used in scripts or Jupyter Notebook.
4 — jupyter_config.py	
5	
6 — data	← data directory contains the input resources.
7 — docker	← docker directory contains Dockerfile.
8 — Dockerfile	← Dockerfile have the container settings. Users modify Dockerfile if additional library is needed for experiments.
9	
10 — model	← model directory store the model files created in the experiments.
11 — my_data_science_project	← cookie-cutter-docker-science creates the directory whose name is same as project name. In this directory users puts python files used in scripts or Jupyter Notebook.
12 — __init__.py	
13	
14 — notebook	← This directory stores the ipynb files saved in Jupyter Notebook.
15 — requirements.txt	← Libraries needed to run experiments. The library listed in this file are installed in the Docker container.
16	
17 — scripts	← Users add the script files to generate model files or run evaluation.
18	

<https://github.com/docker-science/cookiecutter-docker-science>

<https://drivendata.github.io/cookiecutter-data-science/>

```
$ conda install jupyter repo2docker
$ jupyter-repo2docker “.”
```




DO NOT REINVENT THE WHEEL

Leverage the existence and usage
of tools like **repo2docker**.

Already configured and optimised
for Data Science / Scientific
computing.

<https://repo2docker.readthedocs.io/en/latest>



ixek | <https://bit.ly/europython-ml-docker>

DO NOT REINVENT THE WHEEL

Leverage the existence and usage of tools like **repo2docker**.

Already configured and optimised for Data Science / Scientific computing.

- Configuration Files
 - **environment.yml** - Install a Python environment
 - **Pipfile** and/or **Pipfile.lock** - Install a Python environment
 - **requirements.txt** - Install a Python environment
 - **setup.py** - Install Python packages
 - **Project.toml** - Install a Julia environment
 - **REQUIRE** - Install a Julia environment (legacy)
 - **install.R** - Install an R/RStudio environment
 - **apt.txt** - Install packages with apt-get
 - **DESCRIPTION** - Install an R package
 - **manifest.xml** - Install Stencila
 - **postBuild** - Run code after installing the environment
 - **start** - Run code before the user sessions starts
 - **runtime.txt** - Specifying runtimes
 - **default.nix** - the nix package manager
 - **Dockerfile** - Advanced environments

DELEGATE TO YOUR CONTINUOUS INTEGRATION TOOL

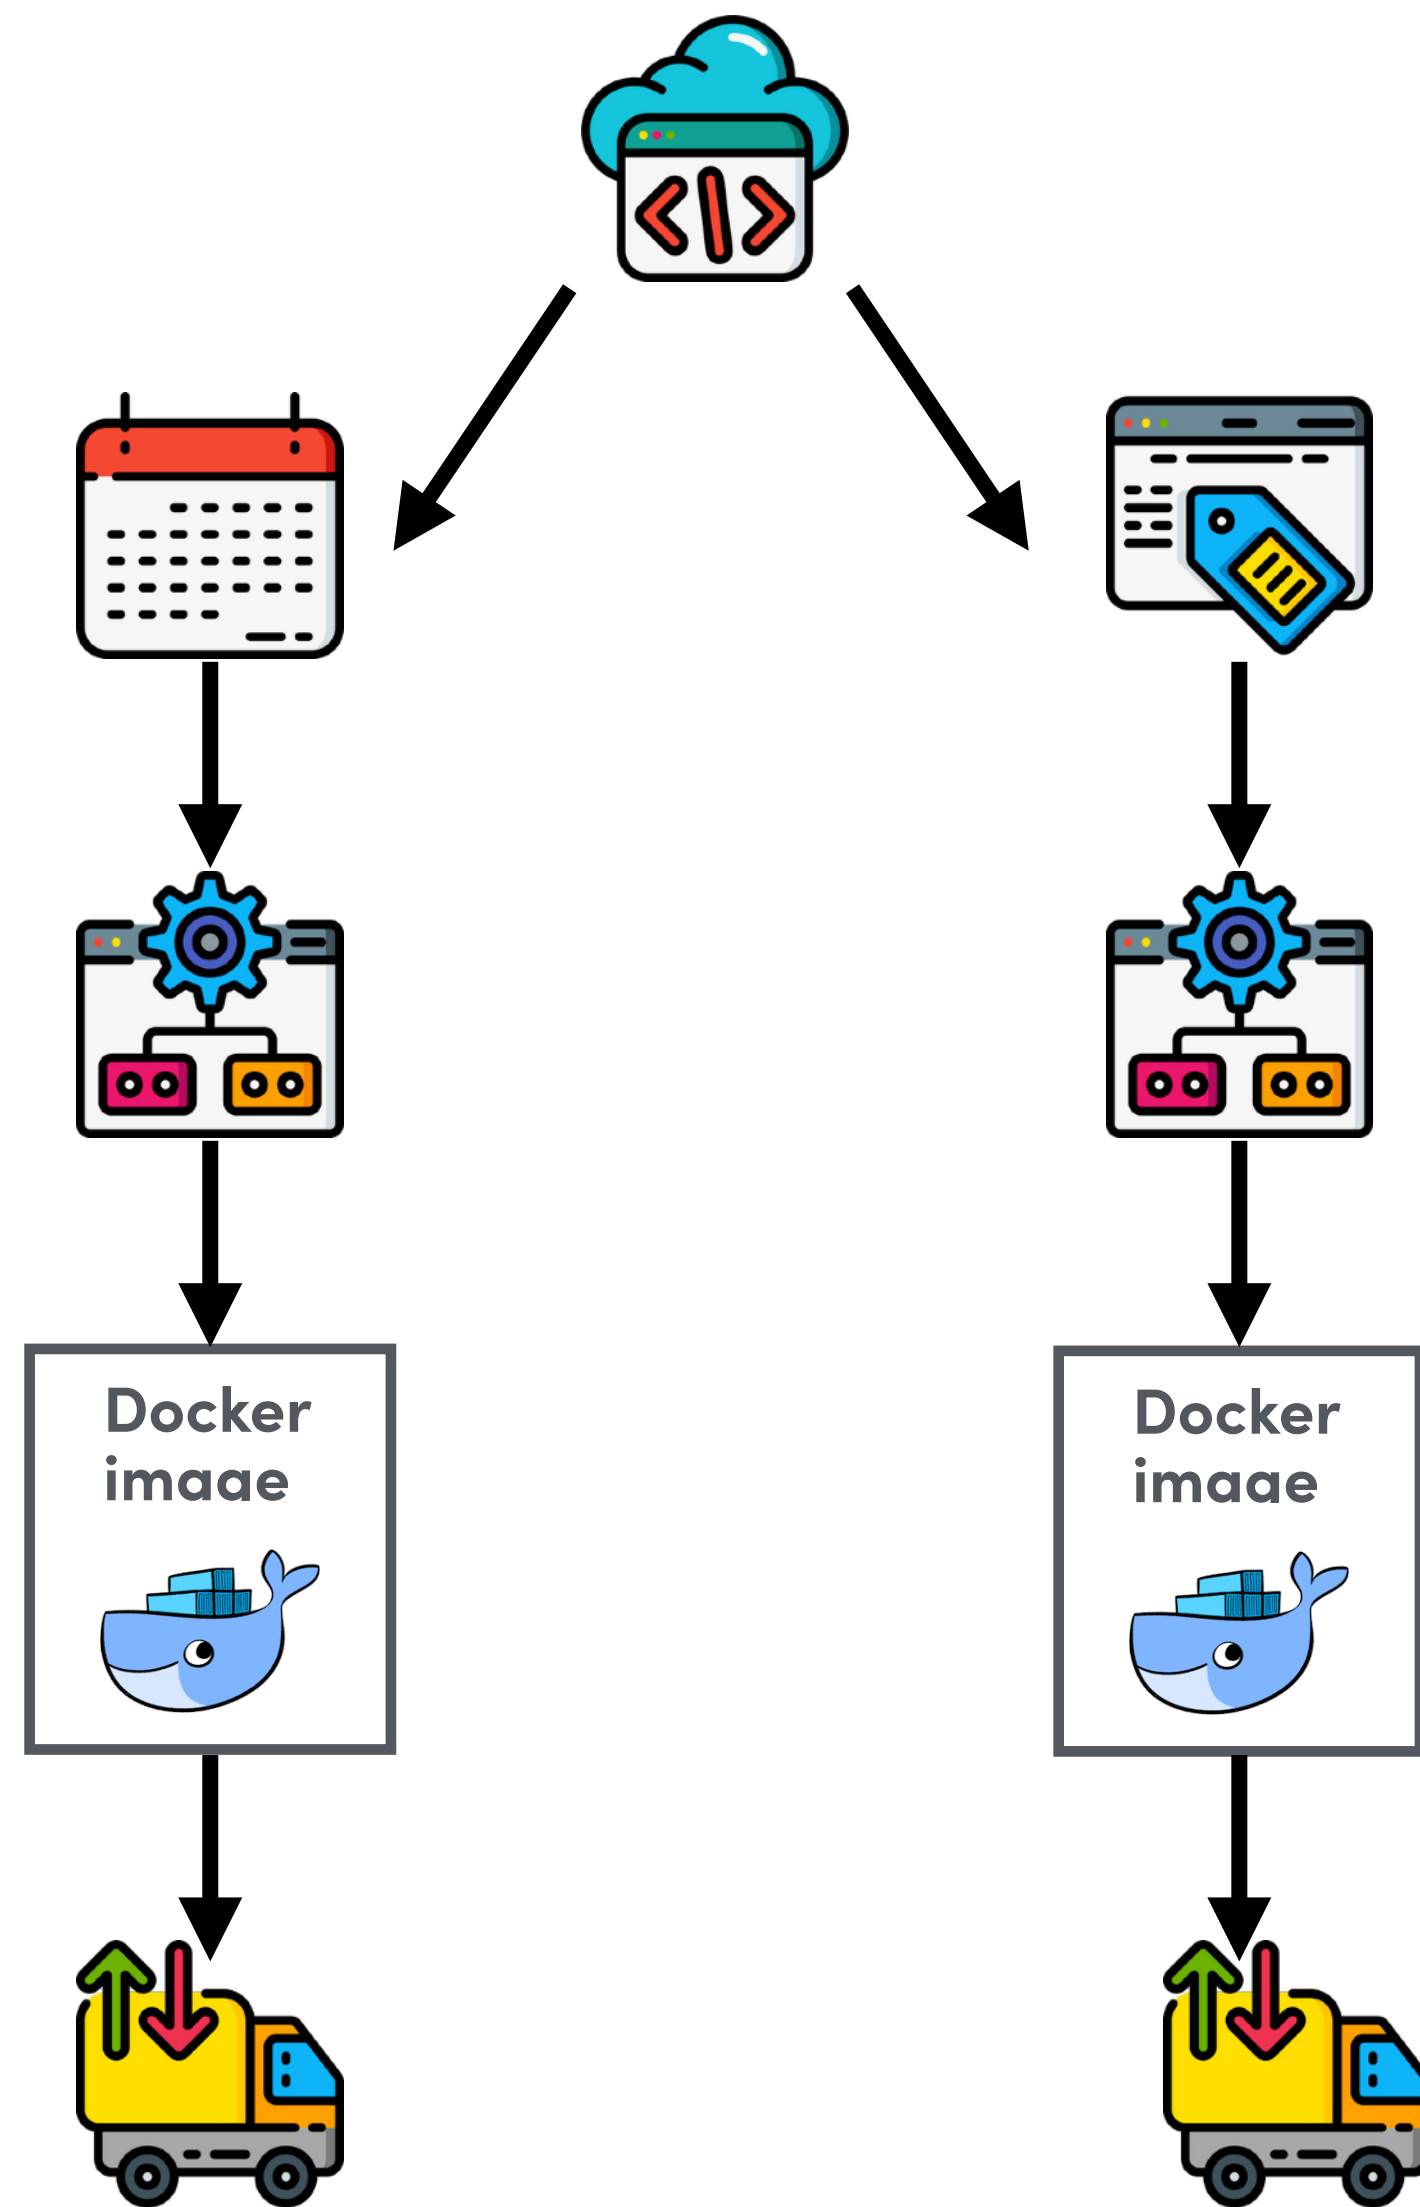
Set Continuous integration
(Travis, GitHub Actions, whatever
you prefer).

And delegate your build – also
build often.

docker-ds - workflow.yml

```
1 name: Publish to Registry
2 on:
3   release:
4     types: [published]
5   schedule:
6     # Build your images frequently
7     - cron: "0 2 * * 0" # Weekly on Sundays at 02:00
8 jobs:
9   update:
10    runs-on: ubuntu-latest
11    steps:
12      - uses: actions/checkout@master
13      - name: Get release version
14        id: get_version
15        run: echo ::set-env name=RELEASE_VERSION::$(echo ${GITHUB_REF:10})
16      - name: Build and publish
17        uses: docker/build-push-action@v1
18        with:
19          username: ${ secrets.DOCKER_USERNAME }
20          password: ${ secrets.DOCKER_PASSWORD }
21          repository: myorg/myrepository
22          tag_with_ref: true
23          tag_with_sha: true
24
```


THIS WORKFLOW



- Code in version control
- Trigger on tag / Also scheduled trigger
- Build image
- Push image





TOP TIPS

TOP TIPS


1. Rebuild your images frequently – get security updates for system packages
2. Never work as root / minimise the privileges
3. You do not want to use Alpine Linux (go for buster, stretch or the Jupyter stack)
4. Always know what you are expecting: pin / version **EVERYTHING** (use pip-tools, conda, poetry or pipenv)
5. Leverage build cache

TOP TIPS

6. Use one Dockerfile per project
7. Use multi-stage builds - need to compile code? Need to reduce your image size?
8. Make your images identifiable (test, production, R&D) - also be careful when accessing databases and using ENV variables / build variables
9. Do not reinvent the wheel! Use repo2docker
10. Automate - no need to build and push manually
11. Use a linter



THANK YOU



@ixek



@trallard



trallard.dev