

I am a researcher, and I work with AI every day. I can argue that everyone in my position is excited like a dog that stares at an ice cream cone. 🍦

This is the reason:

Open AI's Chat GPT is awesome.

For those who don't know what I'm talking about, Chat GPT is an artificial intelligence chatbot that can do, well, pretty much everything.

It can code, it can write articles, it can help you decorate your home, it can make up a recipe (I don't recommend that if you are Italian), and the list goes on.

Yes, we can argue that it will cause ethical (and not only ethical) problems in the future. My mother is a high school teacher, and she is **terrified** about the idea that her students will use Chat GPT to cheat on their tests, and this is just one of the many examples of how things can go wrong with this incredibly powerful technology.

But again, the problem is the **use**, not the **product**. If we strictly talk about the **technological aspect** (that is, frankly, the one I am more interested in, as I am a certified nerd), it is **freaking amazing**.

Now, a lot of developers have used and tested this chatbot to try and develop their codes and their AI ideas, and of course, the usage of this chatbot **strictly depends on your background**. For example, if you are a web developer, you would ask ChatGPT to build a website using HTML. If you are a tester, you could ask ChatGPT to help you find that bug in that specific system.

In my specific case, I am a **researcher**. In particular, what I do for a living is build some **surrogate models** using AI. Let's say that you want to conduct research on "A," but to do "A," you need a **lot of money**, a **lot of power**, and a **lot of computational time**. The idea behind this surrogate model is to replace it with a data-driven approach using artificial intelligence.

Now let's completely change the subject for a moment.

Let's say I am an **entrepreneur**, and I have tons of hotels all over the USA. Given a certain review of a given hotel, I want to know if that review is a good one or a bad one for that hotel. How do I do that? I have three options:

1. I hire a person that reads millions of reviews and classify them every day and I probably get arrested because it is clearly an abuse of human rights.
2. I hire a person that, among other things reads *hundreds* of reviews and classifies them every day. After months, I am able to build a dataset out of this. I train a Machine Learning model out of that dataset.
3. I automatically generate good and bad reviews. I instantly build a dataset out of this. I train a Machine Learning model out of that dataset.

As I value the time of my reader, let me skip the first option.

The second option is what you would do **before ChatGPT**. You can't know in advance if a review is bad or good, so if you want to build a dataset out of this, you need to hire people and wait until the dataset is ready.

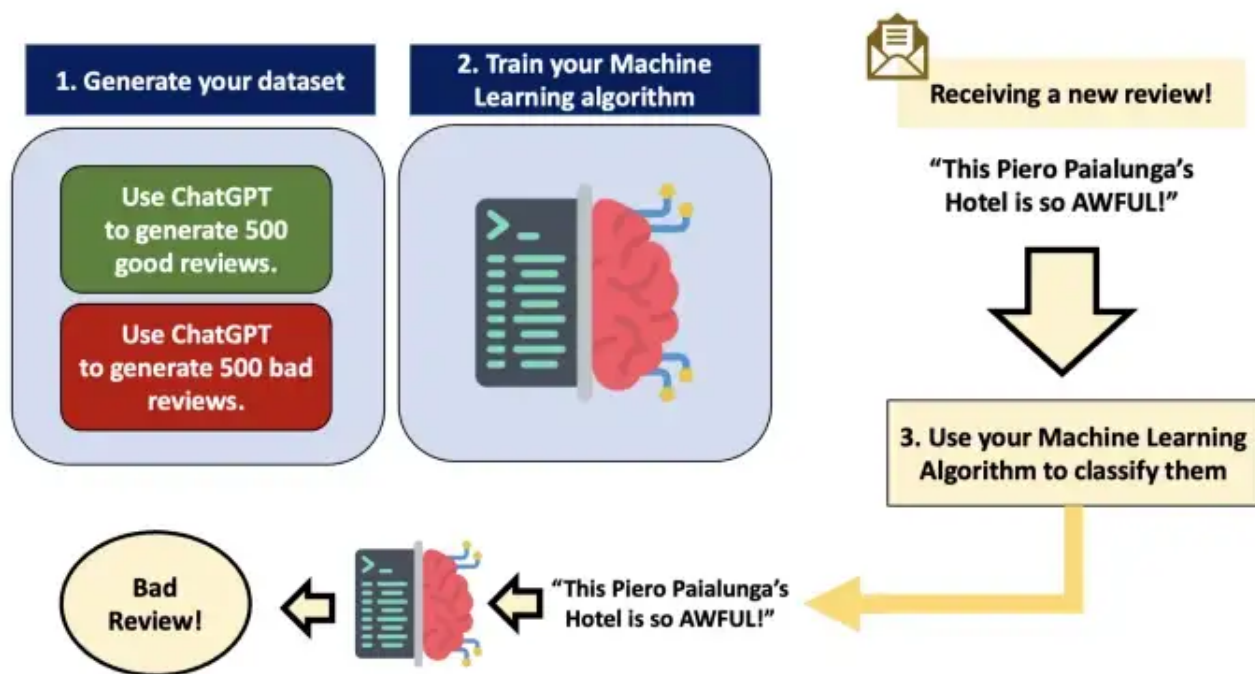
Now that we have ChatGPT, though, we can simply ask them to generate **good and bad reviews!** This would take minutes (rather than months), and it will allow us to build our machine learning algorithm to automatically classify our customer reviews!

Congratulations, this is your first **Surrogate Model**. 😊

Keep in mind that **we will not train ChatGPT or do any fine-tuning**. This model is exceptional for a task like this, and no fine-tuning is required in this instance. Now, the training of the ChatGPT model is, of course, not open source (just like the model itself). All we know is the small description that is in the [Open AI's blog](#). They explain that the model is trained by human AI trainers and a reinforcement-learning supervised algorithm.*

The fact alone that OpenAI's ChatGPT is not open source raises some very tricky and interesting **ethical questions. Should such a powerful model be open source, so that everyone (bad people too) can use it, or should it be not open source so that no one can really trust it?*

Let me recap:



Credits: [Smashicons](#)

Image by author

The little brain-shell thing that you see is the **surrogate model**; as we will see, it will be a **random forest**. But I said it was a **hands-on** article, so let's dive in! (so excited!!!)

**I am sorry, I love spoilers.*

1. Generating our dataset

The first step is to use Open's AI Python API to generate our simulations.

A few things to consider:

1. Open AI is made by geniuses for non-genius users. For this reason, if you want to install it, you just have to do:

```
pip install --upgrade openai
```

(that is LOVABLE)

2. Of course, if you want to send a lot of requests, you will have to pay for a premium service. Assuming we don't want to do that, we'd just have to wait **around 30 minutes** to get our dataset of fake reviews. Again, this is **nothing** compared to the waiting time (and cost) of **months** that we'd have to wait if we did this manually. You'll also have to log in to Open AI and get your Open AI key.

3. We will automatically input whether this is a good review or a bad review by **starting** with the same sentence: “This hotel was terrible.” for a bad review and “This hotel was great.” for a good review. ChatGPT will complete our review for us. Of course, except for these first four words, which we won’t include in our reviews anyway, the rest of the review will be different.

Let me give you an example of a bad review:

```
import openai
import time
openai.api_key = "your_key"

completion = openai.Completion.create(engine="davinci", prompt="This hotel was ter
print(completion.choices[0]['text'])
```

There were broken chairs everywhere and even one broken elevator. The hallways smell

The concier

Hosted on  Deepnote

And of a good one:

```
completion = openai.Completion.create(engine="davinci", prompt="This hotel was great  
print(completion.choices[0]['text'])
```

The breakfast food was excellent and the staff was super friendly. Price was great al

Amazing room, free breakfast and excellent service. Could not in any way find fault wi

We enjoyed the staff. Laste was very helpful during my room change. We will return. -

The Wyndham team worked very hard in replacing our room with ease and professionalism

Global we had an

Hosted on  Deepnote

Now this is the code you will need to generate **your whole dataset**.

```

good_reviews = []
bad_reviews = []
for i in range(0,500):
    completion = openai.Completion.create(engine="davinci", prompt="This hotel was g
    good_reviews.append(completion.choices[0]['text'])
    print('Generating good review number %i'%(i))
    completion = openai.Completion.create(engine="davinci", prompt="This hotel was t
    bad_reviews.append(completion.choices[0]['text'])
    print('Generating bad review number %i'%(i))
    display = np.random.choice([0,1],p=[0.7,0.3])
    time.sleep(3)
    if display ==1:
        display_good = np.random.choice([0,1],p=[0.5,0.5])
        if display_good ==1:
            print('Printing random good review')
            print(good_reviews[-1])
        if display_good ==0:
            print('Printing random bad review')
            print(bad_reviews[-1])

```

Printing random bad review

The bathrooms smelled like mold and the fans shuddered making it next to impossible

Generating good review number 491

Generating bad review number 491

Generating good review number 492

Generating bad review number 492

Generating good review number 493

Generating bad review number 493

Generating good review number 494

Generating bad review number 494

Generating good review number 495

Generating bad review number 495

Printing random bad review

It is not new as advertised and they are renovating as you are there in the main pa

Generating good review number 496

Generating bad review number 496

Printing random bad review

It's apparent that the hotel is trying to save money. The rooms are so small that t

Nicole on

Generating good review number 497

Generating bad review number 497

Generating good review number 498

Generating bad review number 498

Generating good review number 499

Generating bad review number 499

Printing random bad review

I'll just say that our evening entertainment was terrible. This hotel was chosen be

Sam Iwasra Switzerland, - 17 March 2018

We will then store everything in a **dataframe** using **Pandas**.

1. Import pandas and build df:

```
import pandas as pd
import numpy as np
df = pd.DataFrame(np.zeros((1000,2)))
df.columns = ['Reviews', 'Sentiment']
df['Sentiment'].loc[0:499] = 1
```

Hosted on  Deepnote

2. Fill df:

```
df['Reviews'] = good_reviews+bad_reviews
```

Hosted on  Deepnote

3. Export df:

```
df.to_csv('generated_reviews.csv')
```

Hosted on  Deepnote

2. Machine Learning time!

So now we need to build and train a machine learning algorithm.

As we are dealing with **texts**, the first thing that we need to do is use a **vectorizer**. A vectorizer is something that transforms a **text** into a **vector**.

Like:

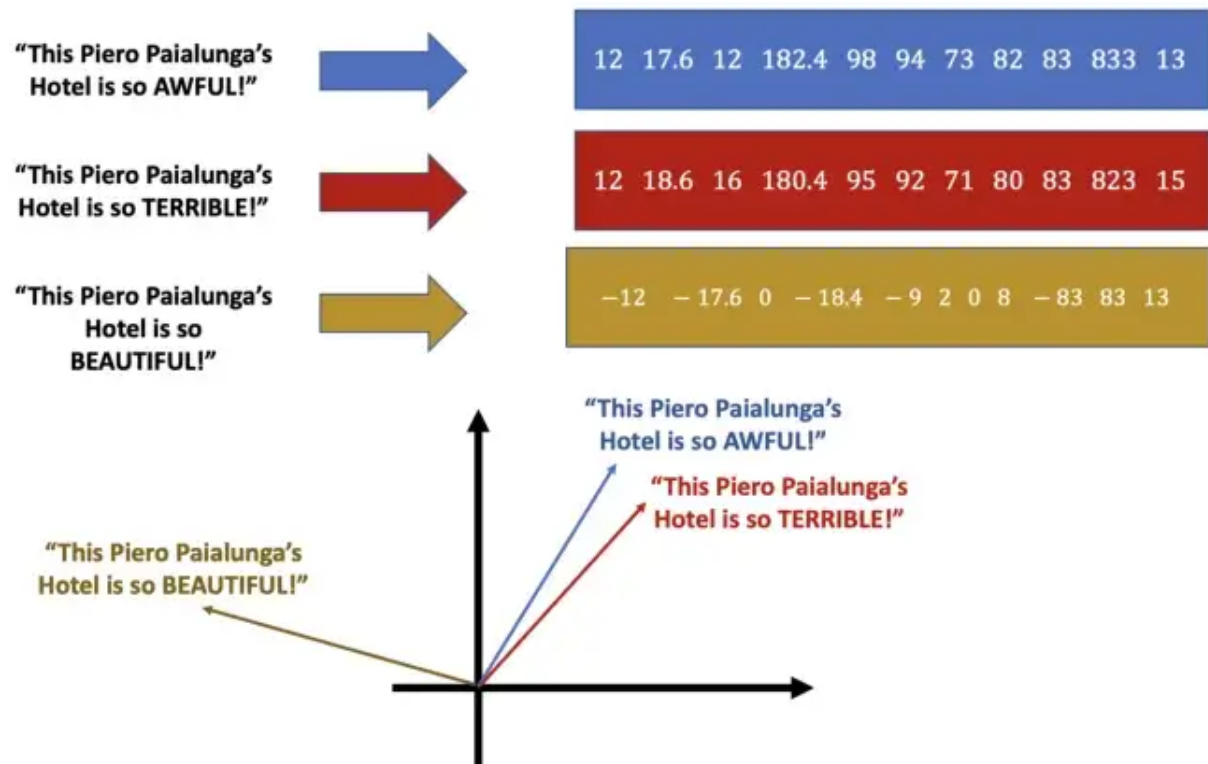


Image by author

As you can see, **similar texts** have **similar vectors** (I know, similarity is a tricky concept, but you know what I mean). and **different texts** have **non-similar vectors**.

There are tons of ways to do the vectorization steps. Some ways are more complex than others; some ways are more **efficient than others**; some ways require machine learning, and some ways don't.

For the purpose of this project (and because I am not a NLP Machine Learning engineer), we will use a fairly simple one, which is called TfIDF vectorizer, ready to use on SkLearn.

Let's start by importing the libraries:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, plot_confusion_matrix
from sklearn.feature_extraction.text import TfidfVectorizer
```


And by importing the dataset we just generated using ChatGPT, with a little bit of preprocessing here and there...

```
labeled_data = pd.read_csv('generated_reviews.csv').drop(columns=['Unnamed: 0'])
labeled_data.Sentiment = labeled_data.Sentiment.astype(int)
labeled_data = labeled_data.dropna().reset_index()
```

Hosted on  Deepnote

And here we go:

```
labeled_data.head()
```

Hosted on  Deepnote

Fantastic. Now let's do this vectorization thing:

```
dataset = labeled_data
from transformers import AutoTokenizer

#tokenizer = AutoTokenizer.from_pretrained("bert-base-uncased")
#tokenized_data = tokenizer(dataset["Reviews"].values.tolist(), return_tensors="np")
vectorizer = TfidfVectorizer(max_features=2500, min_df=7, max_df=0.8)
tokenized_data = vectorizer.fit_transform(dataset['Reviews']).toarray()

labels = np.array(dataset["Sentiment"]) # Label is already an array of 0 and 1
```

```
/Users/pieropaialunga/miniforge3/envs/tf/lib/python3.8/site-packages/tqdm/auto.py:22:
  from .autonotebook import tqdm as notebook_tqdm
```

Hosted on  Deepnote

As I was telling you earlier, the machine learning model we will be using is known as **Random Forest**. What is a random forest? It is a collection of **decision trees**. And

what is a decision tree?

A decision tree is a Machine Learning algorithm that, given a certain information theory criterion, optimizes a tree-search of all the possible splits of the features of your dataset, until it finds a way to distinguish what is 1 and what is 0 based on that split.*

** I am sorry if it is super confusing, but explaining this in 4 lines is a hard task. This [article](#) takes its time to do so and it does that splendidly. Highly recommended.*

Now let's:

1. Define our random forest:

```
rf = RandomForestClassifier(n_estimators=100)
```

Hosted on  Deepnote

2. Split our dataset into training and testing:

```
X = tokenized_data  
y = labels  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

Hosted on  Deepnote

3. Train our model:


```
rf.fit(X_train, y_train)
```

Hosted on  Deepnote

The results are pretty **impressive**, especially considering the lack of hyperparameter tuning.

```
plot_confusion_matrix(rf,X_test,y_test)
```

```
/Users/pieropaialunga/miniforge3/envs/tf/lib/python3.8/site-packages/sklearn/utils/depr  
warnings.warn(msg, category=FutureWarning)
```



Hosted on  Deepnote

3. Sentiment Analysis

As we have our trained model, you can use that on a new, unlabeled dataset. I used a set of New York City hotel reviews I found online, but you can use your own **or** you could make up a review and see how it works.

This dataset is open source ([CC0: Public Domain](#)), very small (2MB) and can be downloaded on [Kaggle](#).

```
target_data = pd.read_csv('NYC_2021_airbnb_reviews_data1.csv')
target_data.drop('url', axis=1)
```

	listing_id	review_posted_date	
0	2595	November 2019	Great location, convenient to ever
1	2595	May 2019	Place was so cute and comfy! Host
2	2595	May 2019	10 / 10 would stay again
3	2595	January 2019	The apartment met expectations to
4	2595	December 2018	Great space in a fun old building
...
17439	1918693	February 2022	Lovely Brownstone in Brooklyn. Cle
17440	1918693	January 2022	We had a great stay at Lorelei & A
17441	1918693	December 2021	This was a perfect spot for mine a
17442	1918693	November 2021	A lovely spot in a lovely neighbor
17443	1918693	November 2021	Overall great stay. Lorelei and Al

17444 rows × 3 columns

Hosted on  Deepnote

Let's preprocess the review column (or your text):

```
dataset = target_data
vectorizer = TfidfVectorizer (max_features=2500, min_df=7, max_df=0.8)
vectorizer.fit(dataset['Reviews'])
new_data_processed = vectorizer.transform(target_data['review']).toarray()
y_pred = rf.predict(new_data_processed)
```

Hosted on  Deepnote

And let's print our predictions:

```
J = np.random.choice(range(0, len(new_data_processed)), 5)
for j in J:
    print('Review number %i: \n'%(j))
    print(target_data['review'].loc[j])
    print('Classified as %i (1=good, 0=bad)' %(y_pred[j]))
```

Review number 13052:

Great stay. Accurate to listing.

Classified as 1 (1=good, 0=bad)

Review number 8373:

The studio was a great place to stay for a couple days while on a work trip. Centrall

Classified as 1 (1=good, 0=bad)

Review number 9744:

Thank you

Classified as 0 (1=good, 0=bad)

Review number 10006:

Zora & Chris are the best host you could expect for. They are always available and ar

The place is even better in real life than the photos : clean, cosy and warm, it feel

If you need a place to stay in NYC, I definitely recommend it!

Classified as 1 (1=good, 0=bad)

Review number 7507:

Incredible area, a little funky that you have to walk through a common area to get to

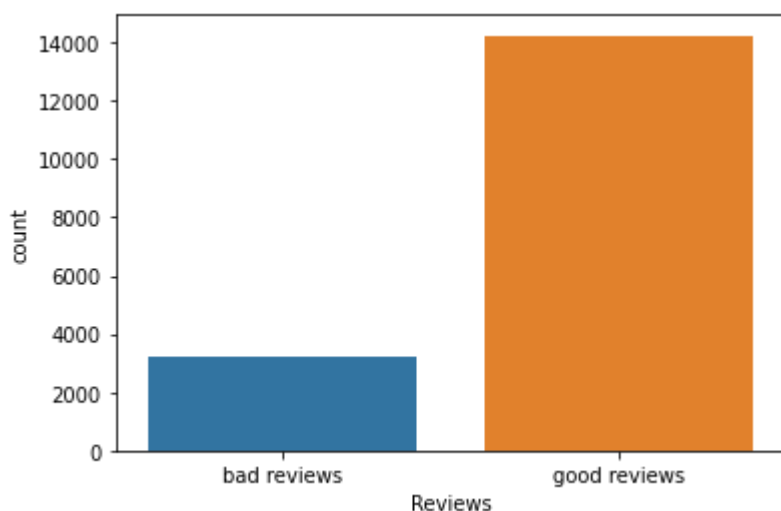
Classified as 1 (1=good, 0=bad)

Hosted on  Deepnote

As we can see, all these 5 random reviews that are classified as 1 are actually good!

Let's show a count plot:

```
sns.countplot(x=pd.DataFrame(y_pred)[0])  
plt.xticks([0,1],['bad reviews','good reviews'])  
plt.xlabel('Reviews')
```



Hosted on  Deepnote

4. Considerations

So what did we do here?

1. We acknowledge that **ChatGPT is awesome**.
2. We used ChatGPT to **build a dataset** for our surrogate model. More specifically, we used ChatGPT to make up good and bad reviews for hotels.
3. We used that **labeled dataset** that we built to train our Machine Learning model. The model that we used is a **Random Forest Classifier**.
4. We tested our trained model on a new dataset getting **promising results**.

Is there any room for improvement? A ton.

Open in app 

[Sign up](#)

[Sign In](#)



2. We can **improve our querying skills** by giving different inputs, maybe also in other languages rather than english only

3. We can **improve our Machine Learning model** by doing some hyperparameter tuning

Now, let me conclude this with a thought.

There are, and there will be, a lot of concerns about how and who is going to use Open AI ChatGPT. While I am not a lawyer (let alone an expert in ethical AI), I can imagine how this tool can be dangerous in many ways and on many different levels.

I strongly disagree with the people who are not impressed by the **performance** of ChatGPT, because I find it pretty amazing and I am so excited to see how this technology is going to evolve. I hope this toy example sparks something in my readers too. ❤️

5. Conclusions

If you liked the article and you want to know more about machine learning, or you just want to ask me something, you can:

A. Follow me on [Linkedin](#), where I publish all my stories

B. Subscribe to my [newsletter](#). It will keep you updated about new stories and give you the chance to text me to receive all the corrections or doubts you may have.

C. Become a [referred member](#), so you won't have any "maximum number of stories for the month" and you can read whatever I (and thousands of other Machine Learning and Data Science top writers) write about the newest technology available.

[Chat](#)

[Open AI](#)

[Machine Learning](#)

[Artificial Intelligence](#)

[NLP](#)



201



4

Sign up for The Variable

By Towards Data Science

Every Thursday, the Variable delivers the very best of Towards Data Science: from hands-on tutorials and cutting-edge research to original features you don't want to miss. [Take a look.](#)