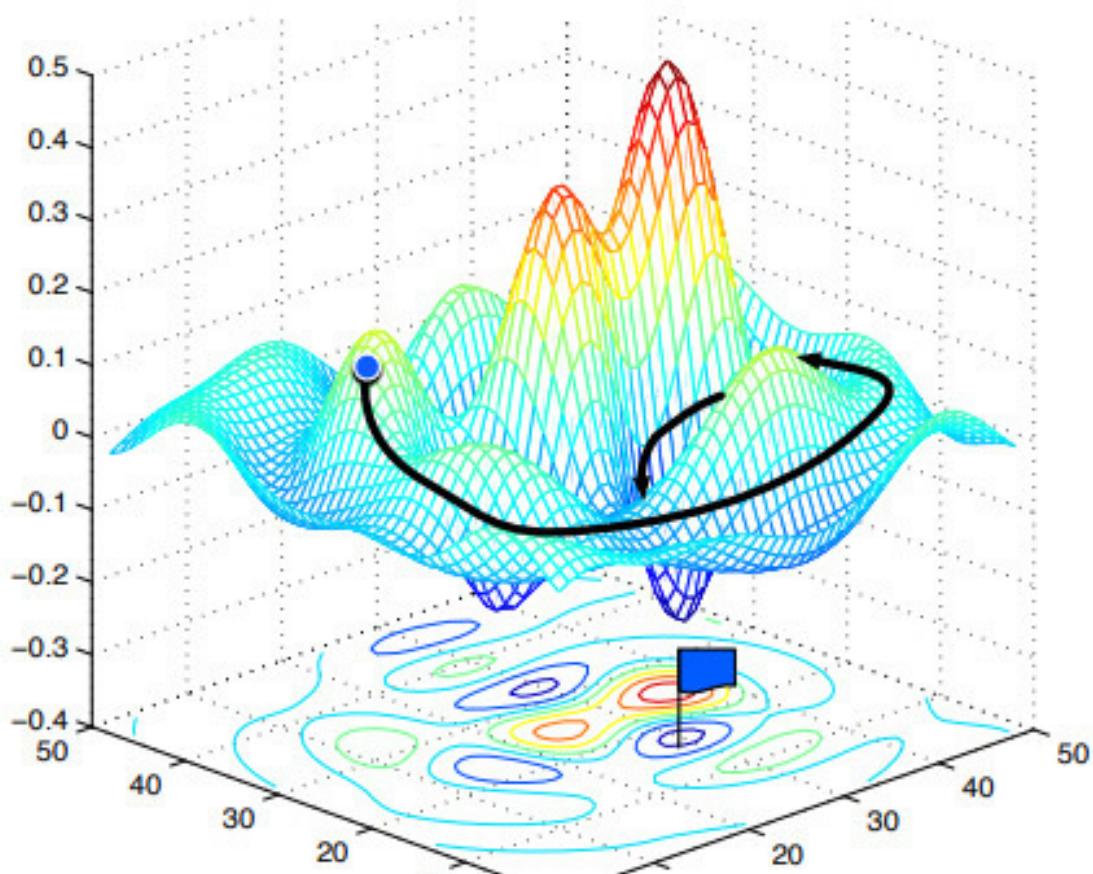


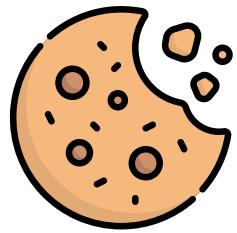
Gradient descent 101 - cookies and local minima



**Ever gotten lost in a mall trying to
find the food court?**



You enter the mall with a clear objective but the irresistible smell of cookies leads you in a different direction.

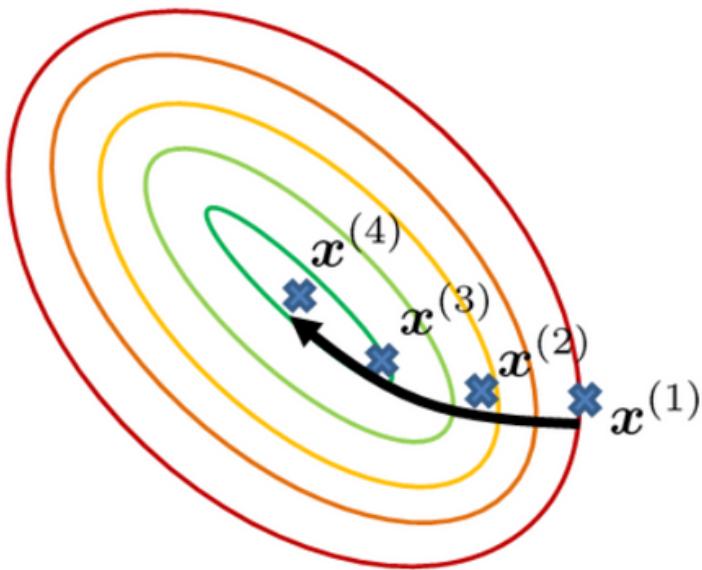


You follow the sweet smell, and it brings you to a bakery on the ground floor. You haven't found the food court. Turns out it's on the top floor!

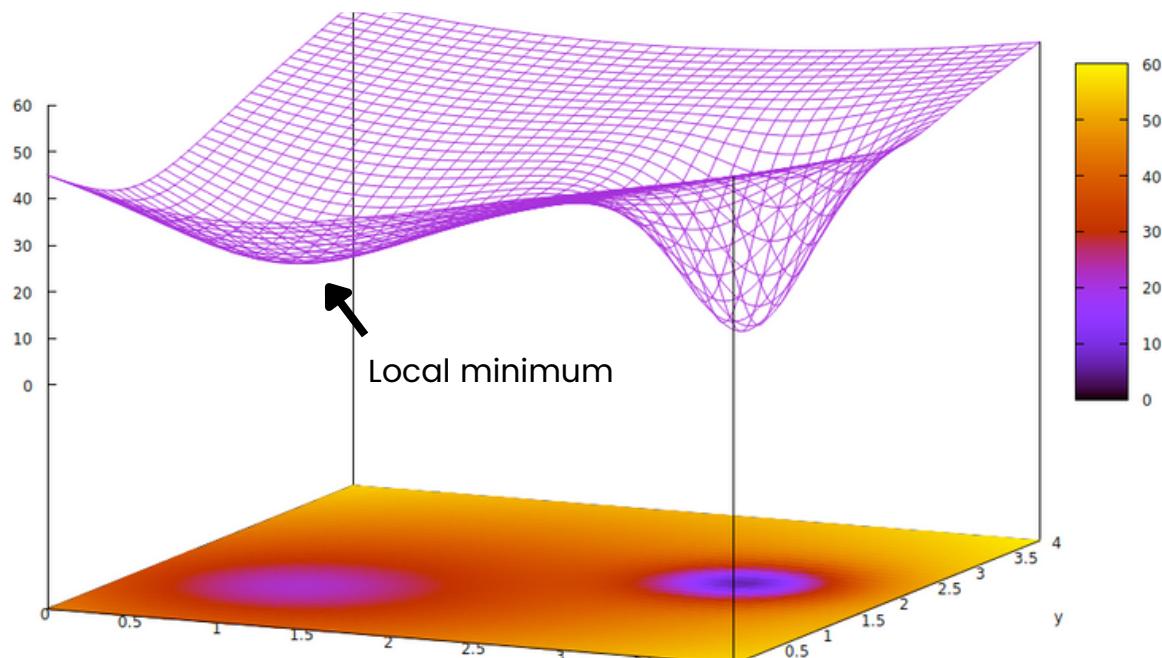


This scenario represents the problem of local minima in gradient descent, a commonly used optimization algorithm in machine learning.

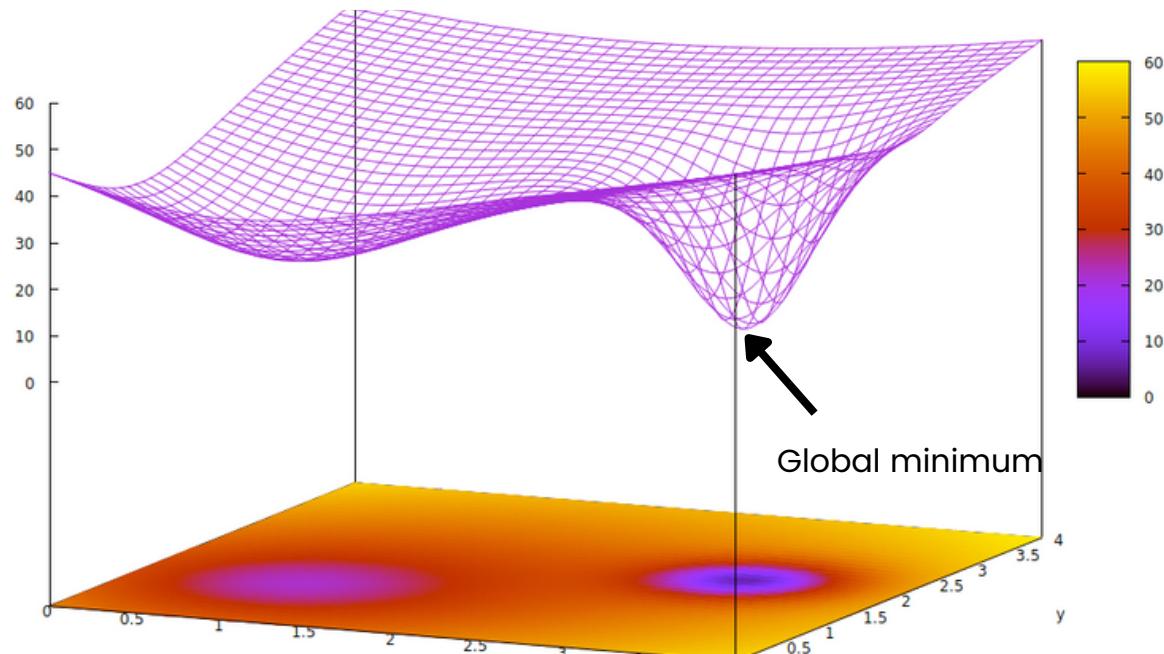
The smell of cookies represents the "pull", or the slope. Naively following the smell represents the greedy approach adopted by gradient descent.



The bakery you ended up in is a local minimum, which means it's the optimal point in your immediate surroundings but not the optimal point overall.

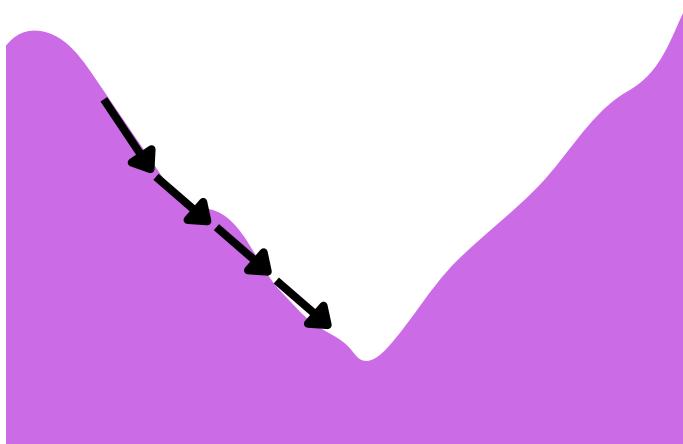


**The food court on the top floor is
the global minimum - the ultimate
objective.**



Applying gradient descent in machine learning involves updating the parameters of a model in the direction of steepest descent of the loss function.

**Just like you unwaveringly followed
the smell of the cookies, gradient
descent follows the slope of the
loss function to reach the optimal
solution.**



However, this approach can often lead to a local minimum (especially in non-convex functions) like the bakery, which is not what we were looking for at all.

To overcome the problem of local minima, techniques such as random restarts and simulated annealing are used to explore different parts of the loss function and eventually reach the global minimum.

**So the next time you either get lost
in a mall or are faced with an
optimization problem, remember the
importance of exploring different
paths and not getting stuck in a
local minimum.**

**Keep your eyes on the global
minimum and don't let the smell of
cookies distract you!**



**What's your favourite optimization
algorithm and why? Let me know** 

