

Two-Phase Commit is Evil

Rob Daigneau

Rob.Daigneau@ArcSage.com



Overview

- Review of Transaction Concepts
- Distributed Transactions
- 2PC Considerations
- The CAP Theorem
- Alternatives to 2PC



About Me

- Principal of **ArcSage**
- Author
 - "**Design Patterns for the Service Layer**
Fundamental Design Solutions for SOA, REST, EDA, and the Cloud"
 - *Estimated release date == 2009, Addison Wesley*
- Host of **www.DesignPatternsFor.Net**
- Chief Architect  
- Director of Architecture 

Review of Transaction Concepts



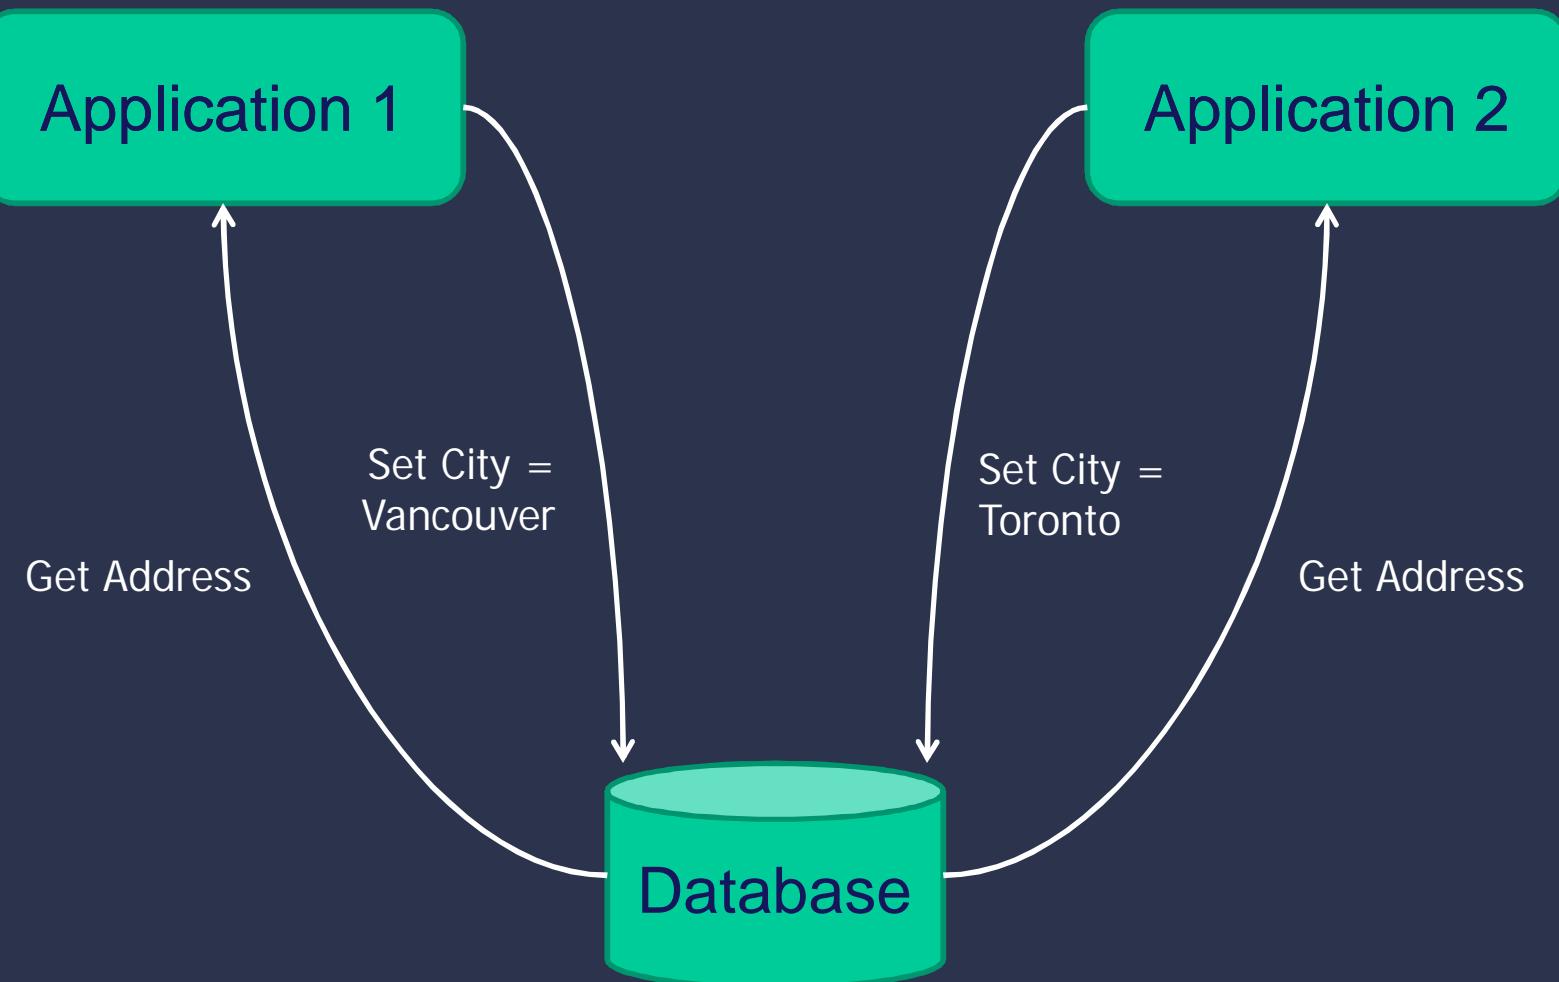
Why Do We Need Transactions Anyway?

- It's all about **Data Integrity**
... while attempting to provide **Concurrent Data Access**
for Read or Write operations
- Classic concurrency problems ...
 - The Lost Update problem
 - Unrepeatable Read



The Lost Update Problem

Toronto? What happened?



The Lost Update Problem

- Last one in wins
 - No warning to loser



Unrepeatable Reads

Application 1

Eastern Region Sales
= \$1 million

Western Region Sales
= \$2 million

Eastern Region Sales
= \$1.5 million

Get Total Sales for
Eastern Region

Get Total Sales for
Western Region

Application 2

Eastern Region
Sales += \$500k

Western Region
Sales += \$250k

Total Sales = \$3.75 million



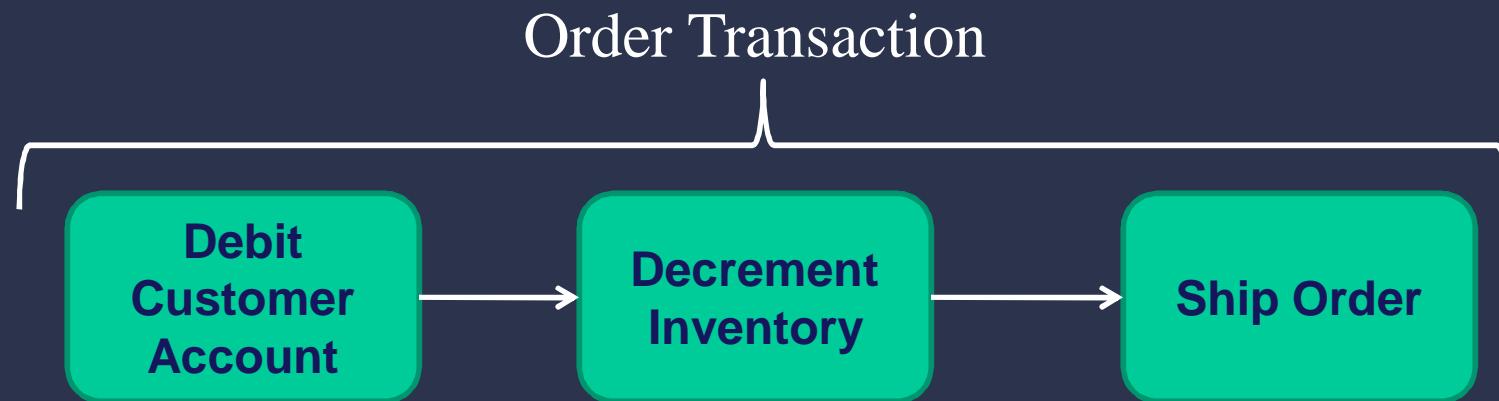
Unrepeatable Read

- “Reader applications” are unable to get the same results with each read of the same data
 - Occurs because “Writer applications” are operating on the same data
 - i.e. Read and Write operations are interleaved



What is a Transaction?

- A logical unit of work
 - May contain one or many steps



ACID

- **Atomicity**
 - “It all goes or none of it goes”
- **Consistency**
 - Data must be valid before and after the operation
- **Isolation**
 - Ensure that intermediate states of data aren’t visible until committed
- **Durability**
 - The result of the data operation is successfully persisted to the resource and can survive crashes or exceptions

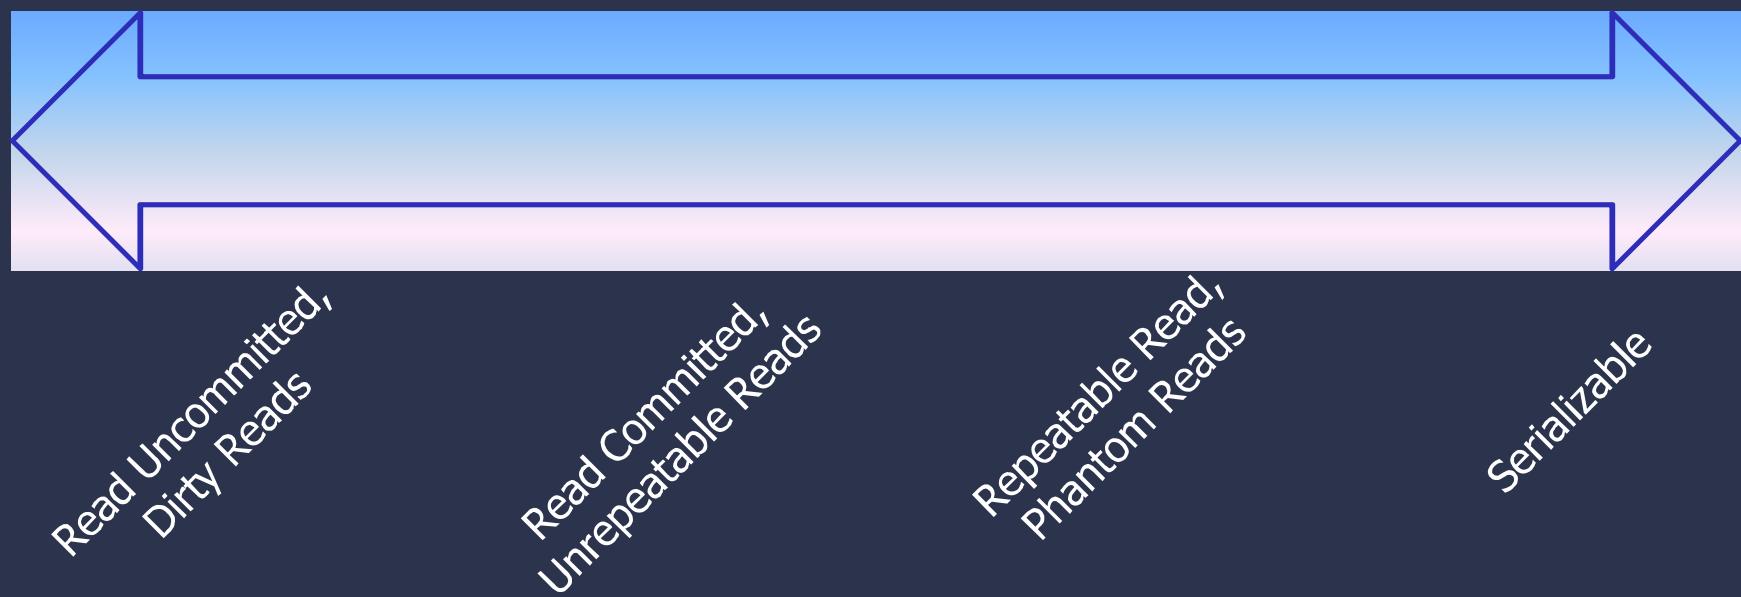


Transaction Isolation Levels

- Typically set at the database or by the ORM

More Concurrency,
Less Data Integrity

Less Concurrency,
More Data Integrity



Patterns for Concurrency

- **Patterns of Enterprise Application Architecture** (*David Rice*)
 - Meant to prevent the Lost Update Problem
 - Controlled by custom code, Datasets, ORMs, etc.
- **Optimistic Offline Lock**
 - Guid or Timestamp is checked at commit
 - Downside = potential for lost work
- **Pessimistic Offline Locks**
 - Preference is given to ensuring no lost work
 - Takes early “Logical Exclusive Locks”
 - Bad for high concurrency, scalability
 - What if you don’t release your locks?





Distributed Transactions

Copyright © 2009, Rob Daigneau, All rights reserved



Distributed Transactions

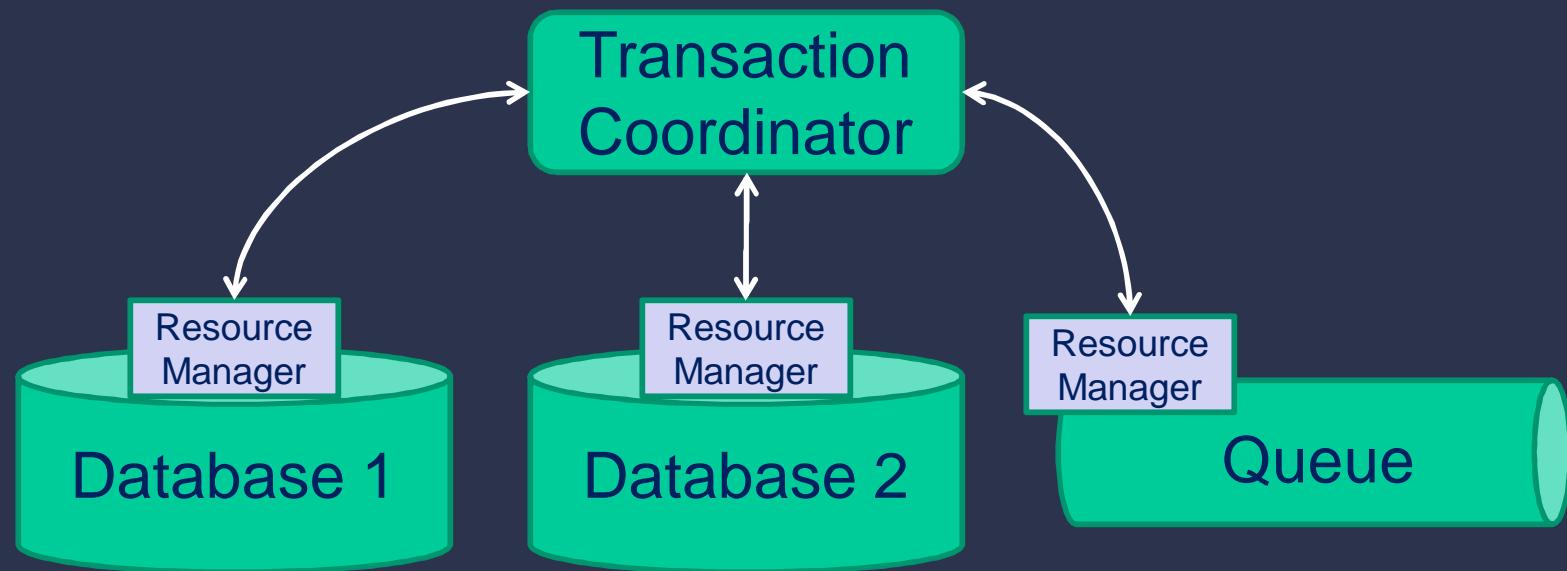
- Transaction that spans multiple resources
 - Databases, queues, files
- Key Concept = Two-Phase Commit (2PC)

- SIDEBAR: WS-Atomic Transactions
 - Defines how distributed transactions are propagated across services
 - WS* Services (i.e. SOAP/WSDL)
 - Keep this in mind if you use WCF !!!!



2PC Roles

- Transaction Coordinator (TC)
- Resource Managers (RM), Transaction Participants

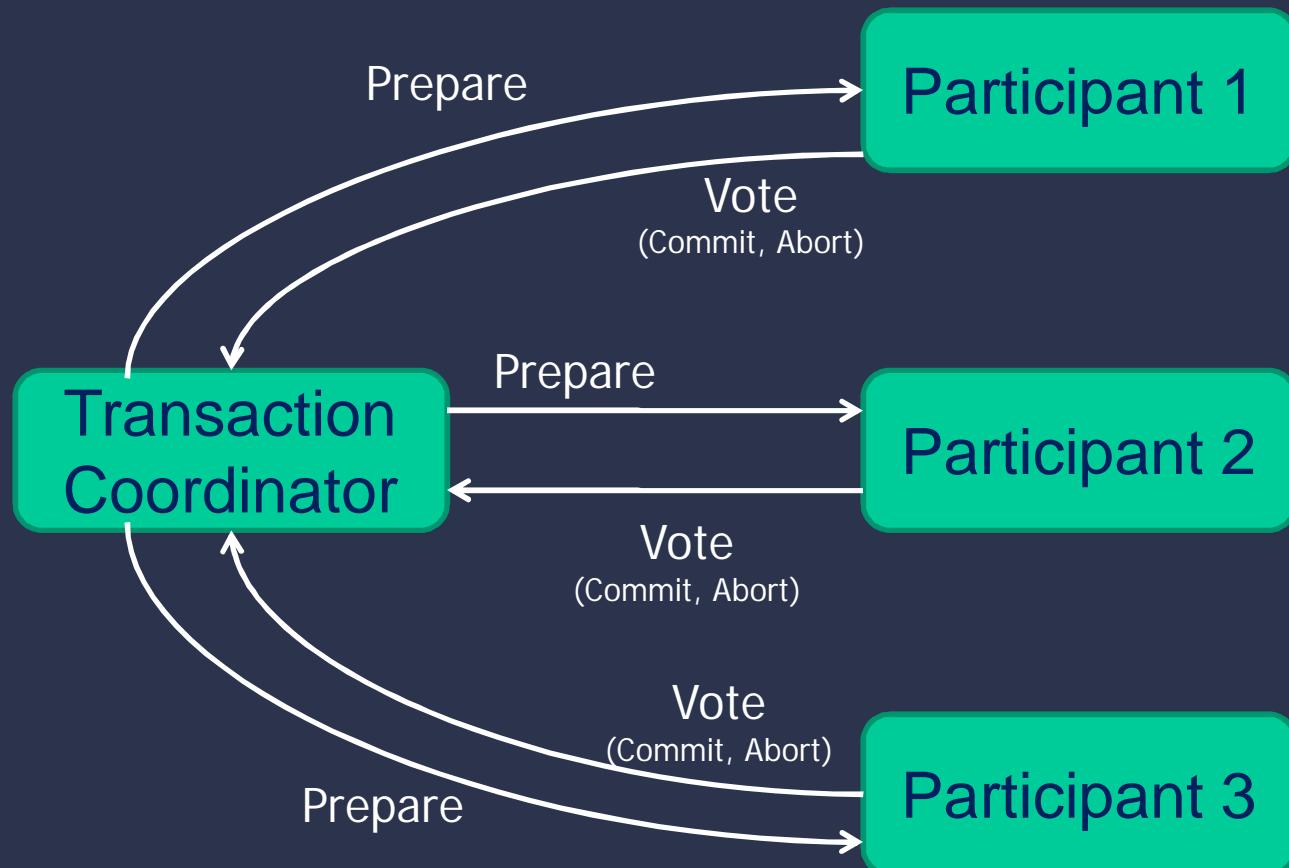




Copyright © 2009, Rob Daigneau, All rights reserved

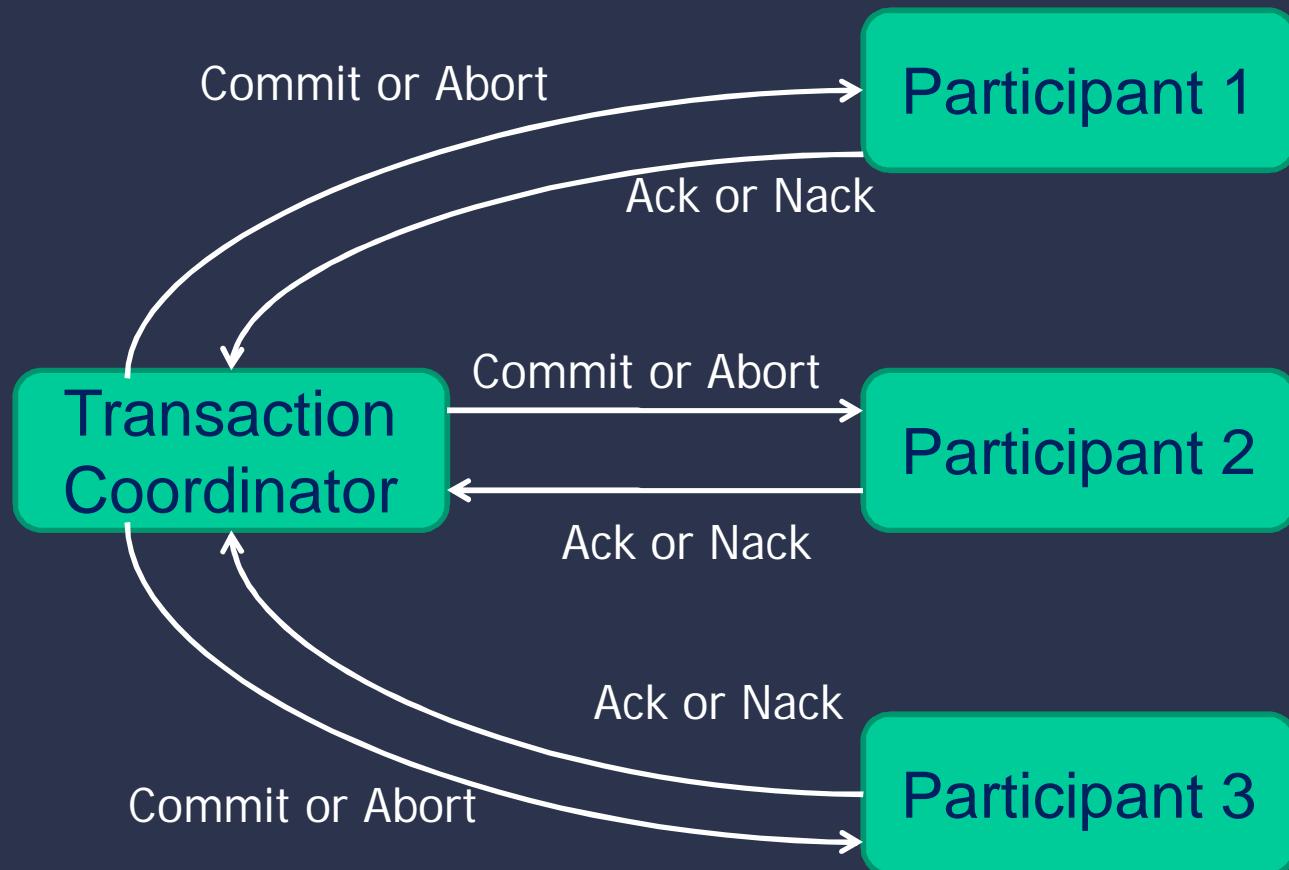
2PC Phase 1

- TC acquires votes from each participant



2PC Phase 2

- TC evaluates votes and tells participants to commit or abort

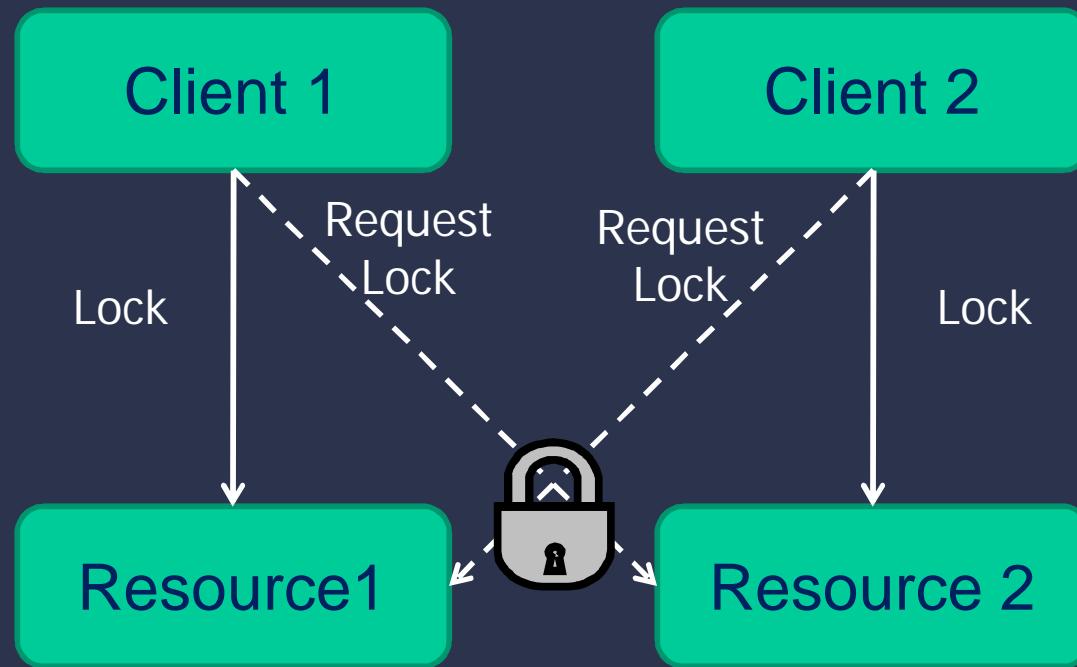


The Key Problem with 2PC

- Resource locks may be held for excessive lengths of time
 - Each participant has to wait for ...
 - Other participants to complete
 - TC to tell each participant to commit or abort
- Result =
 - Client requests can be blocked, time-out
 - Also depends on Isolation level you set
 - Worst-case scenario = Deadlocks

The Deadly Embrace ... Deadlocks

- The odds of deadlocks are increased with 2PC



Other Problems with 2PC

- These problems may occur during any phase ...
- If TC fails, RMs may block indefinitely
 - In many cases you'll have to kill the transaction
- If TC directs participants to commit, and a participant is down ...
 - TC and other participants may hang
- An extremely chatty protocol
 - Lots of network traffic, really poor performance for web services
- When should the TC decide that the transaction should time out?
- Increases coupling, decreases service autonomy





2PC Considerations



Copyright © 2009, Rob Daigneau, All rights reserved

Are You Actually Dealing with Distributed Resources?

- Are you writing to multiple resources?
 - i.e. Physical databases, queues, files, etc.
 - e.g. Have you created a WCF services that uses transactions, yet writes to the same resource?
 - WCF will automatically use 2PC (rather than local transactions) when transactions are flowed, even if against the same resource



WCF Example

```
<bindings>
  <wsHttpBinding>
    <binding name =
      "Dangerous"
      transactionFlow=true />
  </wsHttpBinding>
</bindings>
```

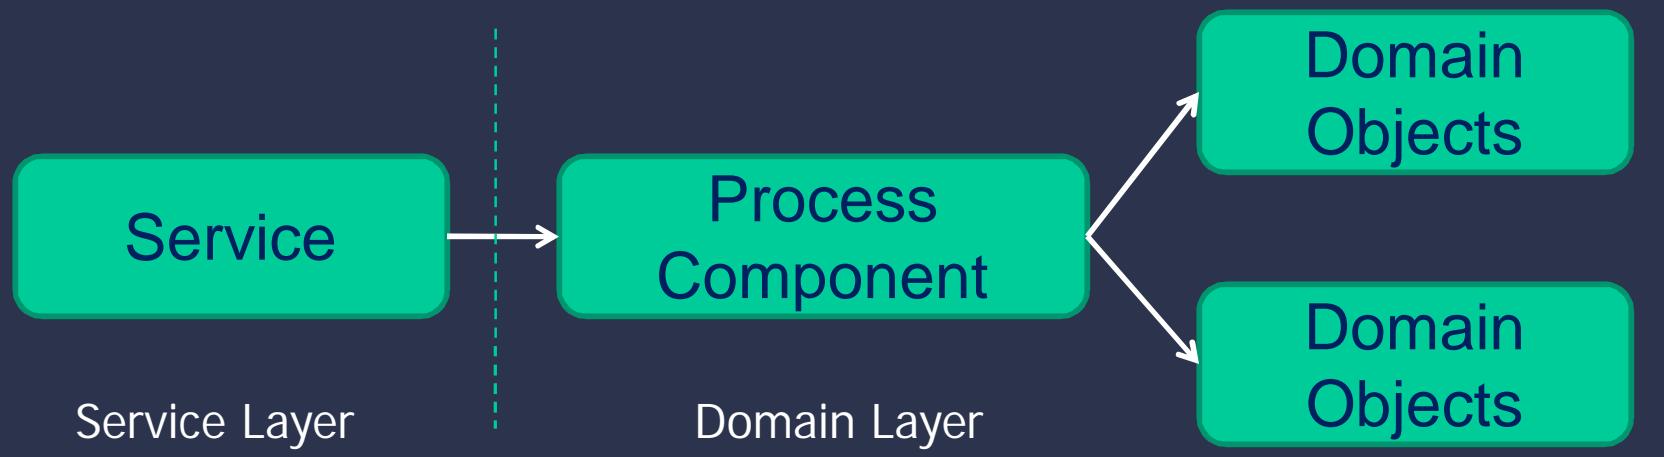
```
[ServiceContract]
interface IFooBar
{
  [OperationContract]
  [TransactionFlow(
    TransactionFlowOption.Allowed)]
  SomeMessage FooBar(SomeMessage request);
}

public class FooBarClass: IFooBar
{
  [OperationBehavior(
    TransactionScopeRequired=true,
    TransactionAutoComplete=true)]
  public SomeMessage FooBar(SomeMessage request)
  {
    // implementation here;
  }
}
```



What to Do?

- Consider Refactoring
 - Rather than calling multiple services ...
 - Refactor and consolidate business logic
 - Reposition common business logic in the Domain Layer rather than the Service Layer
 - Let the Domain Layer components manage the transaction



You are Dealing with Distributed Resources ...

- Do you “own” the resources, or do they exist in some other domain or business?
 - 2PC may be Ok if ...
 - You own the resources AND
 - You make sure that the transactions stay “relatively short” in duration AND
 - Concurrency problems are rare OR
 - You’re Ok with lower Isolation Levels (e.g. Dirty Reads)
 - Example
 - Reading a message from a queue and writing data from the message to a database



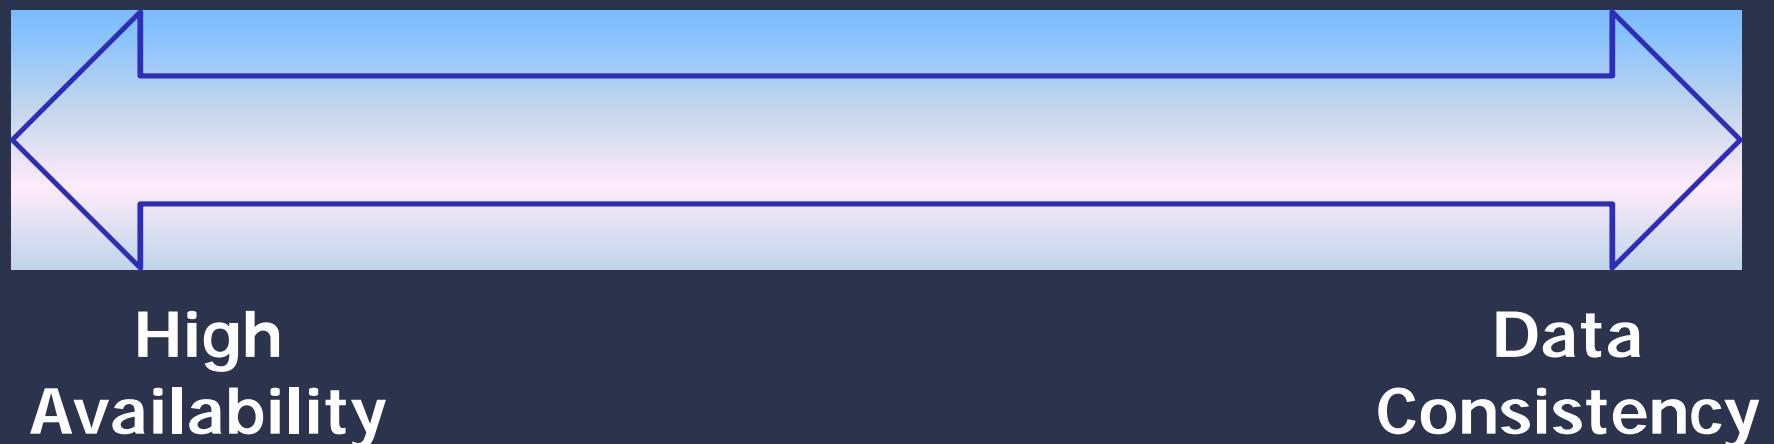


The CAP Theorem



A Fundamental Trade-off

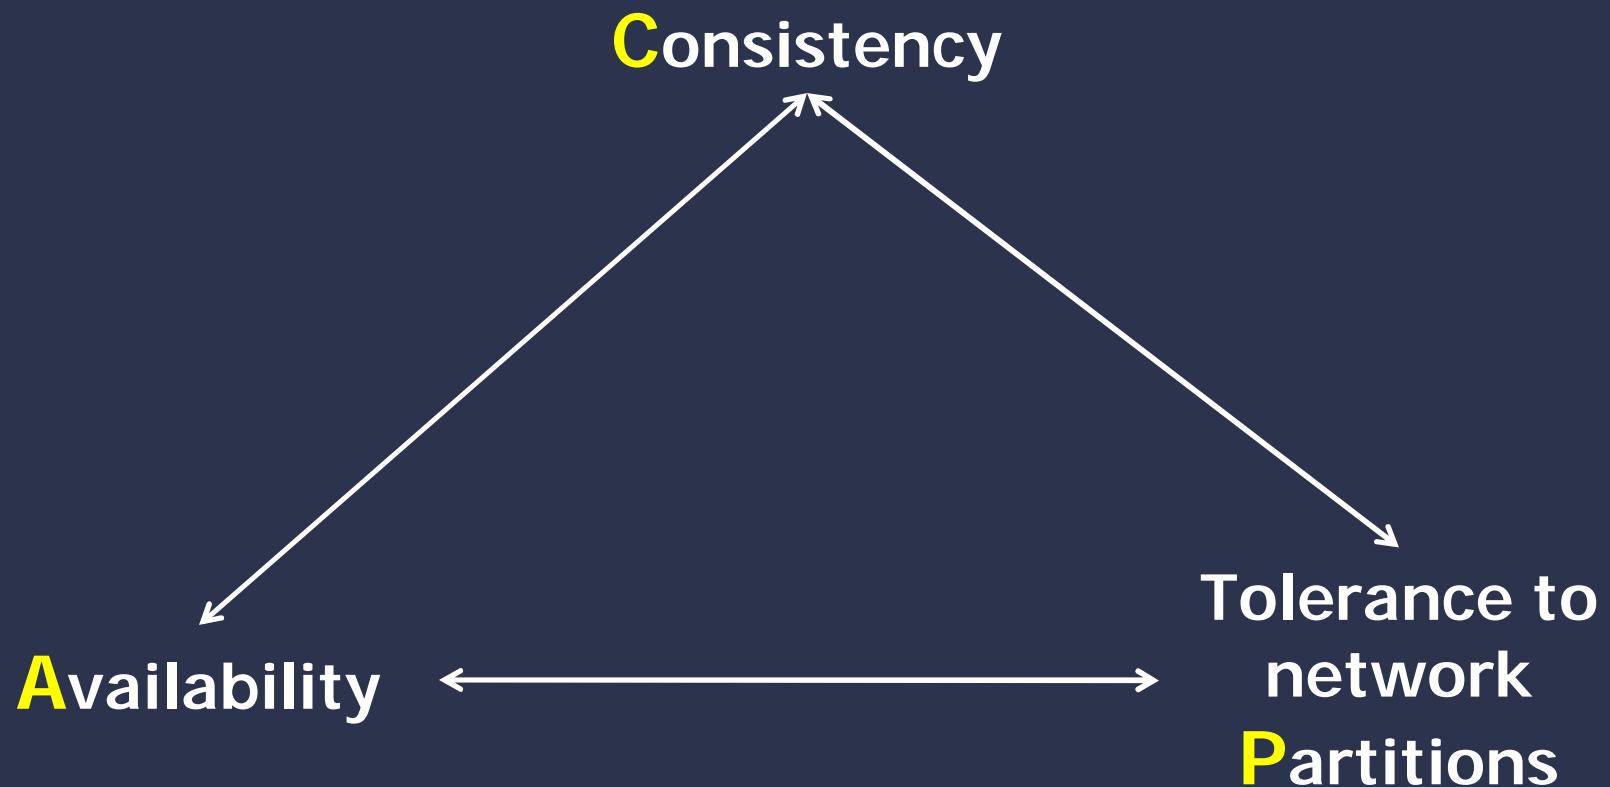
- www.cs.berkeley.edu/~brewer/cs262b-2004/PODC-keynote.pdf
- Dr Brewer contends that a fundamental trade-off exists ...



The CAP Theorem

- Brewer, Dr Eric A; www.cs.berkeley.edu/~brewer/cs262b-2004/PODC-keynote.pdf

Pick two ...



The Basic Trade-off

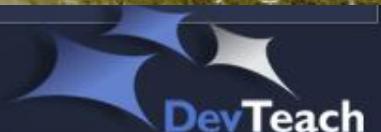
- Highly Available and Scalable distributed systems ...
 - Forgo up-to-the-minute data consistency
 - Stale data can be Ok for a while
 - Assume that data consistency will catch up
- Consider the following ...
 - Travel web sites (e.g. flight reservations)
 - Financial services (e.g. trades)





Alternatives to 2PC

© 2001 ays3@cornell.edu



Copyright © 2009, Rob Daigneau, All rights reserved

What to Do?

- How can distributed systems achieve high availability and scalability?
 - Forget 2PC
 - Do use Optimistic Off-line Locks
 - Set “Expiration Dates” on data and locks
 - Initiate “Data Refresh Requests” on expiration
- Partitioning
 - Provides isolation, independent scaling and tuning for a functional domain
 - Contrasted with large, monolithic systems

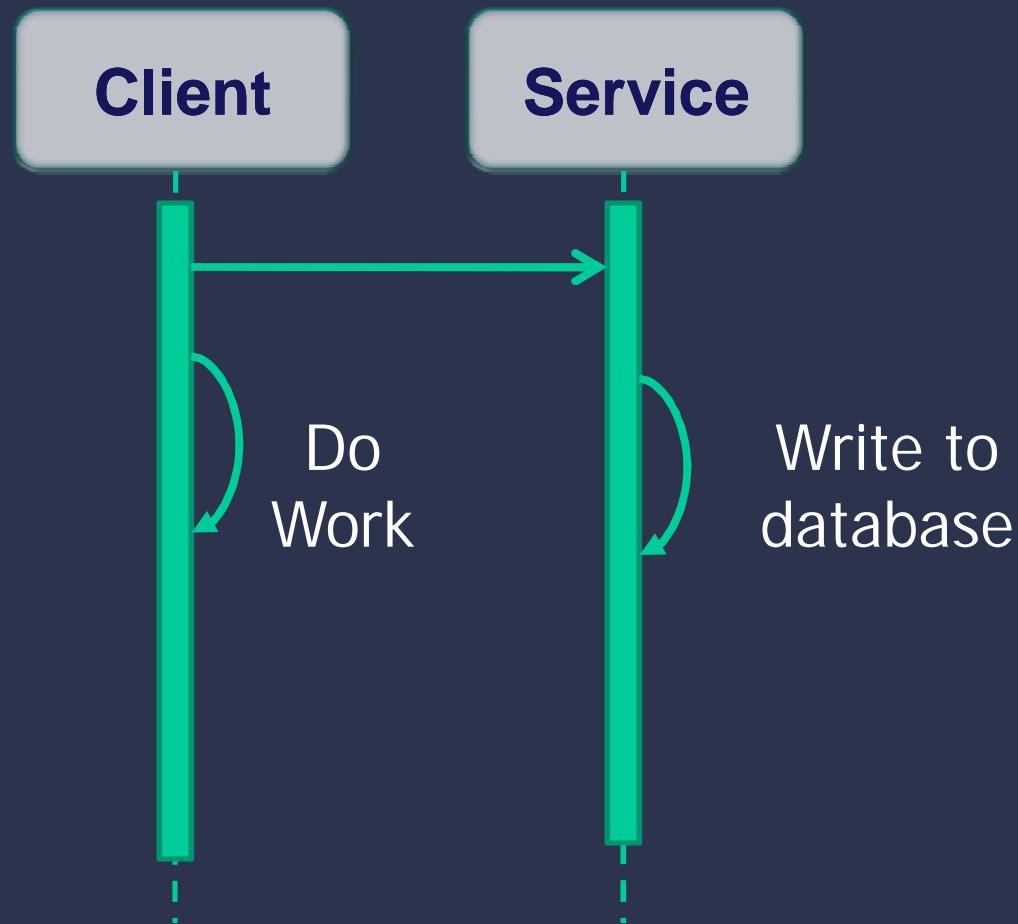


Consider Asynchronous Messaging

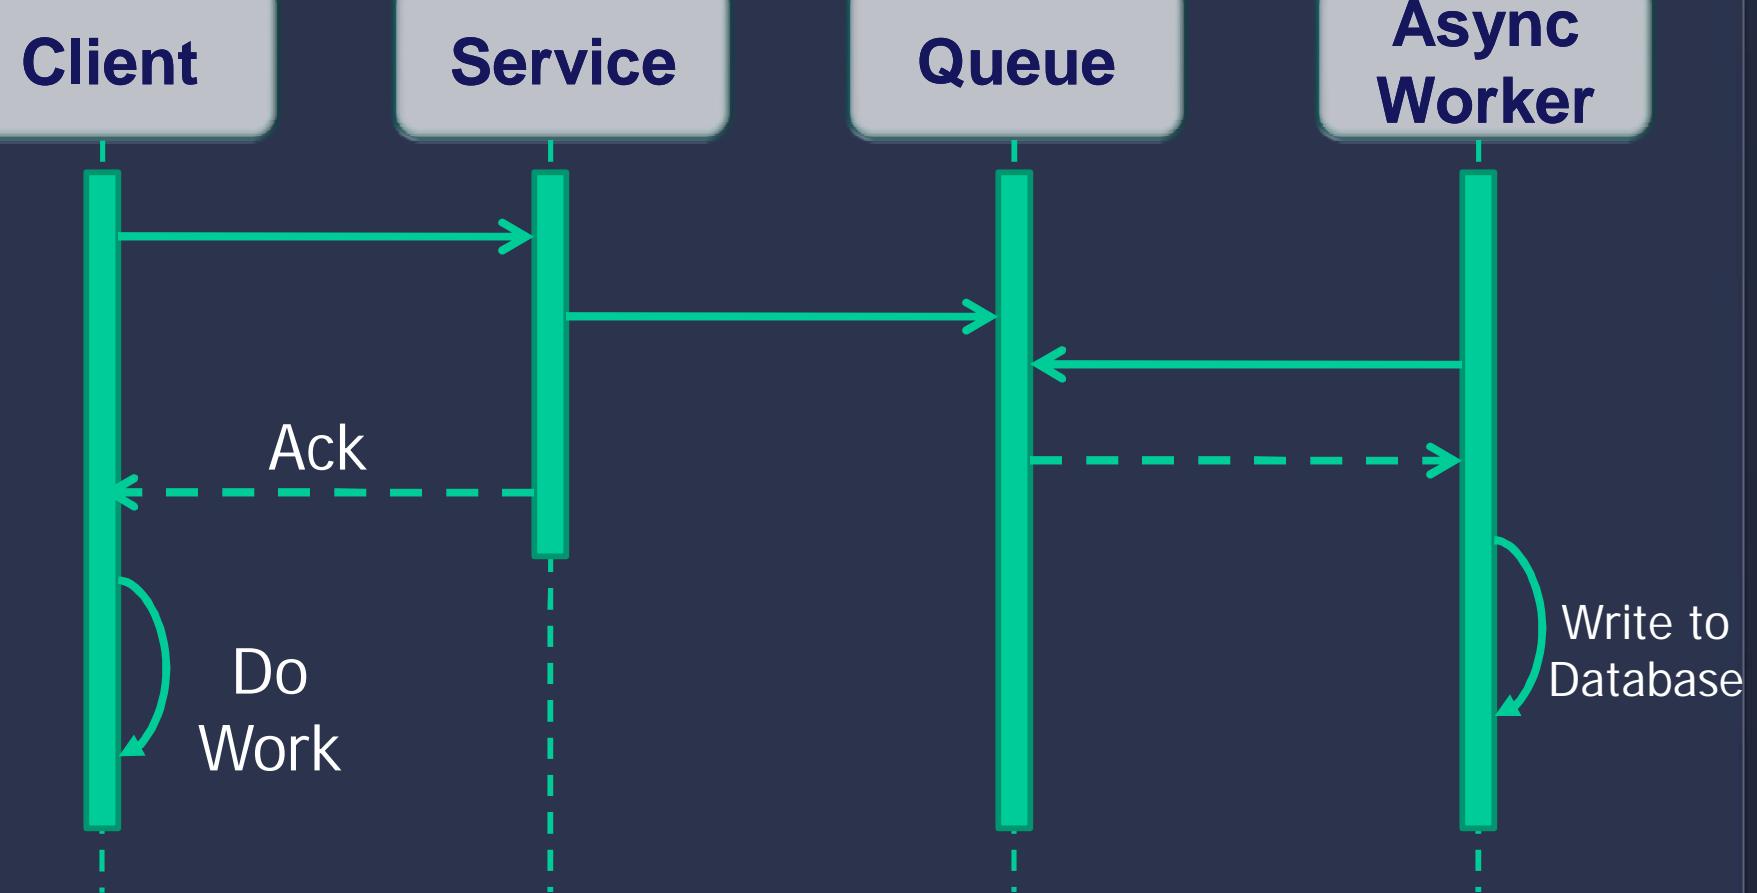
- Key Enablers ...
 - Queues
 - Messages
- Service Design Patterns to Consider ...
 - Event Sink
 - Request/Acknowledge
 - Asynchronous Response-Pull
 - Asynchronous Callback Message
- Compensating Transactions



Event Sink



Request/Acknowledge



Pattern Considerations

- Event Sinks
 - Provide no means for clients to discover errors

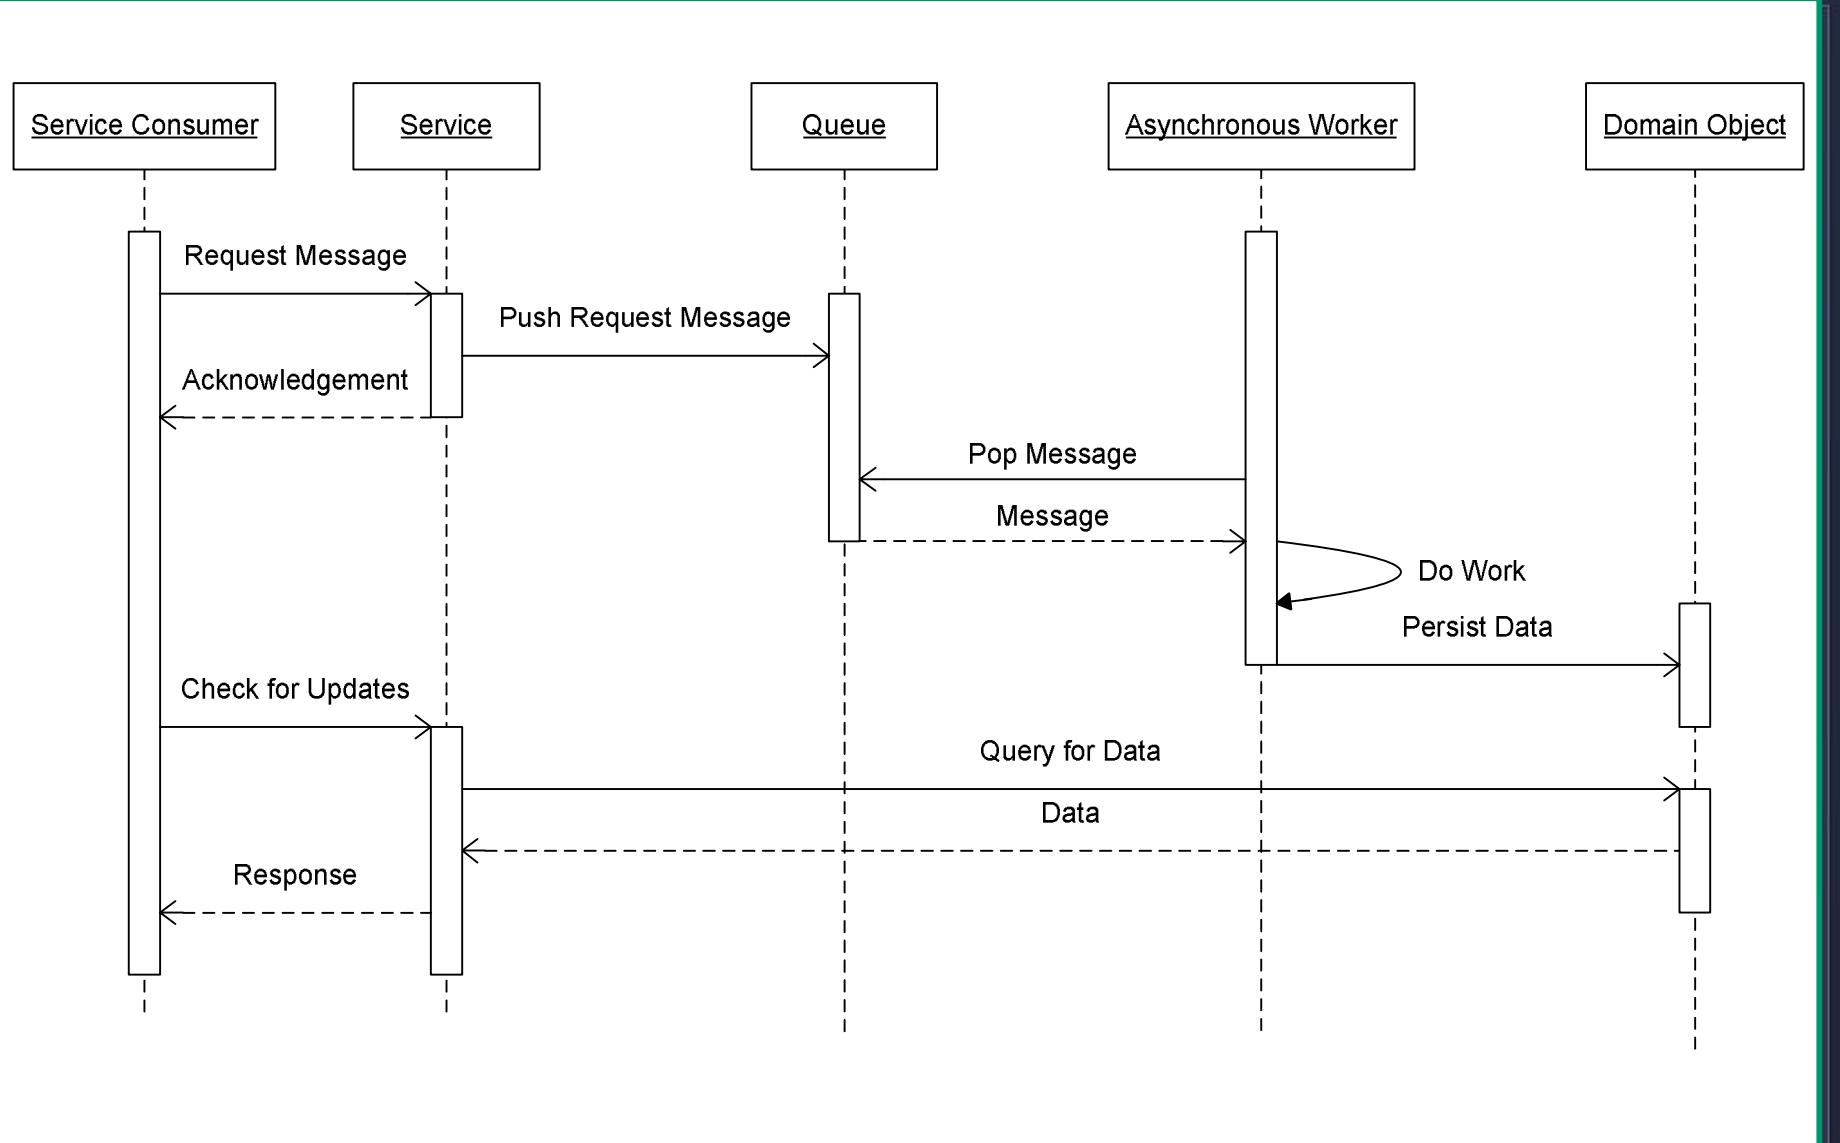
- Request/Acknowledge
 - Nacks (Negative Acknowledgements)

Usually only tell you about problems with Authentication or Data Validation

- Questions
 - How does client learn about errors?
 - How can client undo operations in remote systems?

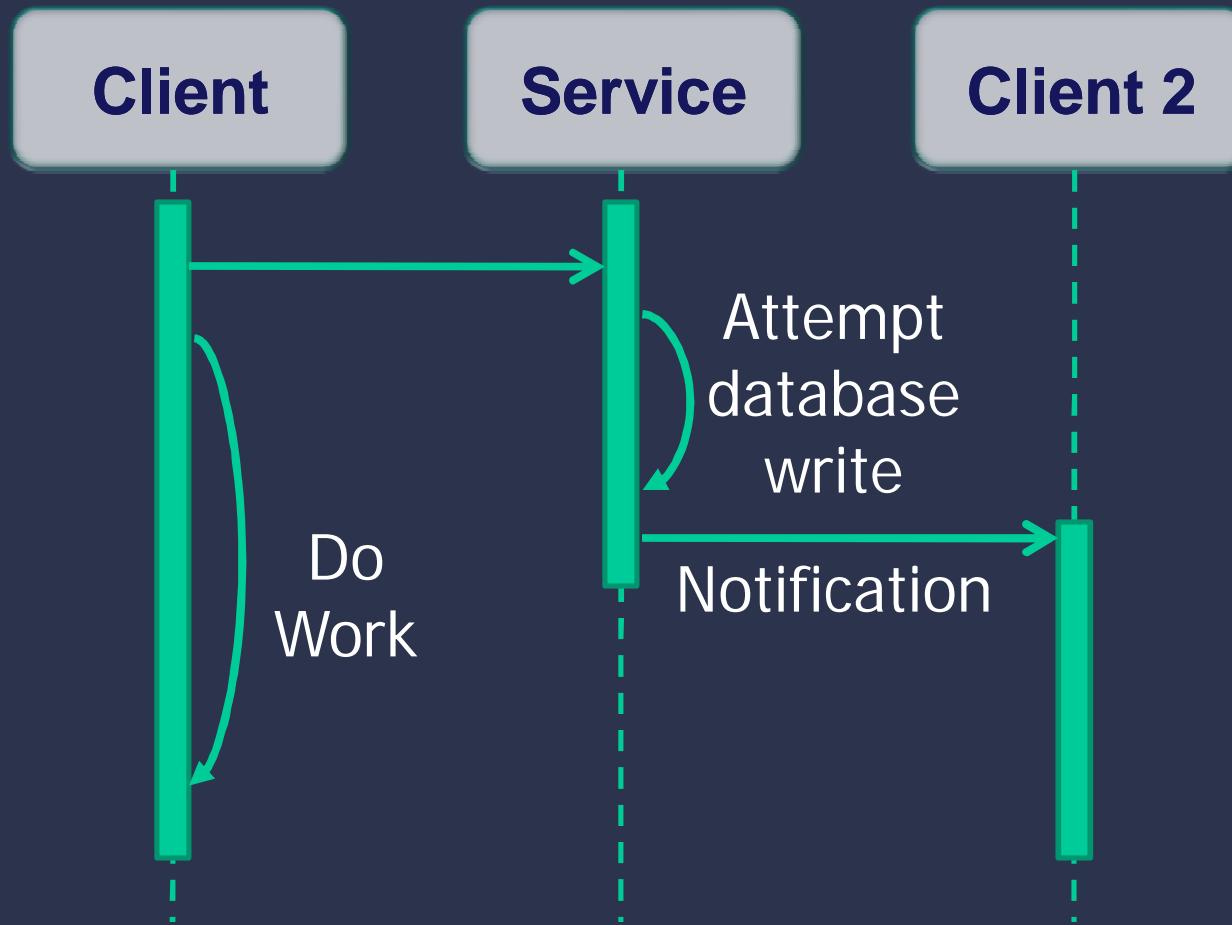


Asynchronous Response-Pull



Asynchronous Callback Message

- *a.k.a. Message Relay*



Compensating Transactions

- A command or message that is the logical inverse of some previous command or message

e.g.

- Reserve a seat on a Flight, Cancel Flight Reservation
 - Debit Account, Credit Account
-
- This is the way “Long-Running Transactions” work in the real world

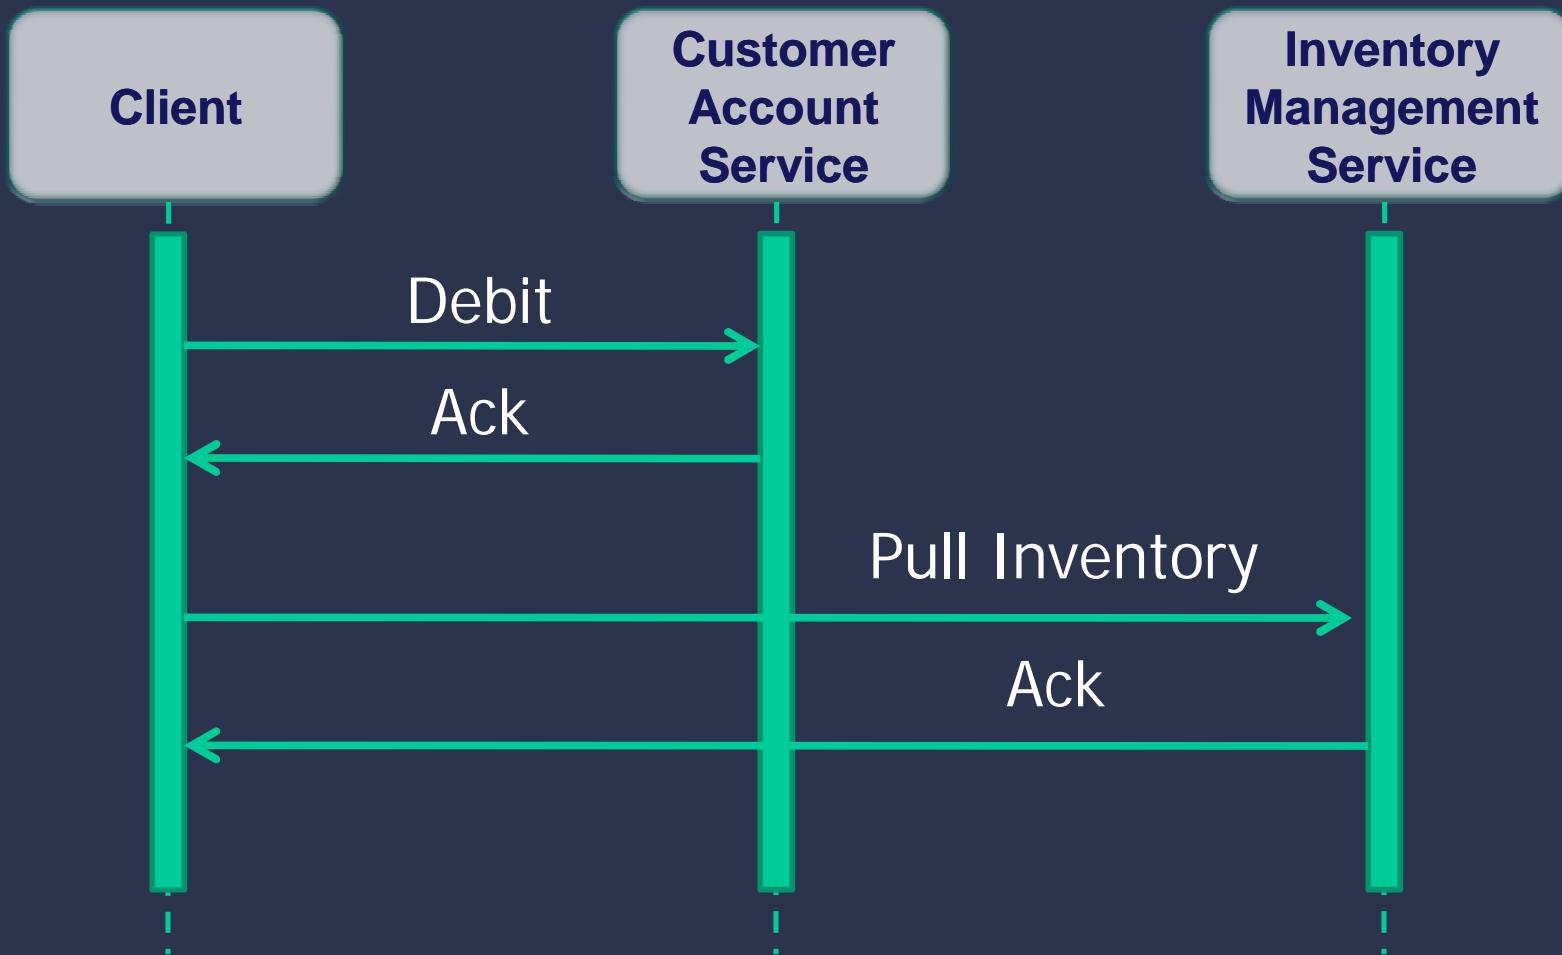
SIDE BAR:

Long-Running business processes aren't really Transactions from the technical perspective

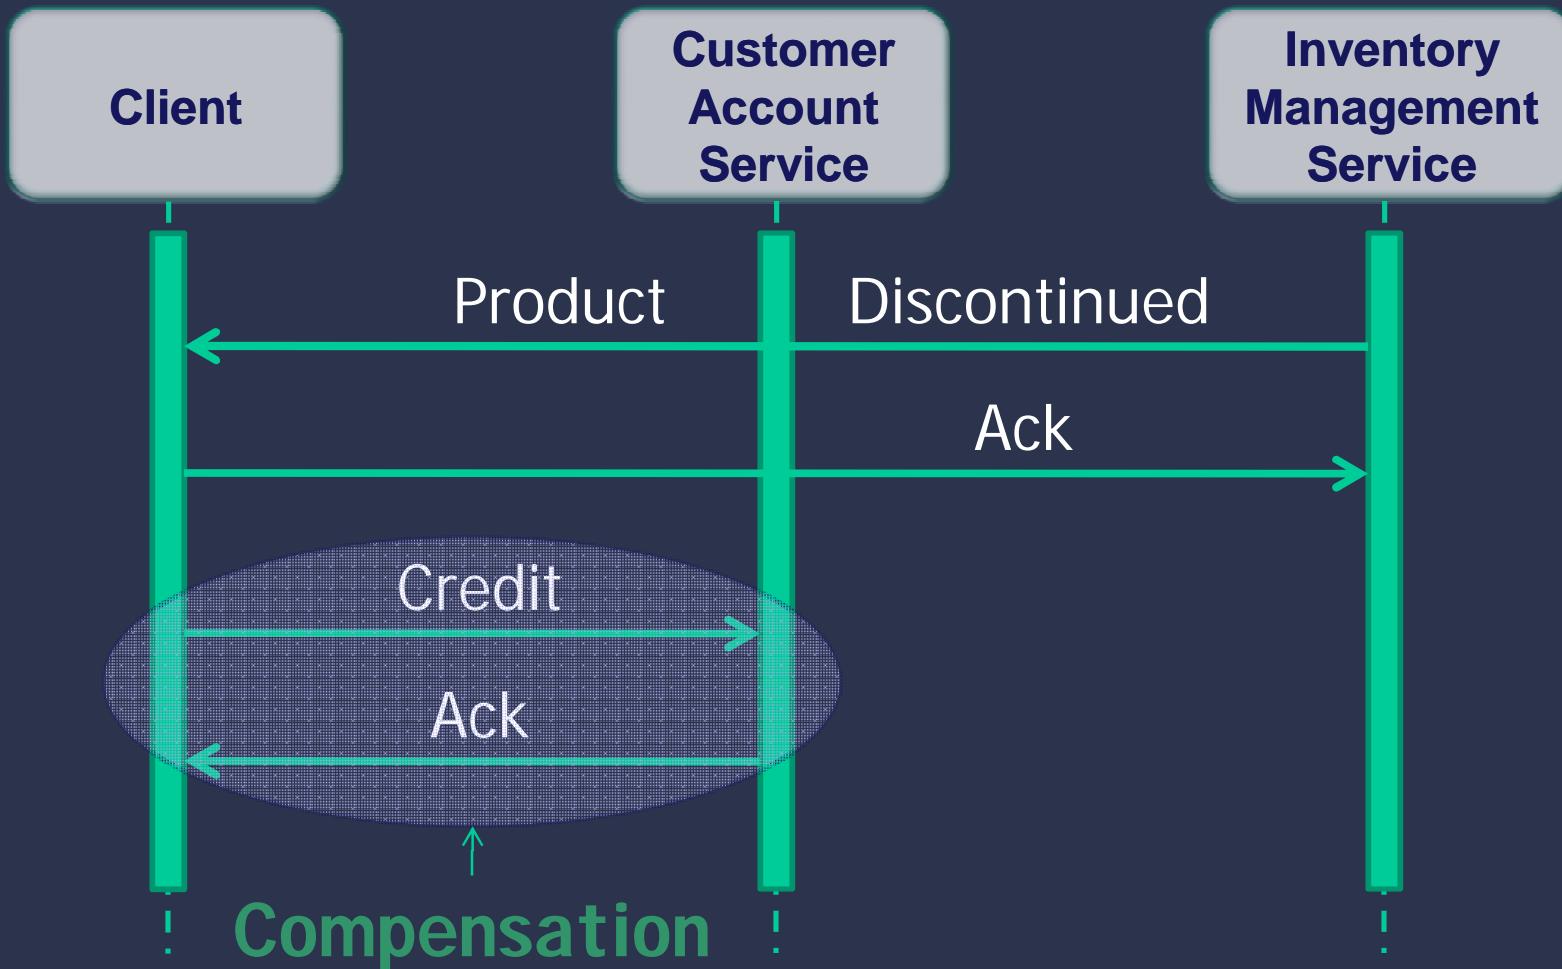
- They are neither Atomic or Isolated.



Compensating Transaction Example, Part 1



Compensating Transaction Example, Part 2



Note: The client is usually not the same instance that sent out the original messages

Prerequisites for Compensating Transactions

- A Coordinator must ...
 - Keep record of data commands or messages submitted to each participant
 - Keep record of the sequence of these calls
 - Have a way to associate each call with a unique Instance of the Long-Running Process (i.e. Conversation ID)
 - Be able to associate any Asynchronous Callback Message with a specific Process Instance
 - Know how to “play back” the appropriate compensations for a given Process Instance with the right data



Sounds Complicated ...

- It is ...
 - Compensations can fail too
 - The data may be visible between the time it is set by the original request and the time that the compensation occurs
- Options ...
 - You can “roll your own”
 - Many companies do
 - On the Microsoft platform you can use ...
 - Micro-Flows (e.g. Windows Workflow)
 - Orchestrations (e.g. Biztalk)





Conclusion



References

- Hohpe, Gregor; Woolf, Bobby; *Enterprise Integration Patterns. Designing, Building, and Deploying Messaging Solutions.* Addison Wesley
- Hohpe, Gregor; *Your Coffee Shop Doesn't Use Two-Phase Commit*
 - http://www.eaipatterns.com/docs/IEEE_Software_Design_2PC.pdf
- http://en.wikipedia.org/wiki/Two-phase_commit_protocol
- Brewer ,Dr. Eric A. *Towards Robust Distributed Systems.*
www.cs.berkeley.edu/~brewer/cs262b-2004/PODC-keynote.pdf
 - This presentation describes the CAP Theorem
- Transaction Management in the R* Distributed Database Management System
 - <http://www.cs.cmu.edu/~natassa/courses/15-823/F02/papers/p378-mohan.pdf>
- <http://www.DesignPatternsFor.Net>