

In [1]:

```
import pandas as pd
```

In [2]:

```
import matplotlib.pyplot as plt
import seaborn as sns
```

In [3]:

```
import warnings
warnings.filterwarnings('ignore')
```

In [4]:

```
df = pd.read_csv('attrition_employee.csv')
```


In [5]:

```
df.head()
```

Out[5]:

	id	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EnvironmentSatisfaction	
0	0	36	Travel_Frequently	599	Research & Development	24	3	Medical	1	4	.
1	1	35	Travel_Rarely	921	Sales	8	3	Other	1	1	.
2	2	32	Travel_Rarely	718	Sales	26	3	Marketing	1	3	.
3	3	38	Travel_Rarely	1488	Research & Development	2	3	Medical	1	3	.
4	4	50	Travel_Rarely	1017	Research & Development	5	4	Medical	1	2	.

5 rows × 35 columns




In [6]:

```
df.tail()
```

Out[6]:

	id	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EnvironmentSatisfac
1672	1672	30	Travel_Rarely	945	Sales	1	3	Life Sciences	1	
1673	1673	32	Travel_Rarely	1303	Research & Development	2	3	Life Sciences	1	
1674	1674	29	Travel_Frequently	1184	Human Resources	24	3	Human Resources	1	
1675	1675	36	Travel_Rarely	441	Sales	9	2	Marketing	1	
1676	1676	36	Travel_Rarely	1141	Research & Development	20	3	Life Sciences	1	

5 rows × 35 columns



In [7]:

```
df.shape
```

Out[7]:

(1677, 35)

In [8]:

```
df.columns
```

Out[8]:

```
Index(['id', 'Age', 'BusinessTravel', 'DailyRate', 'Department',
      'DistanceFromHome', 'Education', 'EducationField', 'EmployeeCount',
      'EnvironmentSatisfaction', 'Gender', 'HourlyRate', 'JobInvolvement',
      'JobLevel', 'JobRole', 'JobSatisfaction', 'MaritalStatus',
      'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked', 'Over18',
      'OverTime', 'PercentSalaryHike', 'PerformanceRating',
      'RelationshipSatisfaction', 'StandardHours', 'StockOptionLevel',
      'TotalWorkingYears', 'TrainingTimesLastYear', 'WorkLifeBalance',
      'YearsAtCompany', 'YearsInCurrentRole', 'YearsSinceLastPromotion',
      'YearsWithCurrManager', 'Attrition'],
      dtype='object')
```

In [9]:

```
df.duplicated().sum()
```

Out[9]:

```
0
```

In [10]:

```
df.isnull().sum()
```

Out[10]:

```
id                0
Age               0
BusinessTravel    0
DailyRate        0
Department        0
DistanceFromHome  0
Education         0
EducationField    0
EmployeeCount     0
EnvironmentSatisfaction  0
Gender            0
HourlyRate        0
JobInvolvement    0
JobLevel          0
JobRole           0
JobSatisfaction   0
MaritalStatus     0
MonthlyIncome     0
MonthlyRate       0
NumCompaniesWorked  0
Over18            0
OverTime          0
PercentSalaryHike  0
PerformanceRating  0
RelationshipSatisfaction  0
StandardHours     0
StockOptionLevel  0
TotalWorkingYears  0
TrainingTimesLastYear  0
WorkLifeBalance   0
YearsAtCompany    0
YearsInCurrentRole  0
YearsSinceLastPromotion  0
YearsWithCurrManager  0
Attrition         0
dtype: int64
```

In [11]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1677 entries, 0 to 1676
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    1677 non-null   int64
1   Age                                   1677 non-null   int64
2   BusinessTravel                       1677 non-null   object
3   DailyRate                            1677 non-null   int64
4   Department                           1677 non-null   object
5   DistanceFromHome                     1677 non-null   int64
6   Education                             1677 non-null   int64
7   EducationField                       1677 non-null   object
8   EmployeeCount                        1677 non-null   int64
9   EnvironmentSatisfaction              1677 non-null   int64
10  Gender                               1677 non-null   object
11  HourlyRate                           1677 non-null   int64
12  JobInvolvement                       1677 non-null   int64
13  JobLevel                             1677 non-null   int64
14  JobRole                              1677 non-null   object
15  JobSatisfaction                      1677 non-null   int64
16  MaritalStatus                       1677 non-null   object
17  MonthlyIncome                       1677 non-null   int64
18  MonthlyRate                          1677 non-null   int64
19  NumCompaniesWorked                  1677 non-null   int64
20  Over18                              1677 non-null   object
21  OverTime                            1677 non-null   object
22  PercentSalaryHike                   1677 non-null   int64
23  PerformanceRating                   1677 non-null   int64
24  RelationshipSatisfaction             1677 non-null   int64
25  StandardHours                       1677 non-null   int64
26  StockOptionLevel                    1677 non-null   int64
27  TotalWorkingYears                   1677 non-null   int64
28  TrainingTimesLastYear               1677 non-null   int64
29  WorkLifeBalance                     1677 non-null   int64
30  YearsAtCompany                      1677 non-null   int64
31  YearsInCurrentRole                  1677 non-null   int64
32  YearsSinceLastPromotion              1677 non-null   int64
33  YearsWithCurrManager                 1677 non-null   int64
34  Attrition                           1677 non-null   int64
dtypes: int64(27), object(8)
memory usage: 458.7+ KB
```

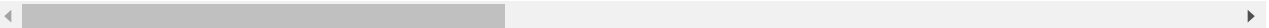
In [12]:

```
df.describe()
```

Out[12]:

	id	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	EnvironmentSatisfaction	HourlyRate	JobInvol
count	1677.000000	1677.000000	1677.000000	1677.000000	1677.000000	1677.0	1677.000000	1677.000000	1677
mean	838.000000	36.036971	892.749553	8.683959	2.937984	1.0	2.757901	67.798450	2
std	484.252517	8.507112	374.496259	7.826143	1.039078	0.0	1.086835	19.435928	0
min	0.000000	18.000000	107.000000	1.000000	1.000000	1.0	1.000000	30.000000	1
25%	419.000000	30.000000	589.000000	2.000000	2.000000	1.0	2.000000	51.000000	2
50%	838.000000	35.000000	890.000000	7.000000	3.000000	1.0	3.000000	69.000000	3
75%	1257.000000	41.000000	1223.000000	12.000000	4.000000	1.0	4.000000	84.000000	3
max	1676.000000	60.000000	3921.000000	29.000000	15.000000	1.0	4.000000	100.000000	4

8 rows × 27 columns



In [13]:

```
df.nunique()
```

Out[13]:

```
id                1677
Age               43
BusinessTravel    3
DailyRate        625
Department        3
DistanceFromHome  29
Education         6
EducationField    6
EmployeeCount     1
EnvironmentSatisfaction  4
Gender            2
HourlyRate        71
JobInvolvement    4
JobLevel          6
JobRole           9
JobSatisfaction   4
MaritalStatus     3
MonthlyIncome     895
MonthlyRate       903
NumCompaniesWorked 10
Over18            1
OverTime          2
PercentSalaryHike 15
PerformanceRating  2
RelationshipSatisfaction  4
StandardHours     1
StockOptionLevel  4
TotalWorkingYears 41
TrainingTimesLastYear  7
WorkLifeBalance   4
YearsAtCompany    34
YearsInCurrentRole 19
YearsSinceLastPromotion 16
YearsWithCurrManager 18
Attrition         2
dtype: int64
```

In [14]:

```
df_new = df[['BusinessTravel', 'Department', 'Education', 'EducationField',
             'EnvironmentSatisfaction', 'Gender', 'JobInvolvement', 'JobRole',
             'JobSatisfaction', 'MaritalStatus', 'NumCompaniesWorked', 'OverTime',
             'PerformanceRating', 'RelationshipSatisfaction', 'StockOptionLevel',
             'TrainingTimesLastYear', 'WorkLifeBalance', 'Attrition']]
```

In [15]:

```
for i in df_new.columns:
    print(i,':')
    print('Unique Values:', df_new[i].unique())
    print('Value Counts:')
    print(df_new[i].value_counts())
    print('\n')
```

```
BusinessTravel :
Unique Values: ['Travel_Frequently' 'Travel_Rarely' 'Non-Travel']
Value Counts:
Travel_Rarely      1290
Travel_Frequently   261
Non-Travel         126
Name: BusinessTravel, dtype: int64
```


```
Department :
Unique Values: ['Research & Development' 'Sales' 'Human Resources']
Value Counts:
Research & Development  1167
Sales                  471
Human Resources         39
Name: Department, dtype: int64
```

```
Education :
Unique Values: ['2' '4' '1' '3' '5' '16']
```

In [16]:

```
for i in df_new.columns:
    plt.figure(figsize=[15,7],)
    print('Countplot for:', i)
    sns.countplot(df_new[i], data = df_new, palette = 'hls')
    plt.xticks(rotation = 0)
    plt.show()
    print('\n')
```

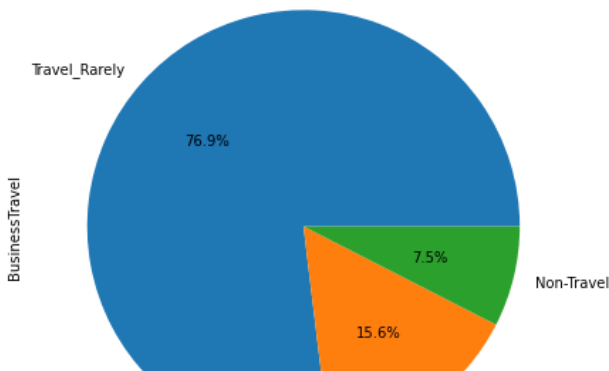
Countplot for: BusinessTravel



In [17]:

```
for i in df_new.columns:
    plt.figure(figsize=[15,7],)
    print('Pieplot for:', i)
    df_new[i].value_counts().plot(kind='pie',autopct='%1.1f%%')
    plt.show()
    print('\n')
```

Pieplot for: BusinessTravel



In [18]:

```
int_cols = [col for col in df.columns if df[col].dtype == 'int64']
print('Integer columns:', int_cols)

print('\n')

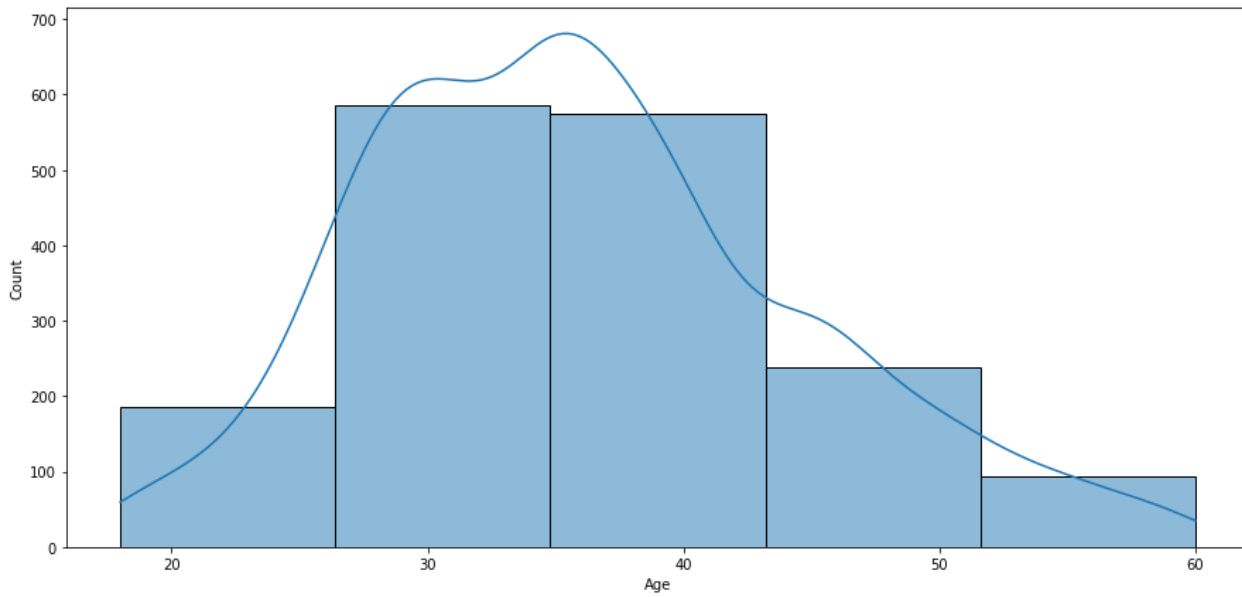
obj_cols = [col for col in df.columns if df[col].dtype == 'object']
print('Object columns:', obj_cols)
```

Integer columns: ['id', 'Age', 'DailyRate', 'DistanceFromHome', 'Education', 'EmployeeCount', 'EnvironmentSatisfaction', 'HourlyRate', 'JobInvolvement', 'JobLevel', 'JobSatisfaction', 'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked', 'PercentSalaryHike', 'PerformanceRating', 'RelationshipSatisfaction', 'StandardHours', 'StockOptionLevel', 'TotalWorkingYears', 'TrainingTimesLastYear', 'WorkLifeBalance', 'YearsAtCompany', 'YearsInCurrentRole', 'YearsSinceLastPromotion', 'YearsWithCurrManager', 'Attrition']

Object columns: ['BusinessTravel', 'Department', 'EducationField', 'Gender', 'JobRole', 'MaritalStatus', 'Over18', 'OverTime']

In [19]:

```
plt.figure(figsize=[15,7],)
sns.histplot(df['Age'], kde = 'True', bins = 5, palette = 'hls')
plt.xticks(rotation = 0)
plt.show()
```




In [20]:

```
import plotly.express as px
```


In [21]:

```
fig = px.box(df, x = 'Attrition', y = 'Age')
fig.show()
```




In [22]:

```
fig = px.box(df, x = 'Attrition', y = 'DailyRate')  
fig.show()
```




In [23]:

```
fig = px.box(df, x = 'Attrition', y = 'DistanceFromHome')  
fig.show()
```



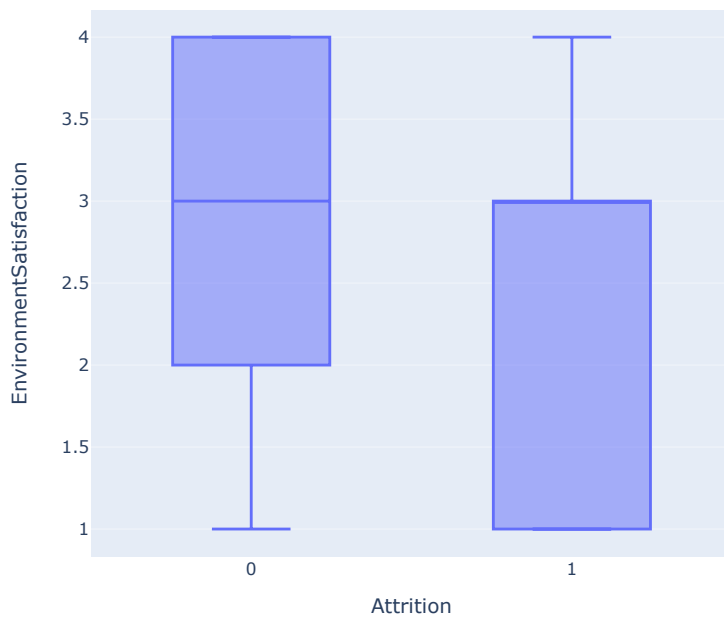
In [24]:

```
fig = px.box(df, x = 'Attrition', y = 'Education')  
fig.show()
```



In [25]:

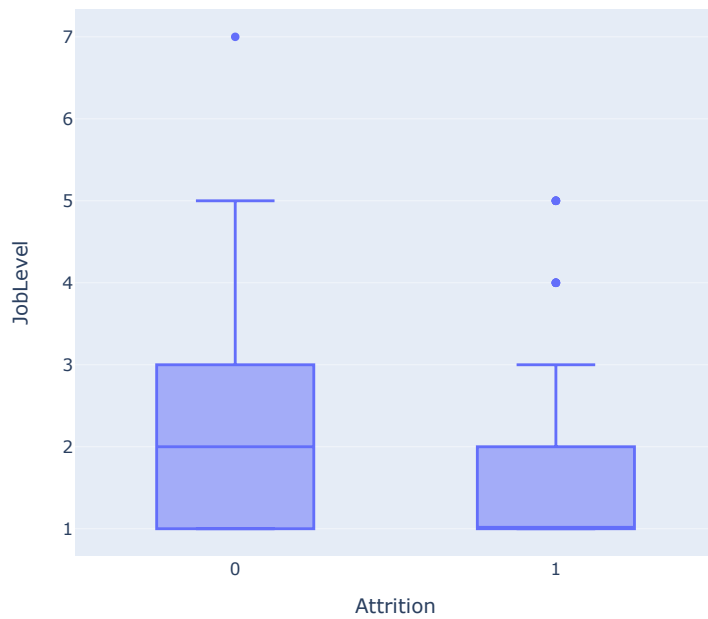
```
fig = px.box(df, x = 'Attrition', y = 'EnvironmentSatisfaction')  
fig.show()
```





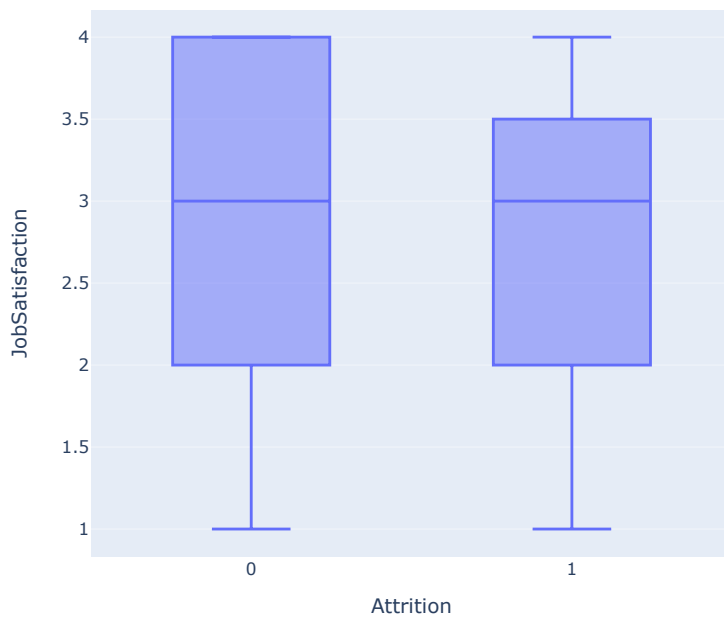
In [26]:

```
fig = px.box(df, x = 'Attrition', y = 'JobLevel')  
fig.show()
```



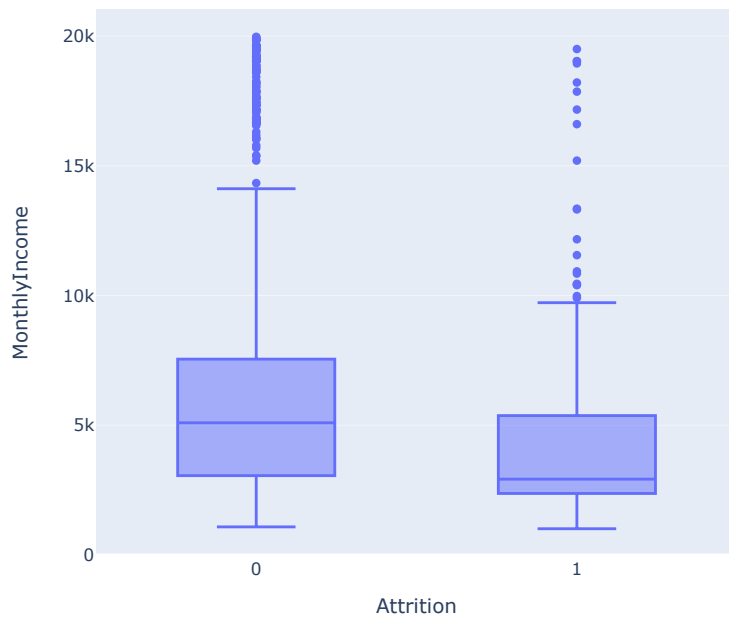
In [27]:

```
fig = px.box(df, x = 'Attrition', y = 'JobSatisfaction')  
fig.show()
```




In [28]:

```
fig = px.box(df, x = 'Attrition', y = 'MonthlyIncome')  
fig.show()
```



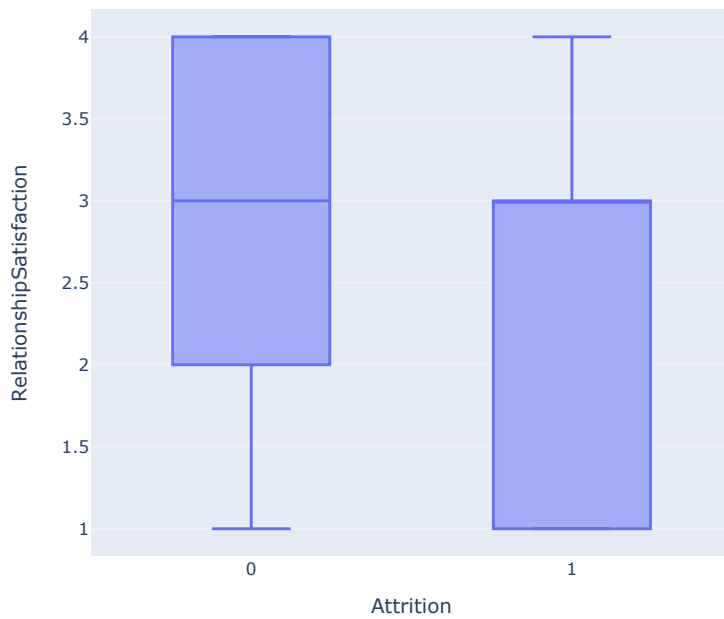
In [29]:

```
fig = px.box(df, x = 'Attrition', y = 'PercentSalaryHike')  
fig.show()
```



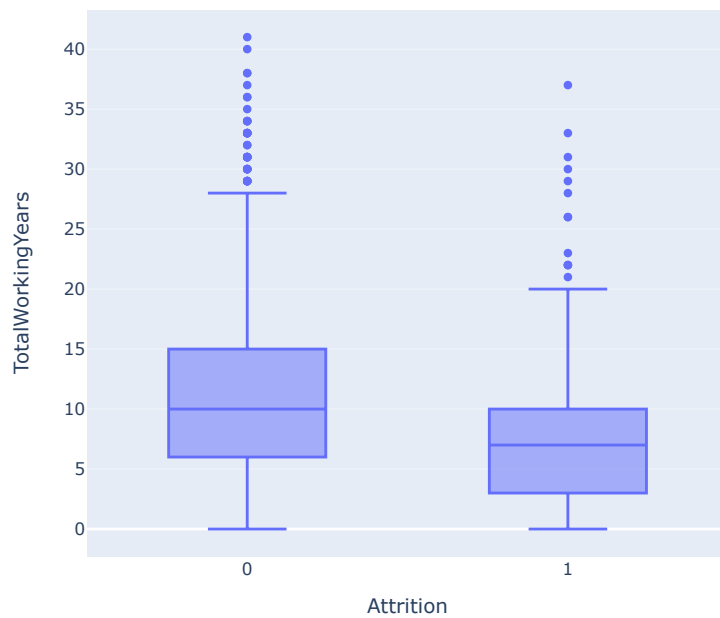
In [30]:

```
fig = px.box(df, x = 'Attrition', y = 'RelationshipSatisfaction')  
fig.show()
```



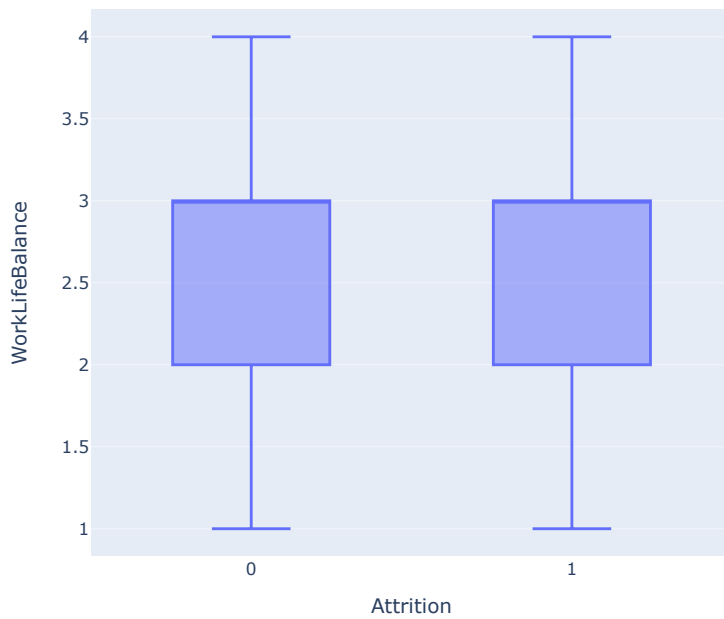
In [31]:

```
fig = px.box(df, x = 'Attrition', y = 'TotalWorkingYears')  
fig.show()
```



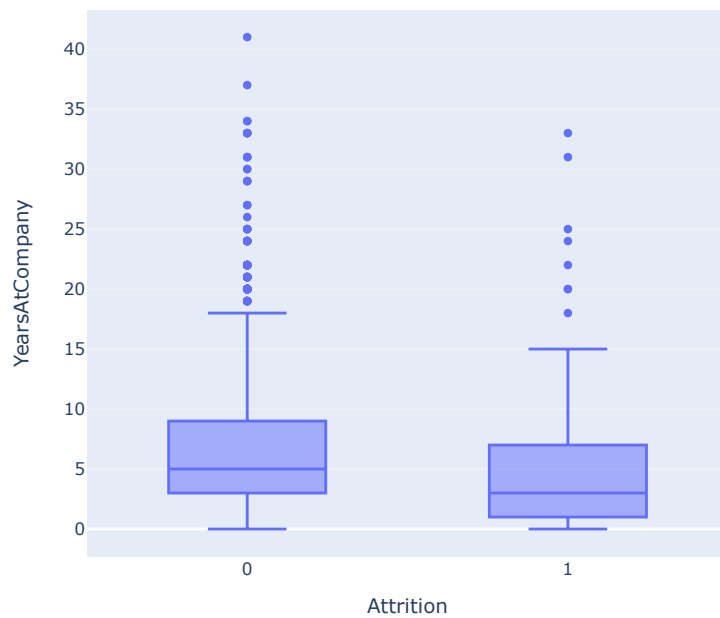
In [32]:

```
fig = px.box(df, x = 'Attrition', y = 'WorkLifeBalance')  
fig.show()
```



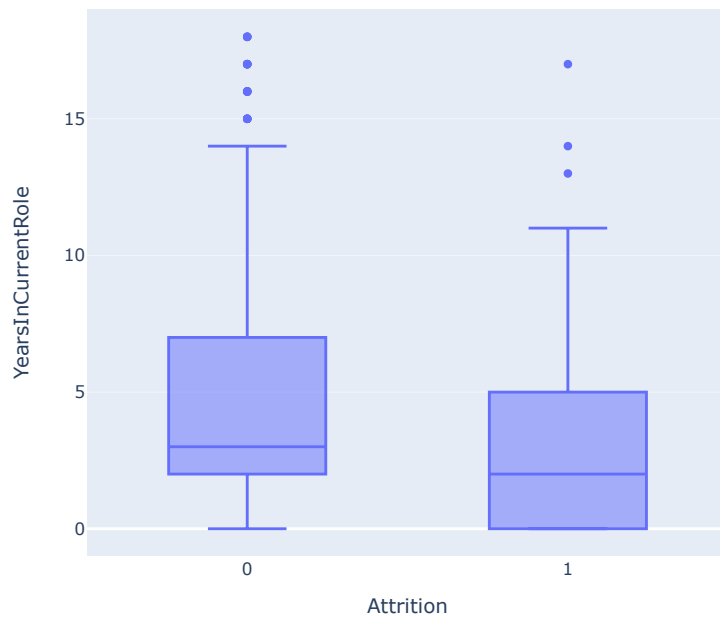
In [33]:

```
fig = px.box(df, x = 'Attrition', y = 'YearsAtCompany')  
fig.show()
```




In [34]:

```
fig = px.box(df, x = 'Attrition', y = 'YearsInCurrentRole')  
fig.show()
```



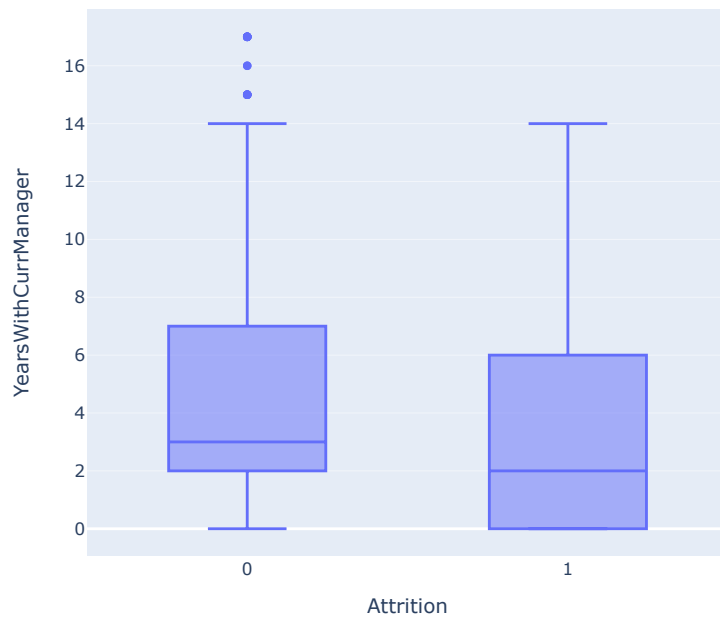
In [35]:

```
fig = px.box(df, x = 'Attrition', y = 'YearsSinceLastPromotion')  
fig.show()
```



In [36]:

```
fig = px.box(df, x = 'Attrition', y = 'YearsWithCurrManager')
fig.show()
```



In [37]:

```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
df['BusinessTravel'] = le.fit_transform(df['BusinessTravel'])
df['Department'] = le.fit_transform(df['Department'])
df['EducationField'] = le.fit_transform(df['EducationField'])
df['JobRole'] = le.fit_transform(df['JobRole'])
df['Gender'] = le.fit_transform(df['Gender'])
df['MaritalStatus'] = le.fit_transform(df['MaritalStatus'])
df['Over18'] = le.fit_transform(df['Over18'])
df['OverTime'] = le.fit_transform(df['OverTime'])
```

In [38]:

```
df.head()
```

Out[38]:

	id	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	EnvironmentSatisfaction	...
0	0	36	1	599	1	24	3	3	1	4	...
1	1	35	2	921	2	8	3	4	1	1	...
2	2	32	2	718	2	26	3	2	1	3	...
3	3	38	2	1488	1	2	3	3	1	3	...
4	4	50	2	1017	1	5	4	3	1	2	...

5 rows × 35 columns

In [39]:

```
corr = df.corr()
```


In [40]:

corr

Out[40]:

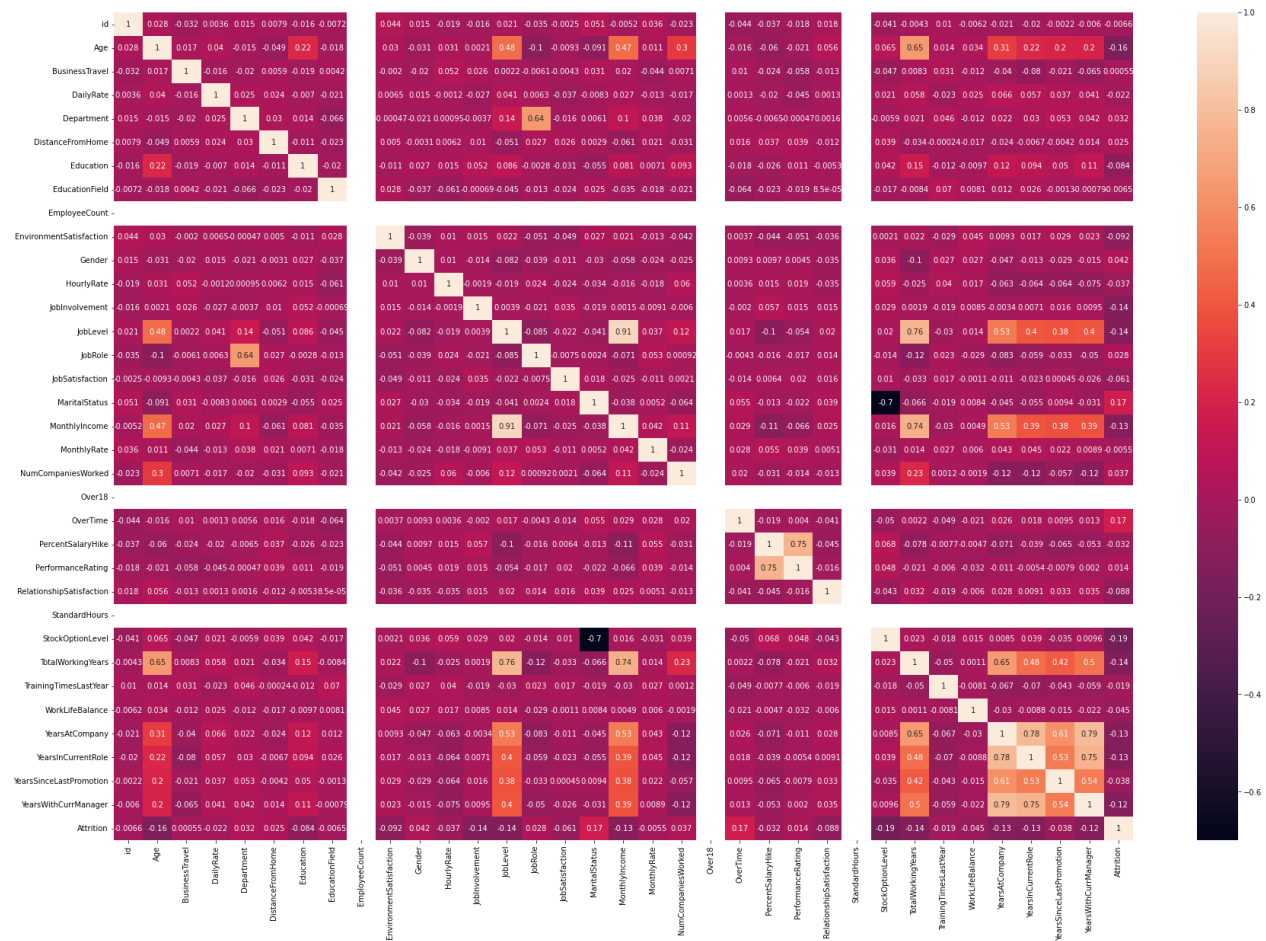
	id	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	Employee
id	1.000000	0.027563	-0.031831	0.003572	0.014652	0.007871	-0.015997	-0.007152	
Age	0.027563	1.000000	0.017462	0.039686	-0.014663	-0.049025	0.223545	-0.018181	
BusinessTravel	-0.031831	0.017462	1.000000	-0.015708	-0.020004	0.005933	-0.018889	0.004240	
DailyRate	0.003572	0.039686	-0.015708	1.000000	0.025107	0.024168	-0.007035	-0.021046	
Department	0.014652	-0.014663	-0.020004	0.025107	1.000000	0.029781	0.013881	-0.065715	
DistanceFromHome	0.007871	-0.049025	0.005933	0.024168	0.029781	1.000000	-0.011436	-0.023405	
Education	-0.015997	0.223545	-0.018889	-0.007035	0.013881	-0.011436	1.000000	-0.019603	
EducationField	-0.007152	-0.018181	0.004240	-0.021046	-0.065715	-0.023405	-0.019603	1.000000	
EmployeeCount	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
EnvironmentSatisfaction	0.044271	0.029557	-0.002002	0.006483	-0.000465	0.004959	-0.011189	0.028366	
Gender	0.014820	-0.031498	-0.019574	0.015338	-0.020542	-0.003123	0.027398	-0.036537	
HourlyRate	-0.018994	0.030628	0.051914	-0.001213	0.000948	0.006191	0.014862	-0.060686	
JobInvolvement	-0.016039	0.002101	0.026406	-0.026725	-0.003713	0.010035	0.052390	-0.000689	
JobLevel	0.021245	0.479015	0.002184	0.041369	0.138378	-0.051008	0.085823	-0.045207	
JobRole	-0.035070	-0.102034	-0.006113	0.006331	0.643463	0.026561	-0.002808	-0.012677	
JobSatisfaction	-0.002511	-0.009273	-0.004297	-0.037459	-0.015714	0.026309	-0.030686	-0.023946	
MaritalStatus	0.051459	-0.091312	0.030960	-0.008254	0.006123	0.002864	-0.054694	0.024995	
MonthlyIncome	-0.005224	0.470758	0.019567	0.027375	0.099545	-0.061019	0.081054	-0.034944	
MonthlyRate	0.036047	0.010959	-0.044104	-0.013332	0.037719	0.020542	0.007133	-0.018051	
NumCompaniesWorked	-0.022619	0.300044	0.007145	-0.017337	-0.020384	-0.031303	0.092789	-0.021372	
Over18	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
OverTime	-0.044391	-0.016417	0.010134	0.001339	0.005621	0.016349	-0.017772	-0.063571	
PercentSalaryHike	-0.036811	-0.060012	-0.024196	-0.020007	-0.006533	0.036970	-0.025858	-0.022959	
PerformanceRating	-0.018226	-0.021206	-0.057984	-0.045213	-0.000466	0.039206	0.010790	-0.019065	
RelationshipSatisfaction	0.018472	0.056115	-0.013382	0.001315	0.001612	-0.011868	-0.005253	0.000085	
StandardHours	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
StockOptionLevel	-0.041337	0.064863	-0.047337	0.021273	-0.005868	0.039086	0.041722	-0.016655	
TotalWorkingYears	-0.004288	0.648047	0.008332	0.058044	0.021490	-0.033573	0.153291	-0.008423	
TrainingTimesLastYear	0.010423	0.014303	0.031099	-0.023140	0.046363	-0.000239	-0.011924	0.070101	
WorkLifeBalance	-0.006236	0.034138	-0.012355	0.025152	-0.012399	-0.017184	-0.009697	0.008050	
YearsAtCompany	-0.020820	0.306628	-0.040091	0.066057	0.022345	-0.023564	0.116723	0.012349	
YearsInCurrentRole	-0.020064	0.219880	-0.079881	0.057011	0.029925	-0.006670	0.094065	0.025646	
YearsSinceLastPromotion	-0.002203	0.204357	-0.021144	0.037035	0.053127	-0.004215	0.050483	-0.001262	
YearsWithCurrManager	-0.005955	0.201601	-0.064886	0.040969	0.041859	0.013749	0.109573	-0.000790	
Attrition	-0.006598	-0.161044	0.000552	-0.022380	0.031996	0.024741	-0.084305	-0.006513	

35 rows × 35 columns



In [41]:

```
plt.figure(figsize=[30,20],)
sns.heatmap(corr, annot = True)
plt.show()
```



In [42]:

```
df1 = df.copy()
```

In [43]:

```
X = df1.drop(['Attrition'], axis = 1)
y = df1['Attrition']
```

In [44]:

```
from sklearn import preprocessing
scaler = preprocessing.StandardScaler()
X = scaler.fit_transform(X)
```

In [45]:

```
from sklearn.ensemble import ExtraTreesClassifier
import matplotlib.pyplot as plt
model = ExtraTreesClassifier()
model.fit(X,y)
print(model.feature_importances_)
```

```
[0.03198568 0.04044102 0.02410835 0.03375663 0.02151207 0.03273632
 0.0303788  0.02818188 0.          0.03522922 0.01800049 0.03412623
 0.04250434 0.02814186 0.03392777 0.03168117 0.03413634 0.03949013
 0.03528703 0.03095158 0.          0.04054669 0.03360471 0.01302054
 0.03729354 0.          0.03881399 0.03549759 0.02990741 0.03018425
 0.03598791 0.03669997 0.027918  0.03394849]
```

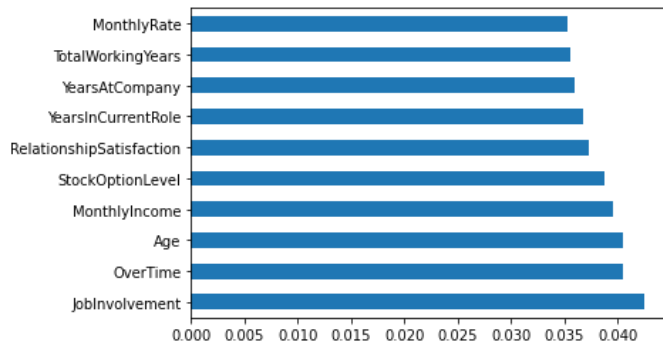


In [46]:

```
X = df1.iloc[:, :-1]
```

In [47]:

```
feat_importances = pd.Series(model.feature_importances_, index=X.columns)
feat_importances.nlargest(10).plot(kind='barh')
plt.show()
```



In [48]:

```
X_new = X[['StockOptionLevel', 'MonthlyIncome', 'JobInvolvement', 'OverTime',
            'Age', 'YearsInCurrentRole', 'YearsAtCompany', 'TotalWorkingYears',
            'JobRole', 'RelationshipSatisfaction']]
```

In [49]:

```
X_new.shape
```

Out[49]:

```
(1677, 10)
```

In [50]:

```
y.shape
```

Out[50]:

```
(1677,)
```

In [51]:

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X_new, y, test_size=0.30, random_state=0)
```

In [52]:

```
from sklearn.linear_model import LogisticRegression
log_r = LogisticRegression(random_state=0)
log_r.fit(X_train, y_train)
```

Out[52]:

```
LogisticRegression
LogisticRegression(random_state=0)
```

In [53]:

```
y_pred_lr = log_r.predict(X_test)
```

In [54]:

```
from sklearn.metrics import accuracy_score
```

In [55]:

```
accuracy_score(y_test, y_pred_lr)
```

Out[55]:

```
0.873015873015873
```

In [56]:

```
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred_lr)
```

In [57]:

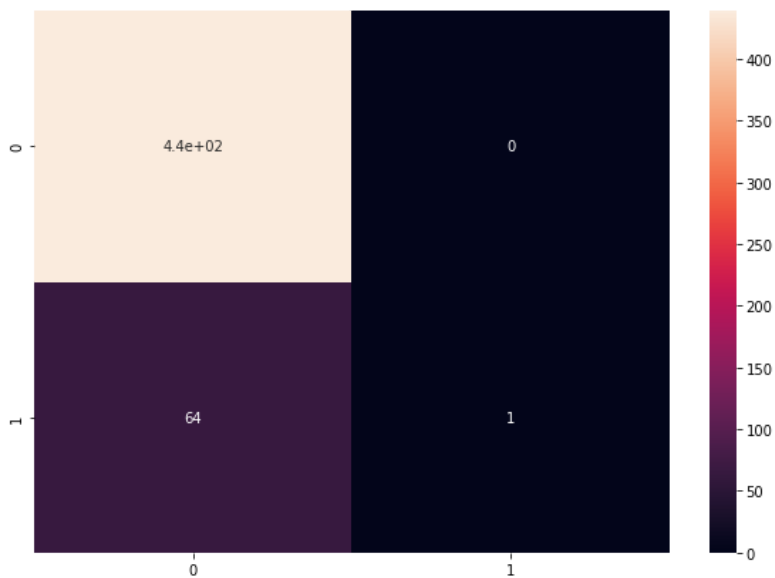
```
cm
```

Out[57]:

```
array([[439,  0],
       [ 64,  1]], dtype=int64)
```

In [58]:

```
plt.figure(figsize=[10,7],)
sns.heatmap(cm, annot = True)
plt.show()
```



In [59]:

```
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred_lr))
```

	precision	recall	f1-score	support
0	0.87	1.00	0.93	439
1	1.00	0.02	0.03	65
accuracy			0.87	504
macro avg	0.94	0.51	0.48	504
weighted avg	0.89	0.87	0.82	504

In [60]:

```
from sklearn.tree import DecisionTreeClassifier
dt = DecisionTreeClassifier()
dt.fit(X_train, y_train)
```

Out[60]:

```
DecisionTreeClassifier
```

In [61]:

```
y_pred_dt= dt.predict(X_test)
```

In [62]:

```
accuracy_score(y_test, y_pred_dt)
```

Out[62]:

```
0.8432539682539683
```

In [63]:

```
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred_dt)
```

In [64]:

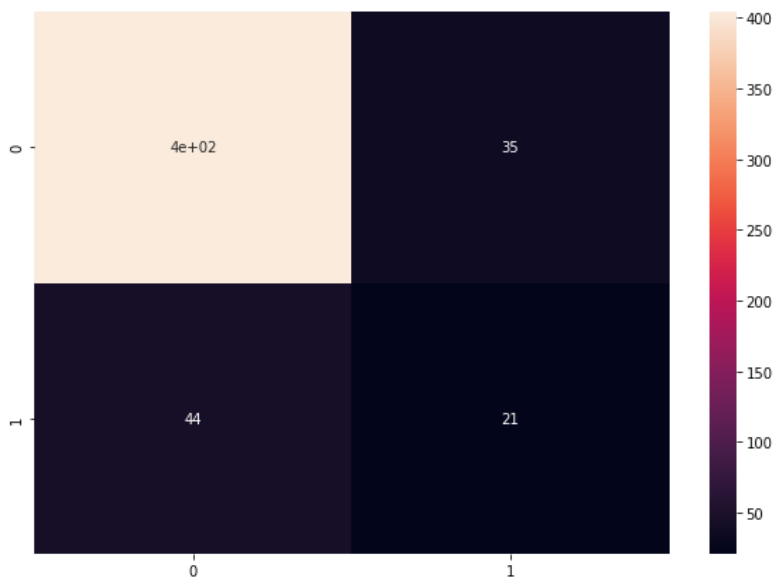
```
cm
```

Out[64]:

```
array([[404, 35],
       [ 44, 21]], dtype=int64)
```

In [65]:

```
plt.figure(figsize=[10,7],)
sns.heatmap(cm, annot = True)
plt.show()
```



In [66]:

```
print(classification_report(y_test, y_pred_dt))
```

	precision	recall	f1-score	support
0	0.90	0.92	0.91	439
1	0.38	0.32	0.35	65
accuracy			0.84	504
macro avg	0.64	0.62	0.63	504
weighted avg	0.83	0.84	0.84	504

In [67]:

```
from sklearn.ensemble import RandomForestClassifier
rf_c = RandomForestClassifier(n_estimators= 10, criterion="entropy")
rf_c.fit(X_train, y_train)
```

Out[67]:

```
RandomForestClassifier
RandomForestClassifier(criterion='entropy', n_estimators=10)
```

In [68]:

```
y_pred_rf_c= rf_c.predict(X_test)
```

In [69]:

```
accuracy_score(y_test, y_pred_rf_c)
```

Out[69]:

```
0.8611111111111112
```

In [70]:

```
cm= confusion_matrix(y_test, y_pred_rf_c)
```

In [71]:

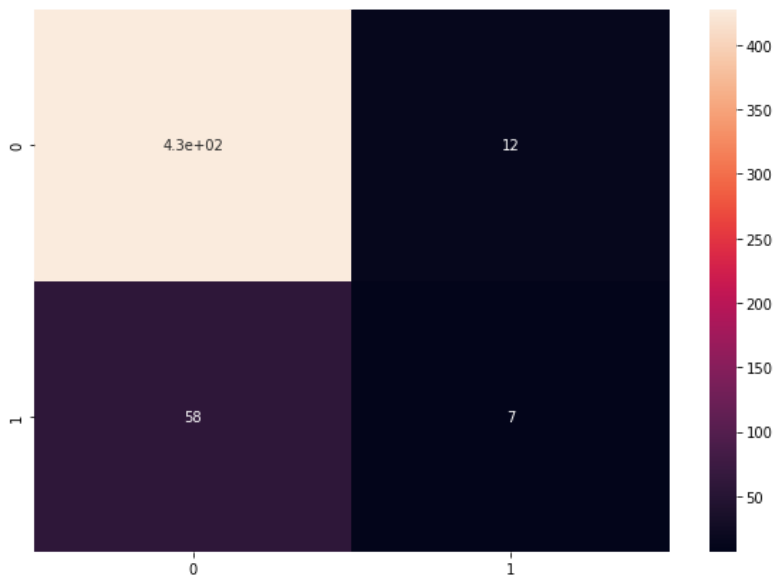
```
cm
```

Out[71]:

```
array([[427, 12],
       [ 58,  7]], dtype=int64)
```

In [72]:

```
plt.figure(figsize=[10,7],)
sns.heatmap(cm, annot = True)
plt.show()
```



In [73]:

```
print(classification_report(y_test, y_pred_rf_c))
```

	precision	recall	f1-score	support
0	0.88	0.97	0.92	439
1	0.37	0.11	0.17	65
accuracy			0.86	504
macro avg	0.62	0.54	0.55	504
weighted avg	0.81	0.86	0.83	504

In [74]:

```
from sklearn.model_selection import StratifiedKFold
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import GridSearchCV
```

In [75]:

```
folds = StratifiedKFold(n_splits = 5, shuffle = True, random_state = 40)
```

In [76]:

```
def grid_search(model, folds, params, scoring):
    grid_search = GridSearchCV(model,
                                cv=folds,
                                param_grid=params,
                                scoring=scoring,
                                n_jobs=-1, verbose=1)

    return grid_search
```

In [77]:

```
def print_best_score_params(model):
    print("Best Score: ", model.best_score_)
    print("Best Hyperparameters: ", model.best_params_)
```

In [78]:

```
log_reg = LogisticRegression()
log_params = {'C': [0.01, 1, 10],
              'penalty': ['l1', 'l2'],
              'solver': ['liblinear', 'newton-cg', 'saga']}
grid_search_log = grid_search(log_reg, folds, log_params, scoring=None)
grid_search_log.fit(X_train, y_train)
print_best_score_params(grid_search_log)
```

Fitting 5 folds for each of 18 candidates, totalling 90 fits  
Best Score: 0.8934315330060011  
Best Hyperparameters: {'C': 1, 'penalty': 'l2', 'solver': 'newton-cg'}

In [79]:

```
dtc = DecisionTreeClassifier(random_state=40)
dtc_params = {
    'max_depth': [5,10,20,30],
    'min_samples_leaf': [5,10,20,30]}
grid_search_dtc = grid_search(dtc, folds, dtc_params, scoring='roc_auc_ovr')
grid_search_dtc.fit(X_train, y_train)
print_best_score_params(grid_search_dtc)
```

Fitting 5 folds for each of 16 candidates, totalling 80 fits  
Best Score: 0.7096681152538641  
Best Hyperparameters: {'max\_depth': 5, 'min\_samples\_leaf': 10}

In [80]:

```
rfc = RandomForestClassifier(random_state=40, n_jobs = -1, oob_score=True)
rfc_params = {'max_depth': [10,20,30,40],
              'min_samples_leaf': [5,10,15,20,30],
              'n_estimators': [100,200,500,700]
            }
grid_search_rfc = grid_search(rfc, folds, rfc_params, scoring='roc_auc_ovr')
grid_search_rfc.fit(X_train, y_train)
print('OOB SCORE :', grid_search_rfc.best_estimator_.oob_score_)
```

Fitting 5 folds for each of 80 candidates, totalling 400 fits  
OOB SCORE : 0.896845694799659

In [81]:

```
from keras.models import Sequential
from keras.layers import Dense

model = Sequential()
model.add(Dense(64, input_dim=X_train.shape[1], activation='relu'))
model.add(Dense(32, activation='relu'))
model.add(Dense(1, activation='sigmoid'))

model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

In [82]:

```
model.fit(X_train, y_train, epochs=50, batch_size=32)
```

```
Epoch 1/50
37/37 [=====] - 1s 2ms/step - loss: 3.0136 - accuracy: 0.7801
Epoch 2/50
37/37 [=====] - 0s 1ms/step - loss: 3.1524 - accuracy: 0.8031
Epoch 3/50
37/37 [=====] - 0s 1ms/step - loss: 1.0043 - accuracy: 0.7621
Epoch 4/50
37/37 [=====] - 0s 1ms/step - loss: 4.4103 - accuracy: 0.8201
Epoch 5/50
37/37 [=====] - 0s 960us/step - loss: 3.6005 - accuracy: 0.8005
Epoch 6/50
37/37 [=====] - 0s 892us/step - loss: 3.6550 - accuracy: 0.8090
Epoch 7/50
37/37 [=====] - 0s 970us/step - loss: 1.6914 - accuracy: 0.8031
Epoch 8/50
37/37 [=====] - 0s 971us/step - loss: 3.9877 - accuracy: 0.7971
Epoch 9/50
37/37 [=====] - 0s 920us/step - loss: 4.7610 - accuracy: 0.8210
Epoch 10/50
37/37 [=====] - 0s 945us/step - loss: 2.3749 - accuracy: 0.8710
```

In [83]:

```
test_loss, test_acc = model.evaluate(X_test, y_test)
print('Test accuracy:', test_acc)
```

16/16 [=====] - 0s 945us/step - loss: 2.3749 - accuracy: 0.8710  
Test accuracy: 0.8710317611694336

In [84]:

```
df2 = pd.read_csv('employee_test.csv')
```

In [85]:

```
df2['BusinessTravel'] = le.fit_transform(df2['BusinessTravel'])
df2['Department'] = le.fit_transform(df2['Department'])
df2['EducationField'] = le.fit_transform(df2['EducationField'])
df2['JobRole'] = le.fit_transform(df2['JobRole'])
df2['Gender'] = le.fit_transform(df2['Gender'])
df2['MaritalStatus'] = le.fit_transform(df2['MaritalStatus'])
df2['Over18'] = le.fit_transform(df2['Over18'])
df2['OverTime'] = le.fit_transform(df2['OverTime'])
```

In [86]:

```
X_new = df2[['StockOptionLevel', 'MonthlyIncome', 'JobInvolvement', 'OverTime',  
            'Age', 'YearsInCurrentRole', 'YearsAtCompany', 'TotalWorkingYears',  
            'JobRole', 'RelationshipSatisfaction']]
```

In [87]:

```
X = scaler.fit_transform(X_new)
```

In [88]:

```
y_lr = grid_search_log.predict(X)
```

In [89]:

```
y_lr
```

Out[89]:

```
array([1, 1, 1, ..., 1, 1, 1], dtype=int64)
```

In [90]:

```
y_dt = grid_search_dtc.predict(X_new)
```

In [91]:

```
y_dt
```

Out[91]:

```
array([0, 0, 0, ..., 0, 0, 0], dtype=int64)
```

In [92]:

```
y_rfc = grid_search_rfc.predict(X_new)
```

In [93]:

```
y_rfc
```

Out[93]:

```
array([0, 0, 0, ..., 0, 0, 0], dtype=int64)
```

In [94]:

```
y_nn = model.predict(X_new)
```

```
35/35 [=====] - 0s 911us/step
```

In [95]:

```
y_nn
```

Out[95]:

```
array([[1.8085128e-04],  
       [1.2094801e-09],  
       [1.0392133e-05],  
       ...,  
       [4.5603679e-10],  
       [2.9436662e-09],  
       [2.5529080e-04]], dtype=float32)
```