

```

1 #from google.colab import drive
2 #drive.mount('/content/drive')

1 import zipfile
2 zip_ref = zipfile.ZipFile('/content/drive/MyDrive/Deep Learning/DogsNCats.zip', 'r')
3 zip_ref.extractall('/content')
4 zip_ref.close()

1 #importing libraries
2 import tensorflow as tf
3 from tensorflow import keras
4 from keras import Sequential
5 from keras.layers import Dense,Conv2D,MaxPooling2D,Flatten,BatchNormalization,Dropout

```

## ▼ Data preprocessing

```

1 train_ds =keras.utils.image_dataset_from_directory(
2     directory = '/content/DogsNCats/train',
3     labels = 'inferred',
4     label_mode = 'int',
5     batch_size = 32,
6     image_size = (256,256)
7 )
8
9 validation_ds = keras.utils.image_dataset_from_directory(
10     directory = '/content/DogsNCats/validation',
11     labels = 'inferred',
12     label_mode = 'int',
13     batch_size = 32,
14     image_size = (256,256)
15 )

Found 2171 files belonging to 2 classes.
Found 829 files belonging to 2 classes.

```

## ▼ normalizing

```

1 #normalize
2 def process(image, label):
3     image = tf.cast(image/255. ,tf.float32)
4     return image, label
5
6 train_ds = train_ds.map(process)
7 validation_ds = validation_ds.map(process)

```

## ▼ Model Selection and training

Creating the CNN model

```

1 # create CNN model
2
3 model = Sequential()
4
5 model.add(Conv2D(32,kernel_size=(3,3),padding='valid',activation='relu',input_shape=(256,256,3)))
6 model.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid'))
7
8 model.add(Conv2D(64,kernel_size=(3,3),padding='valid',activation='relu'))
9 model.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid'))
10
11 model.add(Conv2D(128,kernel_size=(3,3),padding='valid',activation='relu'))
12 model.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid'))
13
14 model.add(Flatten())

```

```
15
16 model.add(Dense(128,activation='relu'))
17 model.add(Dense(64,activation='relu'))
18 model.add(Dense(1,activation='sigmoid'))

1 model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 254, 254, 32)	896
max_pooling2d (MaxPooling2D)	(None, 127, 127, 32)	0
conv2d_1 (Conv2D)	(None, 125, 125, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 62, 62, 64)	0
conv2d_2 (Conv2D)	(None, 60, 60, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 30, 30, 128)	0
flatten (Flatten)	(None, 115200)	0
dense (Dense)	(None, 128)	14745728
dense_1 (Dense)	(None, 64)	8256
dense_2 (Dense)	(None, 1)	65

Total params: 14,847,297  
Trainable params: 14,847,297  
Non-trainable params: 0

```
1 model.compile(optimizer='adam', loss='binary_crossentropy',
2               metrics=['accuracy'])

1 history = model.fit(train_ds, epochs=15, validation_data = validation_ds)
```


```
Epoch 1/15
68/68 [=====] - 23s 205ms/step - loss: 0.7416 - accuracy: 0.5468 - val_loss: 0.7972 - val_accuracy: 0.3969
Epoch 2/15
68/68 [=====] - 8s 106ms/step - loss: 0.6747 - accuracy: 0.5753 - val_loss: 0.6805 - val_accuracy: 0.5271
Epoch 3/15
68/68 [=====] - 7s 90ms/step - loss: 0.5961 - accuracy: 0.6886 - val_loss: 0.9827 - val_accuracy: 0.4741
Epoch 4/15
68/68 [=====] - 8s 115ms/step - loss: 0.5386 - accuracy: 0.7296 - val_loss: 0.6482 - val_accuracy: 0.6478
Epoch 5/15
68/68 [=====] - 6s 90ms/step - loss: 0.4892 - accuracy: 0.7628 - val_loss: 0.7107 - val_accuracy: 0.6707
Epoch 6/15
68/68 [=====] - 8s 114ms/step - loss: 0.3363 - accuracy: 0.8494 - val_loss: 0.8032 - val_accuracy: 0.6960
Epoch 7/15
68/68 [=====] - 8s 106ms/step - loss: 0.2417 - accuracy: 0.9014 - val_loss: 0.9161 - val_accuracy: 0.6828
Epoch 8/15
68/68 [=====] - 8s 107ms/step - loss: 0.1845 - accuracy: 0.9217 - val_loss: 1.2605 - val_accuracy: 0.6550
Epoch 9/15
68/68 [=====] - 8s 112ms/step - loss: 0.1001 - accuracy: 0.9632 - val_loss: 1.8861 - val_accuracy: 0.6586
Epoch 10/15
68/68 [=====] - 8s 107ms/step - loss: 0.0879 - accuracy: 0.9668 - val_loss: 1.4187 - val_accuracy: 0.7021
Epoch 11/15
68/68 [=====] - 9s 127ms/step - loss: 0.0315 - accuracy: 0.9894 - val_loss: 1.8407 - val_accuracy: 0.7057
Epoch 12/15
68/68 [=====] - 10s 145ms/step - loss: 0.0116 - accuracy: 0.9968 - val_loss: 2.2902 - val_accuracy: 0.6996
Epoch 13/15
68/68 [=====] - 8s 106ms/step - loss: 0.0270 - accuracy: 0.9912 - val_loss: 1.7177 - val_accuracy: 0.7081
Epoch 14/15
68/68 [=====] - 9s 128ms/step - loss: 0.0541 - accuracy: 0.9811 - val_loss: 2.0408 - val_accuracy: 0.7214
Epoch 15/15
68/68 [=====] - 7s 105ms/step - loss: 0.0512 - accuracy: 0.9857 - val_loss: 1.9153 - val_accuracy: 0.6574
```

```
1 import matplotlib.pyplot as plt
2
3 plt.plot(history.history['accuracy'], color='red', label = 'train')
```

```

4 plt.plot(history.history['val_accuracy'], color='blue', label='validation')
5 plt.legend()
6 plt.show()


```



```

1 plt.plot(history.history['loss'], color='red', label = 'train')
2 plt.plot(history.history['val_loss'], color='blue', label='validation')
3 plt.legend()
4 plt.show()

```



#### ▼ ways to reduce overfitting:

- add more data
- data augmentation
- L1, L2 regularization
- batch normalization
- dropout

#### ▼ Batch Normalization and dropout

```

1 # create CNN model
2
3 model = Sequential()
4
5 model.add(Conv2D(32, kernel_size=(3,3), padding='valid', activation='relu', input_shape=(256, 256, 3)))
6 model.add(BatchNormalization())

```

```
7 model.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid'))
8
9 model.add(Conv2D(64,kernel_size=(3,3),padding='valid',activation='relu'))
10 model.add(BatchNormalization())
11 model.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid'))
12
13 model.add(Conv2D(128,kernel_size=(3,3),padding='valid',activation='relu'))
14 model.add(BatchNormalization())
15 model.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid'))
16
17 model.add(Flatten())
18
19 model.add(Dense(128,activation='relu'))
20 model.add(Dropout(0.1))
21 model.add(Dense(64,activation='relu'))
22 model.add(Dropout(0.1))
23 model.add(Dense(1,activation='sigmoid'))
```

```
1 model.summary()
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
conv2d_6 (Conv2D)	(None, 254, 254, 32)	896
batch_normalization_3 (Batch Normalization)	(None, 254, 254, 32)	128
max_pooling2d_6 (MaxPooling2D)	(None, 127, 127, 32)	0
conv2d_7 (Conv2D)	(None, 125, 125, 64)	18496
batch_normalization_4 (Batch Normalization)	(None, 125, 125, 64)	256
max_pooling2d_7 (MaxPooling2D)	(None, 62, 62, 64)	0
conv2d_8 (Conv2D)	(None, 60, 60, 128)	73856
batch_normalization_5 (Batch Normalization)	(None, 60, 60, 128)	512
max_pooling2d_8 (MaxPooling2D)	(None, 30, 30, 128)	0
flatten_2 (Flatten)	(None, 115200)	0
dense_6 (Dense)	(None, 128)	14745728
dropout_2 (Dropout)	(None, 128)	0
dense_7 (Dense)	(None, 64)	8256
dropout_3 (Dropout)	(None, 64)	0
dense_8 (Dense)	(None, 1)	65

=====  
Total params: 14,848,193  
Trainable params: 14,847,745  
Non-trainable params: 448  
=====

```
1 model.compile(optimizer='adam', loss='binary_crossentropy',
2               metrics=['accuracy'])
```

```
1 history = model.fit(train_ds, epochs=15, validation_data = validation_ds)
```

Epoch 1/15  
68/68 [=====] - 12s 127ms/step - loss: 4.9238 - accuracy: 0.5491 - val\_loss: 8.1851 - val\_accuracy: 0.4331  
Epoch 2/15  
68/68 [=====] - 9s 129ms/step - loss: 2.3095 - accuracy: 0.5813 - val\_loss: 9.9701 - val\_accuracy: 0.6031  
Epoch 3/15  
68/68 [=====] - 8s 112ms/step - loss: 1.2487 - accuracy: 0.5749 - val\_loss: 1.4024 - val\_accuracy: 0.5356  
Epoch 4/15

```

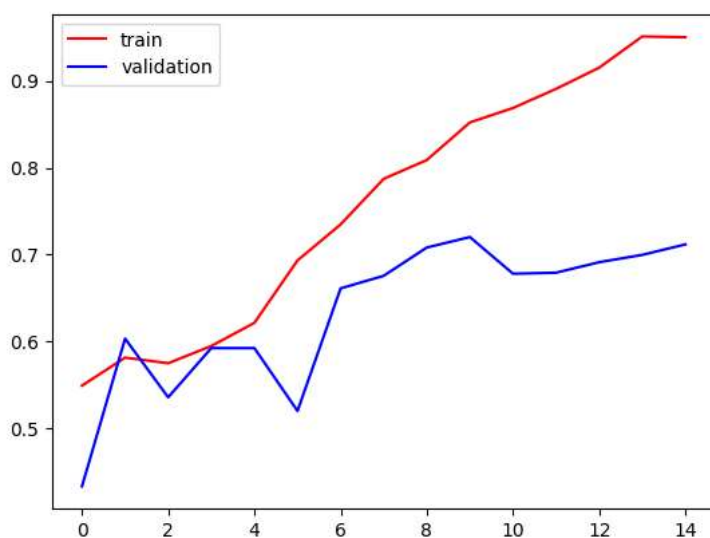
68/68 [=====] - 9s 116ms/step - loss: 1.0509 - accuracy: 0.5947 - val_loss: 1.0247 - val_accuracy: 0.5923
Epoch 5/15
68/68 [=====] - 10s 138ms/step - loss: 0.8082 - accuracy: 0.6214 - val_loss: 1.0102 - val_accuracy: 0.5923
Epoch 6/15
68/68 [=====] - 9s 128ms/step - loss: 0.6395 - accuracy: 0.6932 - val_loss: 0.7884 - val_accuracy: 0.5199
Epoch 7/15
68/68 [=====] - 9s 128ms/step - loss: 0.5472 - accuracy: 0.7347 - val_loss: 0.6379 - val_accuracy: 0.6610
Epoch 8/15
68/68 [=====] - 9s 126ms/step - loss: 0.4806 - accuracy: 0.7872 - val_loss: 0.7122 - val_accuracy: 0.6755
Epoch 9/15
68/68 [=====] - 12s 168ms/step - loss: 0.4413 - accuracy: 0.8088 - val_loss: 0.6571 - val_accuracy: 0.7081
Epoch 10/15
68/68 [=====] - 9s 129ms/step - loss: 0.3350 - accuracy: 0.8521 - val_loss: 0.6973 - val_accuracy: 0.7201
Epoch 11/15
68/68 [=====] - 9s 128ms/step - loss: 0.3074 - accuracy: 0.8687 - val_loss: 0.6488 - val_accuracy: 0.6779
Epoch 12/15
68/68 [=====] - 10s 136ms/step - loss: 0.2739 - accuracy: 0.8908 - val_loss: 0.7076 - val_accuracy: 0.6791
Epoch 13/15
68/68 [=====] - 9s 121ms/step - loss: 0.2190 - accuracy: 0.9152 - val_loss: 0.7889 - val_accuracy: 0.6912
Epoch 14/15
68/68 [=====] - 8s 112ms/step - loss: 0.1327 - accuracy: 0.9512 - val_loss: 0.9872 - val_accuracy: 0.6996
Epoch 15/15
68/68 [=====] - 9s 129ms/step - loss: 0.1339 - accuracy: 0.9503 - val_loss: 1.0129 - val_accuracy: 0.7117

```

```

1 import matplotlib.pyplot as plt
2
3 plt.plot(history.history['accuracy'], color='red', label = 'train')
4 plt.plot(history.history['val_accuracy'], color='blue', label='validation')
5 plt.legend()
6 plt.show()

```



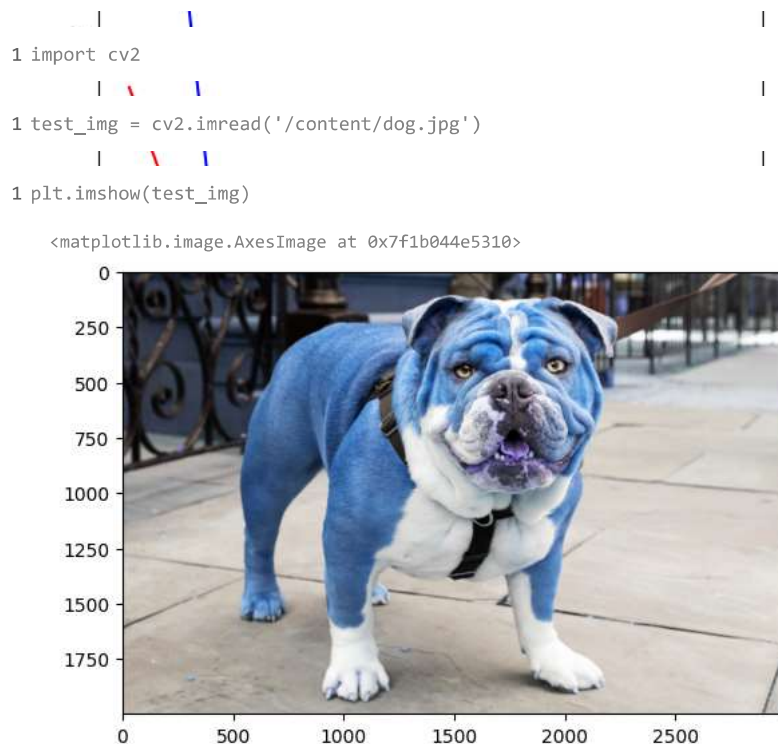
```

1 plt.plot(history.history['loss'], color='red', label = 'train')
2 plt.plot(history.history['val_loss'], color='blue', label='validation')
3 plt.legend()
4 plt.show()

```



## test case1



```
1 test_img.shape
```

(1999, 3000, 3)

```
1 test_img = cv2.resize(test_img, (256, 256))
```

```
1 test_input = test_img.reshape((1,256,256,3))
```

```
1 model.predict(test_input)
```

1/1 [=====] - 0s 19ms/step  
array([[1.]], dtype=float32)

```
1 test = model.predict(test_input)
```

```
2 test = test > 0.5
```

```
3
```

```
4 if (test == 0):
```

```
5     pred = 'cat'
```

```
6 else:
```

```
7     pred = 'dog'
```

```
8
```

```
9 print('our model says its a : ', pred)
```


1/1 [=====] - 0s 24ms/step  
our model says its a : dog

## test case2

```
1 test_img = cv2.imread('/content/cat.jpg')
```

```
1 plt.imshow(test_img)
```

<matplotlib.image.AxesImage at 0x7f1b0471db80>



```
1 test_img.shape
(720, 1280, 3)

1 test_img = cv2.resize(test_img, (256, 256))

1 test_input = test_img.reshape((1,256,256,3))

1 model.predict(test_input)
1/1 [=====] - 0s 92ms/step
array([[0.]], dtype=float32)

1 #model.predict(test_input)
2 test = model.predict(test_input)
3 test = test > 0.5
4
5 if (test == 0):
6     pred = 'cat'
7 else:
8     pred = 'dog'
9
10 print('our model says it a :', pred)

1/1 [=====] - 0s 173ms/step
our model says it a : cat

1
```