

Description

The English dataset is scraped from many different web resources. It consists of 199,002 verses, each of them is labeled with one of these four meters: Iambic, Trochee, Dactyl and Anapaestic. The Iambic class dominates the dataset; they are 186,809 Iambic verses, 5418 Trochee verses, 5378 Anapaestic verses, 1397 Dactyl verses.

Steps Included :

- 1. Checking Of Null Values
- 2. Removal Of StopWords
- 3. Removal Of Rare Words (Optional)
- 4. Cleansing Of Dataset
- 5. Stemming Using Porter Stemmer
- 6. Lemmatization (if Required)
- 7. Use Of Word Cloud
- 8. Finding the Frequency Of Words
- 9. Finding the Frequency Of Bi-Gram Words
- 10. Finding the Frequency Of Tri-Gram Words
- 11. Adding the Review Length and Word Count
- 12. Adding the Polarity
- 13. Rating Vs Polarity
- 14. Removing the Neutral Ratings
- 15. Use Of Count Vectorizer with Logistic Regression
- 16. Use Of TF-IDF Vectorizer with Logistic Regression
- 17. Use Of Ramdom Forest Classifier
- 18. Model Fitting
- 19. Model Evaluation

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
print('setup Completed^____^')
```

setup Completed^\_\_\_\_^

```
###! pip install --upgrade pandas
```

```
np.version.version
```

'1.21.6'

```
import pandas as pd
import numpy as np
import seaborn as sn
import matplotlib.pyplot as plt
import re
import nltk
from collections import Counter
from sklearn.feature_extraction.text import TfidfVectorizer,CountVectorizer
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import GaussianNB,MultinomialNB
```

```
from sklearn.svm import SVC
from sklearn import metrics
```

```
##!mkdir ~/.kaggle
```

```
##!cp /kaggle.json ~/.kaggle/
```

```
##! pip install kaggle
```

```
##! kaggle datasets download -d mohamedkhaledelsafty/english-poem-comprehensive-dataset-apcd
```

```
##! unzip /content/english-poem-comprehensive-dataset-apcd.zip
```

```
plt.style.use('dark_background')
```

```
full_english = pd.read_csv("/content/Full English PCD.csv")
```

```
english_poem = pd.read_csv("/content/Down-sampled English PCD.csv")
```

```
print(full_english.columns, full_english.columns)
```

```
Index(['Unnamed: 0', 'Verse', 'Meter', 'char_count'], dtype='object') Index(['Unnamed: 0', 'Verse', 'Meter', 'char_count'], dtype='object')
```

## ▼ Checking Of Null Values

```
train = full_english.sample(frac=0.8, random_state=0)
test = full_english.drop(train.index)
```

```
print(train.shape, test.shape)
```

```
(159202, 4) (39800, 4)
```

```
train.to_csv("/content/train.csv")
```

```
test.to_csv("/content/test.csv")
```

```
train.isnull().sum()
```

```
Unnamed: 0    0
Verse         0
Meter         0
char_count    0
dtype: int64
```

```
train.rename(columns={"Unnamed: 0" : "index"}, inplace=True)
```

```
test.rename(columns={"Unnamed: 0" : "index"}, inplace=True)
```

```
train['Verse'] = train['Verse'].astype(str)
test['Verse'] = test['Verse'].astype(str)
```

```
c = train.Meter.astype('category')
```

```
d = dict(enumerate(c.cat.categories))
```

```
print(d)

{0: 'anapaestic', 1: 'dactyl', 2: 'iambic', 3: 'trochaic'}
```

```
train['Meter'] = train.Meter.astype('category').cat.codes
```

```
test['Meter'] = test.Meter.astype('category').cat.codes
```

## ▼ Cleansing The Text - Making all text lowercase, remove text in square brackets,remove links,remove punctuation

```
print('the column data types is:',train['Verse'].dtypes)
```

```
the column data types is: object
```

```
print('the column data types is:',test['Verse'].dtypes)
```

```
the column data types is: object
```

```
train['Verse'] = train['Verse'].astype(str)
test['Verse'] = test['Verse'].astype(str)
```

```
import string
import re
def clean_text(text):
    '''Make text lowercase, remove text in square brackets,remove links,remove punctuation
    and remove words containing numbers.'''
    text = text.lower()
    text = re.sub('\[.*?\]', '', text)
    text = re.sub('https?://\S+|www\.\S+', '', text)
    text = re.sub('<.*?>+', '', text)
    text = re.sub('[%s]' % re.escape(string.punctuation), '', text)
    text = re.sub('\n', '', text)
    text = re.sub('\w*\d\w*', '', text)
    return text
```

```
train['Cleaned_Verse'] = train['Verse'].apply(lambda x: clean_text(x))
```

```
test['Cleaned_Verse'] = test['Verse'].apply(lambda x: clean_text(x))
```

## ▼ Removing StopWords

```
###! pip install nltk
```

```
import nltk
nltk.download('stopwords')
from nltk.corpus import stopwords
stop = stopwords.words('english')
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip.
```

```
train['Cleaned_Verse'] = train['Cleaned_Verse'].apply(lambda x: " ".join(x for x in x.split() if x not in stop))
train['Cleaned_Verse'].head()
```

```
69602          untainted vice art
133404      stung envy spleen diseasd
123426          plenty wantond cheek
28956      deck foster son fit need daring toil
```

```
30517      blushing sisterb saw pace along
Name: Cleaned_Verse, dtype: object

test['Cleaned_Verse'] = test['Cleaned_Verse'].apply(lambda x: " ".join(x for x in x.split() if x not in stop))
test['Cleaned_Verse'].head()
```

```
2          mean regardless yon midnight bell
10   bids heavns bright guard paraclete remove
21          anguish muses horror broods
27          taught heart glow god
30          fond soul impassiond rapt unveild
Name: Cleaned_Verse, dtype: object
```

## ▼ Remove the Rare Words

```
freq = pd.Series(' '.join(train['Cleaned_Verse']).split()).value_counts()
less_freq = list(freq[freq == 1].index)
```

```
train['Cleaned_Verse'] = train['Cleaned_Verse'].apply(lambda x: " ".join(x for x in x.split() if x not in less_freq))
train['Cleaned_Verse'].head(2)
```

```
69602      untainted vice art
133404      stung envy spleen
Name: Cleaned_Verse, dtype: object
```

```
freq = pd.Series(' '.join(train['Cleaned_Verse']).split()).value_counts()
less_freq = list(freq[freq == 1].index)
```

```
test['Cleaned_Verse'] = test['Cleaned_Verse'].apply(lambda x: " ".join(x for x in x.split() if x not in less_freq))
test['Cleaned_Verse'].head(2)
```

```
2          mean regardless yon midnight bell
10   bids heavns bright guard paraclete remove
Name: Cleaned_Verse, dtype: object
```

```
freq = pd.Series(' '.join(test['Cleaned_Verse']).split()).value_counts()
less_freq = list(freq[freq == 1].index)
```

```
from textblob import TextBlob, Word, Blobber
from nltk.stem import PorterStemmer
st = PorterStemmer()
```

```
import nltk
nltk.download('wordnet')
```

```
[nltk_data] Downloading package wordnet to /root/nltk_data...
True
```

```
from nltk.stem import WordNetLemmatizer
wordnet_lemmatizer = WordNetLemmatizer()
```

```
train['Cleaned_Verse'] = train['Cleaned_Verse'].apply(lambda x: " ".join([st.stem(word) for word in x.split()])))
```

```
test['Cleaned_Verse'] = test['Cleaned_Verse'].apply(lambda x: " ".join([st.stem(word) for word in x.split()])))
```

## ▼ Adding the length of the review and the word count of each Verse

```
train['Cleaned_Verse_len'] = train['Cleaned_Verse'].astype(str).apply(len)
train['word_count'] = train['Cleaned_Verse'].apply(lambda x: len(str(x).split()))
```

```
test['Cleaned_Verse_len'] = test['Cleaned_Verse'].astype(str).apply(len)
test['word_count'] = test['Cleaned_Verse'].apply(lambda x: len(str(x).split()))
```

Adding Polarity

Add one more feature called polarity. **Polarity** shows the sentiment of a piece of text. It counts the negative and positive words and determines the polarity. The value ranges from -1 to 1 where -1 represents the negative sentiment, 0 represents neutral and 1 represent positive sentiment.

```
train['polarity'] = train['Cleaned_Verse'].map(lambda text: TextBlob(text).sentiment.polarity)
train.head(2)
```

	index	Verse	Meter	char_count	Cleaned_Verse	Cleaned_Verse_len	word_count	polarity
69602	70451	and you untainted by the vice of art	2	6	untaint vice art	16	3	0.0
133404	136080	nor stung with envy nor with spleen diseases'd	2	6	stung envi spleen	17	3	0.0




```
test['polarity'] = test['Cleaned_Verse'].map(lambda text: TextBlob(text).sentiment.polarity)
test.head(2)
```

	index	Verse	Meter	char_count	Cleaned_Verse	Cleaned_Verse_len	word_count	polarity
2	2	what mean regardless of yon midnight bell	2	6	mean regardless yon midnight bell	33	5	-0.3125
10	10	bids heav'n's bright guard from paraclete remove	2	6	bid heavn bright guard paraclet remov	37	6	0.7000



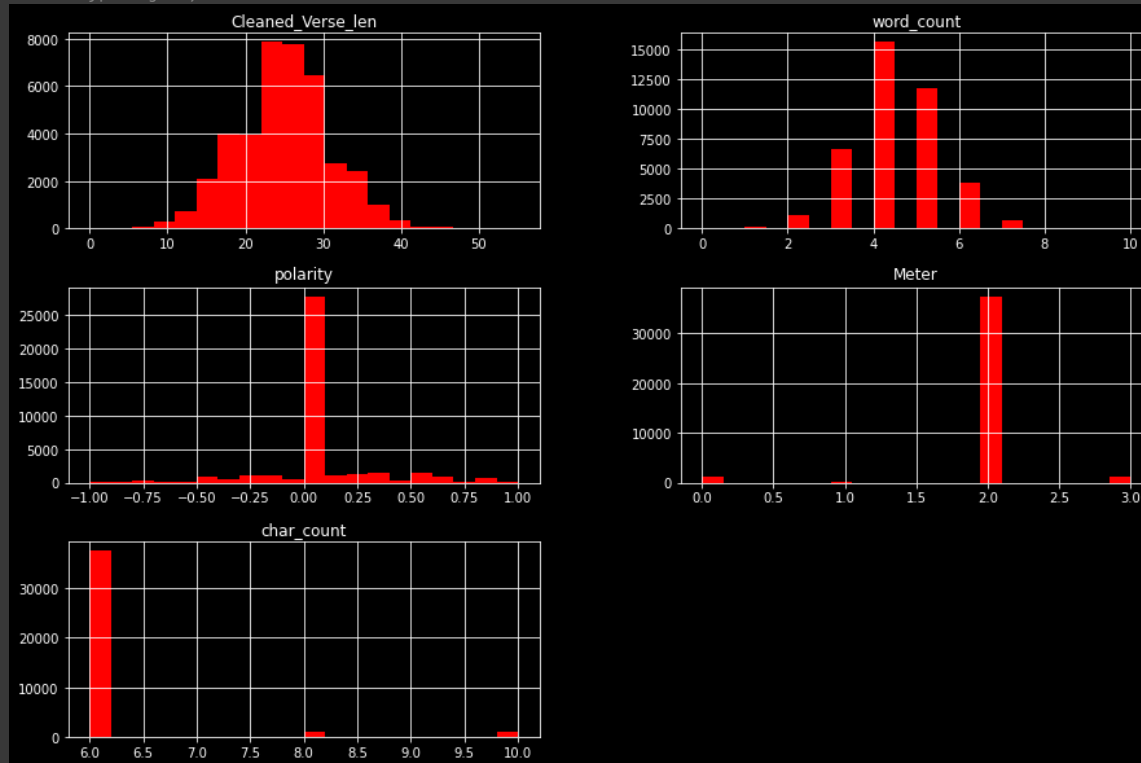
```
train[['Cleaned_Verse_len', 'word_count', 'polarity', 'Meter', 'char_count']].hist(bins=20, figsize=(15, 10), color='red')
```

```
array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7f27e0d9f250>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f27e193e6a0>],
      [<matplotlib.axes._subplots.AxesSubplot object at 0x7f27e1877ac0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f27e1663eb0>],
      [<matplotlib.axes._subplots.AxesSubplot object at 0x7f27e181f2b0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f27e1821700>]],
      dtype=object)
```



```
test[['Cleaned_Verse_len', 'word_count', 'polarity', 'Meter', 'char_count']].hist(bins=20, figsize=(15, 10), color='red')
```

```
array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7f27e18242b0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f27e11f0940>],
      [<matplotlib.axes._subplots.AxesSubplot object at 0x7f27e0e5bd90>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f27e179e1c0>],
      [<matplotlib.axes._subplots.AxesSubplot object at 0x7f27e17825b0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f27e128ca00>]],
      dtype=object)
```



```
train.columns
```

```
Index(['index', 'Verse', 'Meter', 'char_count', 'Cleaned_Verse',
      'Cleaned_Verse_len', 'word_count', 'polarity'],
      dtype='object')
```

```
condition_pol = train.groupby('Meter')['polarity'].agg([np.mean])
condition_pol.columns = ['polarity']
condition_pol = condition_pol.sort_values('polarity', ascending=False)
```

```
condition_pol = condition_pol.head(30)
```

```
condition_pol
```

	polarity
0	0.067479
3	0.058498
2	0.047953
1	0.038299

```
condition_pol = test.groupby('Meter')['polarity'].agg([np.mean])
condition_pol.columns = ['polarity']
condition_pol = condition_pol.sort_values('polarity', ascending=False)
condition_pol = condition_pol.head(30)
condition_pol
```

```
polarity
```

	polarity
1	0.080843
0	0.070277
2	0.047639
3	0.041828

```
Meter
```

1	0.080843
---	----------

0	0.070277
---	----------

2	0.047639
---	----------

3	0.041828
---	----------

## WordCloud:

Wordcloud is a common and beautiful visualization for text data to plot the frequency of words. You may need to install wordcloud if you do not have it already, using this command:

```
####! pip install wordcloud
```

```
train.columns
```

```
Index(['index', 'Verse', 'Meter', 'char_count', 'Cleaned_Verse',  
      'Cleaned_Verse_len', 'word_count', 'polarity'],  
      dtype='object')
```

```
text_train = " ".join(review for review in train.Cleaned_Verse)
```

```
text_test = " ".join(review for review in test.Cleaned_Verse)
```

```
from wordcloud import WordCloud, STOPWORDS, ImageColorGenerator
stopwords = set(STOPWORDS)
stopwords = stopwords.union(["ha", "thi", "now", "onli", "im", "becaus", "wa", "will", "even", "go", "realli", "didnt", "abl"])
wordcl = WordCloud(stopwords = stopwords, background_color='white', max_font_size = 50, max_words = 5000).generate(text_train)
plt.figure(figsize=(14, 12))
plt.imshow(wordcl, interpolation='bilinear')
plt.axis('off')
plt.show()
```



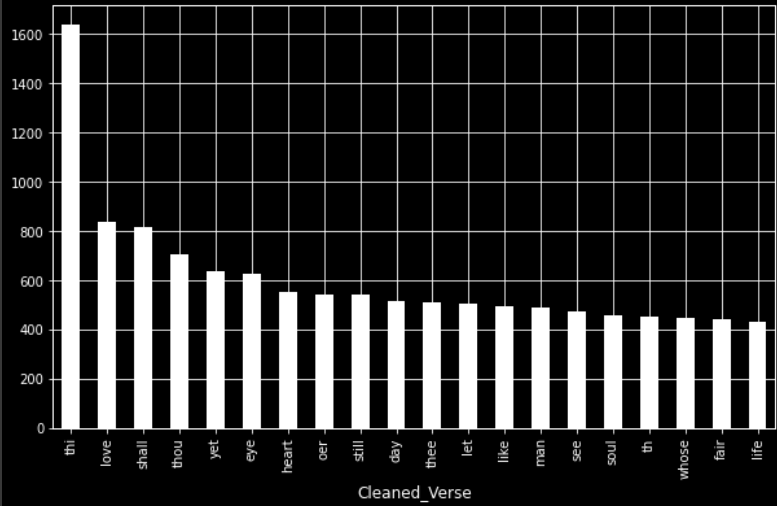


```
def get_top_n_words(corpus, n=None):
    vec=CountVectorizer().fit(corpus)
    bag_of_words = vec.transform(corpus)
    sum_words = bag_of_words.sum(axis=0)
    words_freq = [(word, sum_words[0, idx]) for word, idx in vec.vocabulary_.items()]
    words_freq =sorted(words_freq, key = lambda x: x[1], reverse=True)
    return words_freq[:n]

common_words = get_top_n_words(test['Cleaned_Verse'], 20)
df2 = pd.DataFrame(common_words, columns = ['Cleaned_Verse', 'count'])
df2.head()
```

	Cleaned_Verse	count
0	thi	1637
1	love	840
2	shall	817
3	thou	704
4	yet	635

```
df2.groupby('Cleaned_Verse').sum()['count'].sort_values(ascending=False).plot(kind='bar',color='white',figsize = (10, 6))
xlabel = 'Top Words'
ylabel = 'Count'
title = 'BarChart represent the Top Words Frequency'
plt.show()
```



## ▼ Bi-Grams


```
def get_top_n_bigram(corpus, n=None):
    vec = CountVectorizer(ngram_range=(2,2)).fit(corpus)
    bag_of_words = vec.transform(corpus)
    sum_words = bag_of_words.sum(axis=0)
    words_freq = [(word, sum_words[0, idx]) for word, idx in vec.vocabulary_.items()]
    words_freq =sorted(words_freq, key = lambda x: x[1], reverse=True)
    return words_freq[:n]

common_words2 = get_top_n_bigram(train['Cleaned_Verse'], 30)
df3 = pd.DataFrame(common_words2, columns=['Cleaned_Verse', "Count"])
df3.head()
```



	Cleaned_Verse	Count
0	thou art	170
1	ten thousand	126
2	let us	123
3	yet still	116
4	thou hast	102

```
df3.groupby('Cleaned_Verse').sum()['Count'].sort_values(ascending=False).plot(kind='bar',figsize=(12,6), color='white')
xlabel = "Bigram Words"
ylabel = "Count"
title = "Bar chart of Bigrams Frequency"
plt.show()
```



```
def get_top_n_bigram(corpus, n=None):
    vec = CountVectorizer(ngram_range=(2,2)).fit(corpus)
    bag_of_words = vec.transform(corpus)
    sum_words = bag_of_words.sum(axis=0)
    words_freq = [(word, sum_words[0, idx]) for word, idx in vec.vocabulary_.items()]
    words_freq =sorted(words_freq, key = lambda x: x[1], reverse=True)
    return words_freq[:n]


common_words3 = get_top_n_bigram(test['Cleaned_Verse'], 30)
df4 = pd.DataFrame(common_words3, columns=['Cleaned_Verse', "Count"])
df4.head()
```



	Cleaned_Verse	Count
0	thou art	54
1	let us	32
2	thou canst	30
3	art thou	28
4	hast thou	25

```
df4.groupby('Cleaned_Verse').sum()['Count'].sort_values(ascending=False).plot(kind='bar',figsize=(12,6), color='white')
xlabel = "Bigram Words"
ylabel = "Count"
```

```
title = "Bar chart of Bigrams Frequency"
```



```
train.columns
```

```
Index(['index', 'Verse', 'Meter', 'char_count', 'Cleaned_Verse',
      'Cleaned_Verse_len', 'word_count', 'polarity'],
      dtype='object')
```

```
X_train = train["Cleaned_Verse"]
```

```
y_train = train["Meter"]
```

```
X_test = test["Cleaned_Verse"]
```

```
y_test = test["Meter"]
```

```
print(X_train.shape, y_train.shape)
print(X_test.shape, y_test.shape)
```

```
(159202,) (159202,)
(39800,) (39800,)
```

## Count Vectorizer

```
from sklearn.feature_extraction.text import CountVectorizer
# fit the countvectorizer to the training data:
vect = CountVectorizer().fit(X_train)
```

```
len(vect.get_feature_names()[:2000])
#len(vect.get_feature_names())
```

```
/usr/local/lib/python3.8/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function get_feature_names is deprecated; get_feature_names is deprecated in 1.0 and will be removed in 1.2. Please use
warnings.warn(msg, category=FutureWarning)
2000
```

```
X_train_vectorized = vect.transform(X_train)
```

```
from sklearn.linear_model import LogisticRegression
```

```
# Train the model
model = LogisticRegression()
model.fit(X_train_vectorized, y_train)
```

```
/usr/local/lib/python3.8/dist-packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

```
Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
    n_iter_i = _check_optimize_result(
LogisticRegression()
```

```
from sklearn.metrics import accuracy_score
```

```
# Predict the transformed test documents
predictions = model.predict(vect.transform(X_test))
print('Accuracy: ', accuracy_score(y_test, predictions))
```

```
Accuracy: 0.9407286432160804
```

## ▼ TF-IDF Vectorizer

```
from sklearn.feature_extraction.text import TfidfVectorizer
```

```
vect = TfidfVectorizer(min_df=5).fit(X_train)
len(vect.get_feature_names())
```

```
/usr/local/lib/python3.8/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function get_feature_names is deprecated; get_feature_names is deprecated in 1.0 and will be removed in 1.2. Please
warnings.warn(msg, category=FutureWarning)
8931
```

```
X_train_vectorized = vect.transform(X_train)
```

```
model = LogisticRegression()
model.fit(X_train_vectorized, y_train)
```

```
predictions = model.predict(vect.transform(X_test))
```

```
print('Accuracy: ', accuracy_score(y_test, predictions))
```

```
/usr/local/lib/python3.8/dist-packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

```
Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
    n_iter_i = _check_optimize_result(
Accuracy: 0.9397989949748744
```


```
feature_names = np.array(vect.get_feature_names())
```

```
sorted_tfidf_index = X_train_vectorized.max(0).toarray()[0].argsort()
```

```
print('Smallest tfidf:\n{}\n'.format(feature_names[sorted_tfidf_index[:10]]))
print('Largest tfidf: \n{}'.format(feature_names[sorted_tfidf_index[:-11:-1]]))
```

```
Smallest tfidf:
['casa' 'pr' 'basil' 'notari' 'forgd' 'talon' 'udder' 'raw' 'somer'
 'lybia']
```

```
Largest tfidf:
['dabbl' 'rather' 'street' 'turn' 'turnd' 'flew' 'avaric' 'fli' 'ill'
 'seat']
/usr/local/lib/python3.8/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function get_feature_names is deprecated; get_feature_names is deprecated in 1.0 and will be removed in 1.2. Please
warnings.warn(msg, category=FutureWarning)
```



## ▼ Random Forest Classifier

```
from sklearn.ensemble import RandomForestClassifier
```

```
model = RandomForestClassifier()
model.fit(X_train_vectorized, y_train)
```

```
RandomForestClassifier()
```

```
predictions = model.predict(vect.transform(X_test))
```

```
print('Accuracy: ', accuracy_score(y_test, predictions))
```

```
Accuracy:  0.9353015075376885
```

```
predictions = pd.DataFrame(predictions)
```

## ▼ Gradient Boosting Classifier

```
from sklearn.ensemble import GradientBoostingClassifier
```

```
model_gb = GradientBoostingClassifier()
model_gb.fit(X_train_vectorized, y_train)
```

```
GradientBoostingClassifier()
```

```
predictions_gb = model_gb.predict(vect.transform(X_test))
```

```
print('Accuracy: ', accuracy_score(y_test, predictions))
```

```
Accuracy:  0.9353015075376885
```

```
predictions_gb.shape
```

```
(39800,)
```

```
testvalue = pd.read_csv("/content/test.csv")
```

```
testvalue.shape
```

```
(39800, 5)
```

```
output = pd.concat([test, predictions], axis = 1)
```

```
output.rename(columns = { 0 : 'Predict'}, inplace = True)
```

```
output.head(2)
```

	index		Verse	Meter	char_count		Cleaned_Verse	Cleaned_Verse_len	word_count	polarity	Predict
	2	2.0	what mean regardless of yon midnight bell	2.0	6.0		mean regardless yon midnight bell	33.0	5.0	-0.3125	2.0
	10	10.0	bids heav'n's bright guard from paraclete remove	2.0	6.0		bid heavn bright guard paraclet remov	37.0	6.0	0.7000	2.0

```
output.Predict.value_counts()
```

```
2.0    39357
3.0      290
0.0     121
1.0       32
Name: Predict, dtype: int64
```

```
bw = output["Predict"].value_counts()
```

```
plt.figure(figsize=(10,6), dpi=80, facecolor='white')
plt.style.use('seaborn-white')
explode = (0.1, 0.1, 0.2, 0.2)
)
labels = ['iambic', 'trochaic', 'anapaestic', 'dactyl']
colors = ['red', 'magenta', 'orange', 'blue']
bw.plot.pie(shadow = True, colors=colors, autopct='%1.1f%%', wedgeprops = {'linewidth' : 6}, startangle=140)
plt.title('Predictions Results')
plt.legend(labels, loc="best", fontsize=15)
plt.axis('equal')
plt.show()
```

