

# Outlier Detection & Log Transformation



## Table of Contents

- 1. Data Processing and EDA [Train Data]
- 2. Univariate Analysis [numerical data]
- 3. Label Encoding [categorical data]
- 4. Processing Test data
- 5. Data Processing [Test Data]
- 6. Model Training
- 6. Model Evaluation [Confusion Matrix]
- 7. Submission



## Preprocessing and EDA

- 1. Preprocessing
- 2. EDA

```
In [16]: import numpy as np
import pandas as pd
import os
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.graph_objects as go
from scipy.stats import gaussian_kde
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier,
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
import plotly.express as px
```

```
import plotly.figure_factory as ff
from sklearn.metrics import confusion_matrix, classification_report
import warnings
warnings.filterwarnings('ignore')

RUN_CNT = 0
MAIN_FOLDER = "/kaggle/input/playground-series-s4e10"
TEST = os.path.join(MAIN_FOLDER, 'test.csv')
TRAIN = os.path.join(MAIN_FOLDER, 'train.csv')
```

```
In [17]: def IQR(train_df, column, range1, range2) -> pd.DataFrame:
    Q1 = train_df[column].quantile(range1)
    Q3 = train_df[column].quantile(range2)
    IQR = Q3 - Q1

    print("Q1 = {}, Q3 = {}, IQR = {}".format(Q1, Q3, IQR))

    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    print("lower_bound = {}, upper_bound = {}".format(lower_bound, upper_bound))
    train_df_filtered = train_df[(train_df[column] >= lower_bound) & (train_df[c
    return train_df_filtered

def bar_plotter(df, column_name):

    plt.figure(figsize=(8, 6))
    sns.histplot(df[column_name], bins=30, kde=True, color='blue')
    plt.title(f'Distribution of {column_name}')
    plt.xlabel(column_name)
    plt.ylabel('Frequency')
    plt.show()

def scatter(df, column):
    col = df[column]
    plt.scatter(x=list(range(len(col))), y=list(col), marker='.', color='red', a
    plt.xlabel("Datapoint")
    plt.ylabel("frequemcy of person age")
    plt.show()

def outlier_remove(df, column, range1=0.22, range2=0.99) -> pd.DataFrame:
    scatter(df, column)
    final_df = IQR(df, column, range1=range1, range2=range2)
    print("Scatter after outlier removal")
    scatter(final_df, column)
    return final_df

def compare_series_barchart_compact(series1, series2, column_name):
    if series1.shape != series2.shape:
        raise ValueError("Series must have the same length")

    # Combine the series into a DataFrame for plotting
    df = pd.DataFrame({'DF1': series1, 'DF2': series2})

    # Melt the DataFrame to Long format for plotting
    melted_df = df.reset_index().melt(id_vars='index', var_name='DataFrame', val

    # Plotting
    plt.figure(figsize=(10, 6))
```

```
sns.barplot(x='index', y='Values', hue='DataFrame', data=melted_df, palette=

# Adjust bar width and layout
plt.title(f"train vs test dataset comparison with {column_name}")
plt.xticks(rotation=45)
plt.xlabel('Index')
plt.tight_layout()
plt.show()
```

```
In [25]: import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd

# Apply a dark background globally
plt.style.use('dark_background')

def IQR(train_df, column, range1, range2) -> pd.DataFrame:
    Q1 = train_df[column].quantile(range1)
    Q3 = train_df[column].quantile(range2)
    IQR = Q3 - Q1

    print("Q1 = {}, Q3 = {}, IQR = {}".format(Q1, Q3, IQR))

    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    print("lower_bound = {}, upper_bound = {}".format(lower_bound, upper_bound))
    train_df_filtered = train_df[(train_df[column] >= lower_bound) & (train_df[c
    return train_df_filtered

def bar_plotter(df, column_name):
    plt.figure(figsize=(8, 6))
    sns.set_style("dark")
    sns.histplot(df[column_name], bins=30, kde=True, color='gray') # Gray color
    plt.title(f'Distribution of {column_name}', color='gray') # Gray title
    plt.xlabel(column_name, color='gray') # Gray label
    plt.ylabel('Frequency', color='gray') # Gray label
    plt.xticks(color='gray') # Gray ticks
    plt.yticks(color='gray') # Gray ticks
    plt.show()

def scatter(df, column):
    col = df[column]
    plt.scatter(x=list(range(len(col))), y=list(col), marker='.', color='gray',
    plt.xlabel("Datapoint", color='gray') # Gray label
    plt.ylabel("Frequency of person age", color='gray') # Gray label
    plt.xticks(color='gray') # Gray ticks
    plt.yticks(color='gray') # Gray ticks
    plt.show()

def outlier_removal(df, column, range1=0.22, range2=0.99) -> pd.DataFrame:
    scatter(df, column)
    final_df = IQR(df, column, range1=range1, range2=range2)
    print("Scatter after outlier removal")
    scatter(final_df, column)
    return final_df

def compare_series_barchart_compact(series1, series2, column_name):
    if series1.shape != series2.shape:
        raise ValueError("Series must have the same length")
```

```

# Combine the series into a DataFrame for plotting
df = pd.DataFrame({'DF1': series1, 'DF2': series2})

# Melt the DataFrame to Long format for plotting
melted_df = df.reset_index().melt(id_vars='index', var_name='DataFrame', val

# Plotting
plt.figure(figsize=(10, 6))
sns.set_style("dark")
sns.barplot(x='index', y='Values', hue='DataFrame', data=melted_df, palette=

# Adjust bar width and Layout
plt.title(f"Train vs Test Dataset Comparison with {column_name}", color='gra
plt.xlabel('Index', color='gray')
plt.xticks(rotation=45, color='gray') # Gray ticks
plt.yticks(color='gray') # Gray ticks
plt.tight_layout()
plt.show()

```

```

In [26]: original_df_train = pd.read_csv(TRAIN)
train_df = original_df_train.applymap(lambda x: x + 0.00001 if x == 0 else x)
original_df_test = pd.read_csv(TEST)
test_df = original_df_test.applymap(lambda x: x + 0.00001 if x == 0 else x)
train_df_copy = train_df.copy()
test_df_copy = test_df.copy()

train_df.head()

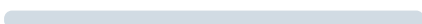
```

```

Out[26]:

```

	id	person_age	person_income	person_home_ownership	person_emp_length	loan_status
0	0.00001	37	35000	RENT	0.00001	EMERGENCY
1	1.00000	22	56000	OWN	6.00000	PAID
2	2.00000	29	28800	OWN	8.00000	PAID
3	3.00000	30	70000	RENT	14.00000	PAID
4	4.00000	22	60000	RENT	2.00000	PAID

<  >

---->  Preprocessing 

```

In [27]: train_df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 58645 entries, 0 to 58644
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    58645 non-null  float64
1   person_age                           58645 non-null  int64
2   person_income                         58645 non-null  int64
3   person_home_ownership                 58645 non-null  object
4   person_emp_length                     58645 non-null  float64
5   loan_intent                           58645 non-null  object
6   loan_grade                           58645 non-null  object
7   loan_amnt                            58645 non-null  int64
8   loan_int_rate                         58645 non-null  float64
9   loan_percent_income                   58645 non-null  float64
10  cb_person_default_on_file             58645 non-null  object
11  cb_person_cred_hist_length            58645 non-null  int64
12  loan_status                           58645 non-null  float64
dtypes: float64(5), int64(4), object(4)
memory usage: 5.8+ MB

```

In [28]: `test_df.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 39098 entries, 0 to 39097
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    39098 non-null  int64
1   person_age                           39098 non-null  int64
2   person_income                         39098 non-null  int64
3   person_home_ownership                 39098 non-null  object
4   person_emp_length                     39098 non-null  float64
5   loan_intent                           39098 non-null  object
6   loan_grade                           39098 non-null  object
7   loan_amnt                            39098 non-null  int64
8   loan_int_rate                         39098 non-null  float64
9   loan_percent_income                   39098 non-null  float64
10  cb_person_default_on_file             39098 non-null  object
11  cb_person_cred_hist_length            39098 non-null  int64
dtypes: float64(3), int64(5), object(4)
memory usage: 3.6+ MB

```

In [29]: `train_df.describe()`

Out[29]:

	id	person_age	person_income	person_emp_length	loan_amnt	l
<b>count</b>	58645.000000	58645.000000	5.864500e+04	58645.000000	58645.000000	5
<b>mean</b>	29322.000000	27.550857	6.404617e+04	4.701016	9217.556518	
<b>std</b>	16929.497605	6.033216	3.793111e+04	3.959783	5563.807384	
<b>min</b>	0.000010	20.000000	4.200000e+03	0.000010	500.000000	
<b>25%</b>	14661.000000	23.000000	4.200000e+04	2.000000	5000.000000	
<b>50%</b>	29322.000000	26.000000	5.800000e+04	4.000000	8000.000000	
<b>75%</b>	43983.000000	30.000000	7.560000e+04	7.000000	12000.000000	
<b>max</b>	58644.000000	123.000000	1.900000e+06	123.000000	35000.000000	


In [30]: `test_df.describe()`

Out[30]:

	id	person_age	person_income	person_emp_length	loan_amnt	l
<b>count</b>	39098.000000	39098.000000	3.909800e+04	39098.000000	39098.000000	3
<b>mean</b>	78193.500000	27.566781	6.406046e+04	4.687070	9251.466188	
<b>std</b>	11286.764749	6.032761	3.795583e+04	3.868393	5576.254680	
<b>min</b>	58645.000000	20.000000	4.000000e+03	0.000010	700.000000	
<b>25%</b>	68419.250000	23.000000	4.200000e+04	2.000000	5000.000000	
<b>50%</b>	78193.500000	26.000000	5.800000e+04	4.000000	8000.000000	
<b>75%</b>	87967.750000	30.000000	7.588500e+04	7.000000	12000.000000	
<b>max</b>	97742.000000	94.000000	1.900000e+06	42.000000	35000.000000	

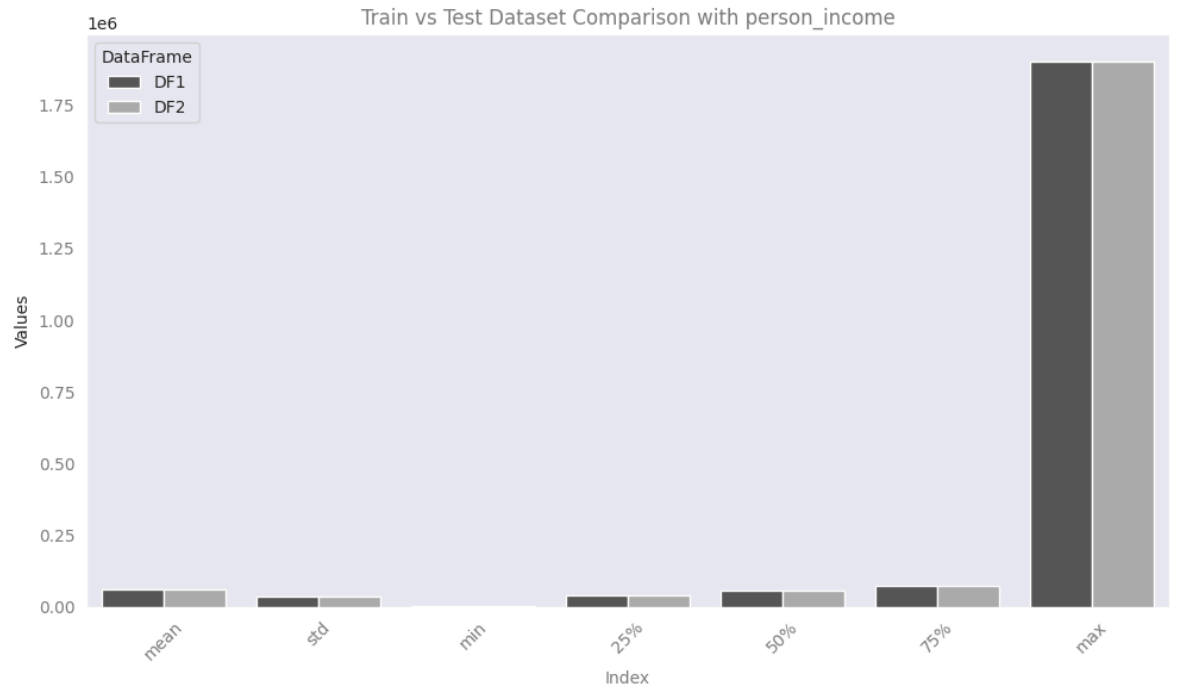
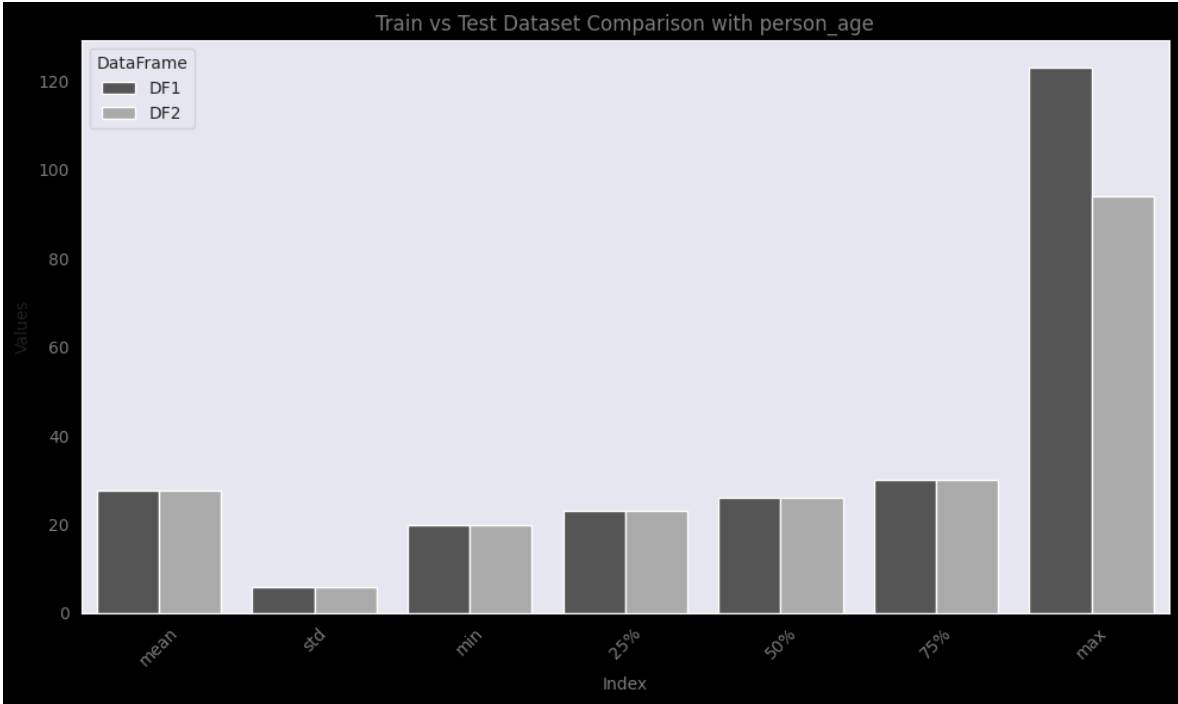


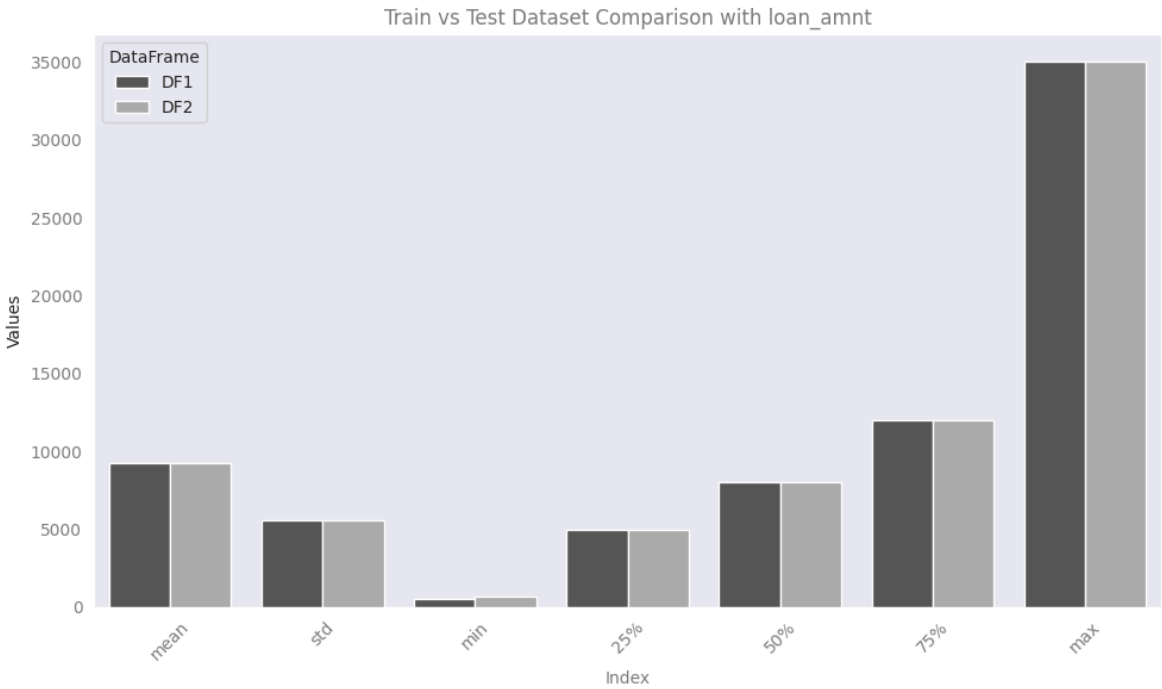
```
In [31]: train_report = train_df.describe().drop(['id', 'loan_status'], axis='columns')
test_report = test_df.describe().drop(['id'], axis='columns')
test_report.index
```

```
Out[31]: Index(['count', 'mean', 'std', 'min', '25%', '50%', '75%', 'max'], dtype='object')
```

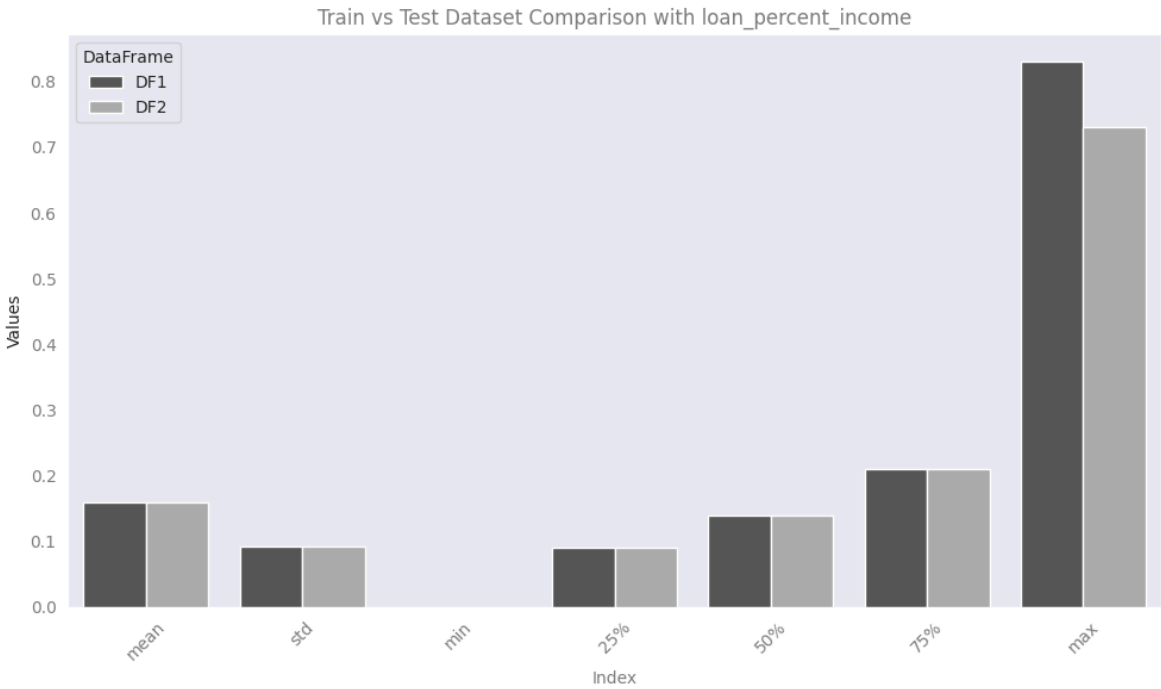
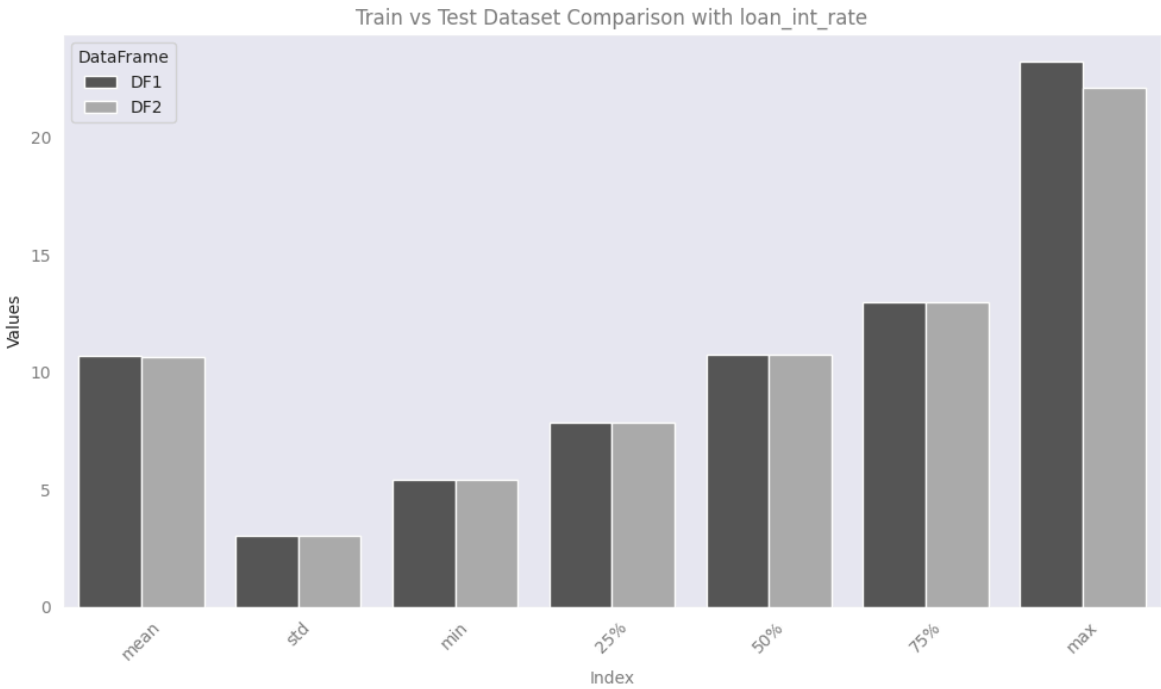
---->  **Train Data vs Test Data** 

```
In [32]: for row in train_report.columns:
compare_series_barchart_compact(train_report[row][1:], test_report[row][1:],
```












## Univariate Analysis


- **Outlier Detection:** Identifying extreme values in the dataset that may skew model performance.
- **Log Transformation:** Applying a logarithmic function to reduce skewness in data distribution.

---->  **Person Age Column** 

```
In [33]: df1 = outlier_remover(train_df, 'person_age', 0.22, 0.99)
```



Q1 = 23.0, Q3 = 49.0, IQR = 26.0  
lower\_bound = -16.0, upper\_bound = 88.0  
Scatter after outlier removal



## Outlier Detection

The Interquartile Range (IQR) is a measure of statistical dispersion that describes the spread of the middle 50% of a dataset. It is calculated as the difference between the third


quartile (Q3) and the first quartile (Q1):

1. **First Quartile (Q1):** The 25th percentile, meaning 25% of the data points fall below this value.
2. **Third Quartile (Q3):** The 75th percentile, meaning 75% of the data points fall below this value.
3. **IQR Formula:**  
$$\text{IQR} = Q3 - Q1$$

## Purpose:

- IQR is used to identify outliers, with values falling below  $(Q1 - 1.5 \times \text{IQR})$  or above  $(Q3 + 1.5 \times \text{IQR})$  considered as outliers.

```
In [34]: bar_plotter(df1, 'person_age')
```



## Right skewed data handling


To handle a right-skewed histogram, you can apply the following techniques:

1. **Log Transformation:** Apply  $\log(x)$  to compress larger values and stretch smaller ones.
2. **Square Root Transformation:** Use  $\sqrt{x}$  to reduce skewness.
3. **Reciprocal Transformation:** Apply  $1/x$  to handle extreme skewness.
4. **Power Transformation (Box-Cox):** Use `scipy.stats.boxcox` to normalize the data.

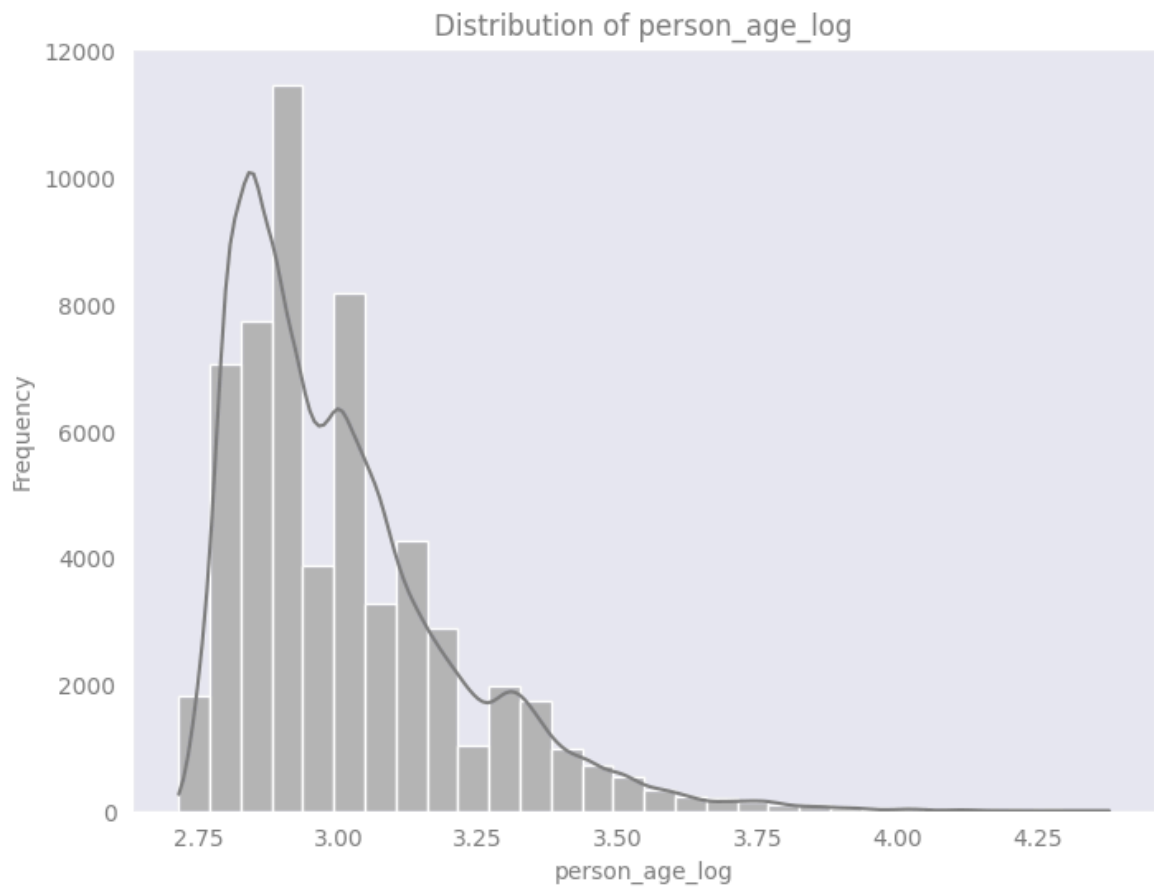
5. **Binning**: Group data into intervals to reduce skewness.

These techniques help in making the distribution more symmetric for statistical analysis.

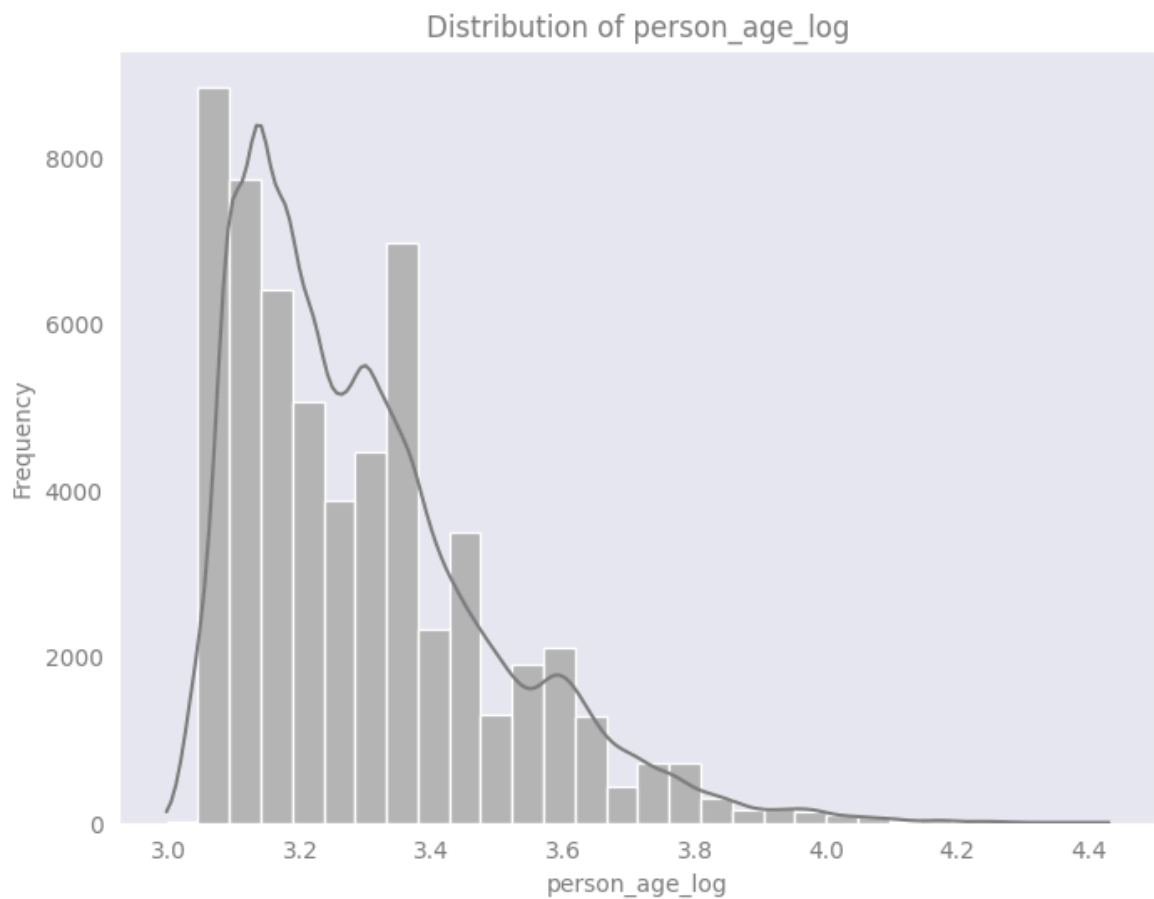
```
In [35]: # square root
df1.loc[:, 'person_age_log'] = df1['person_age'].apply(lambda x: np.sqrt(x))
bar_plotter(df=df1, column_name='person_age_log')
```



```
In [36]: # cube root
df1.loc[:, 'person_age_log'] = df1['person_age'].apply(lambda x: x**(1/3))
bar_plotter(df=df1, column_name='person_age_log')
```




```
In [37]: # log transformation"
df1.loc[:, 'person_age_log'] = df1['person_age'].apply(lambda x: np.log(x))
bar_plotter(df=df1, column_name='person_age_log')
```




----->  **Person Income Column** 

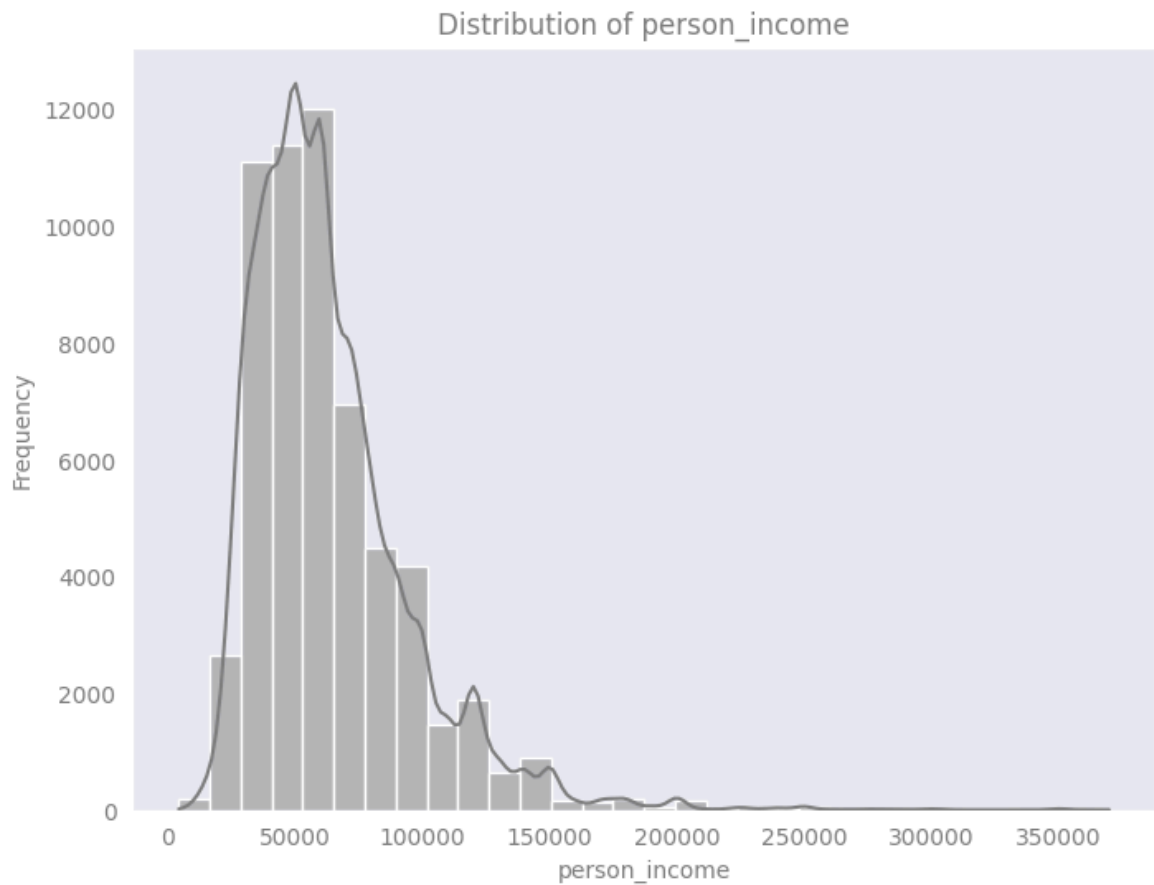
```
In [38]: df2 = outlier_removal(df1, 'person_income')
```



Q1 = 40000.0, Q3 = 175000.0, IQR = 135000.0  
lower\_bound = -162500.0, upper\_bound = 377500.0  
Scatter after outlier removal

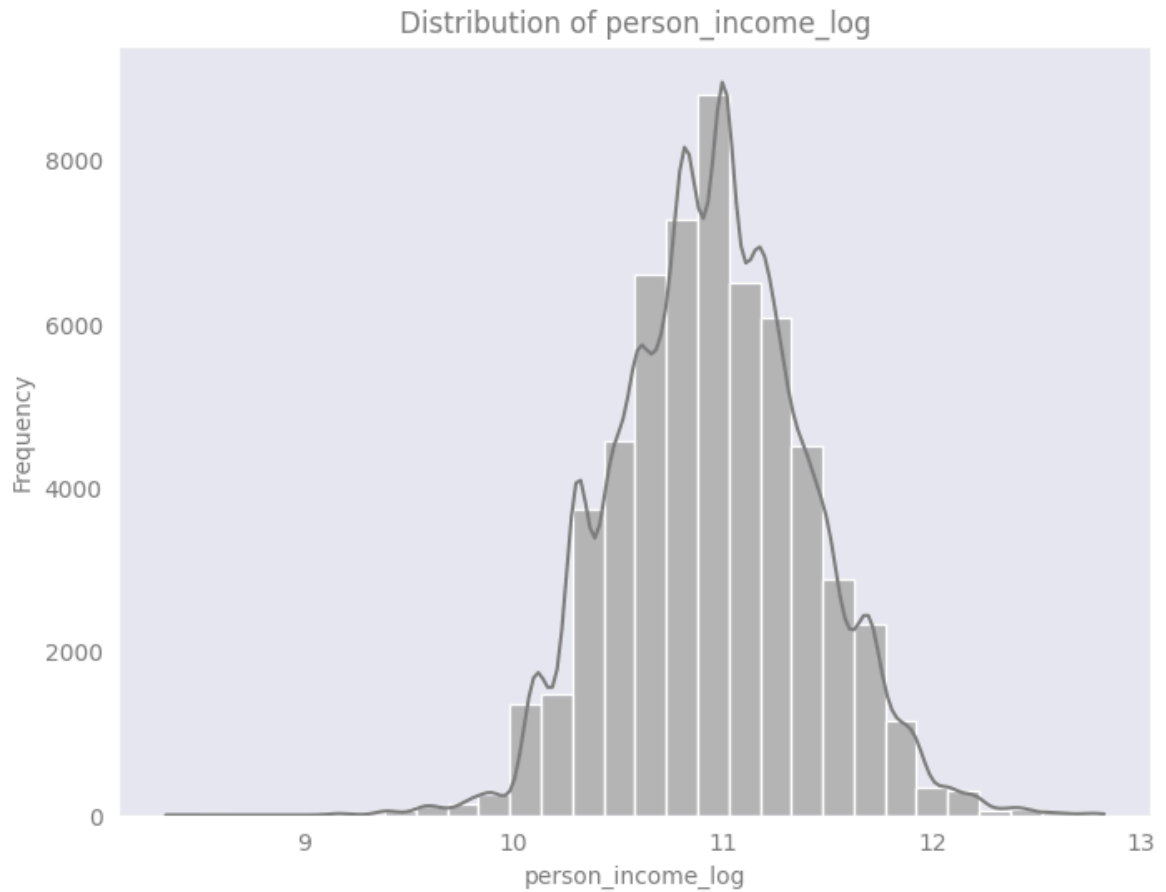


```
In [39]: bar_plotter(df2, 'person_income')
```



```
In [40]: # log transformation
df2['person_income_log'] = df2['person_income'].apply(lambda x: np.log(x))
bar_plotter(df2, 'person_income_log')
```






----> *person\_emp\_length* column


## Data Transformation Techniques

- **Outlier Detection:** Identifying extreme values in the dataset that may skew model performance.
- **Log Transformation:** Applying a logarithmic function to reduce skewness in data distribution.

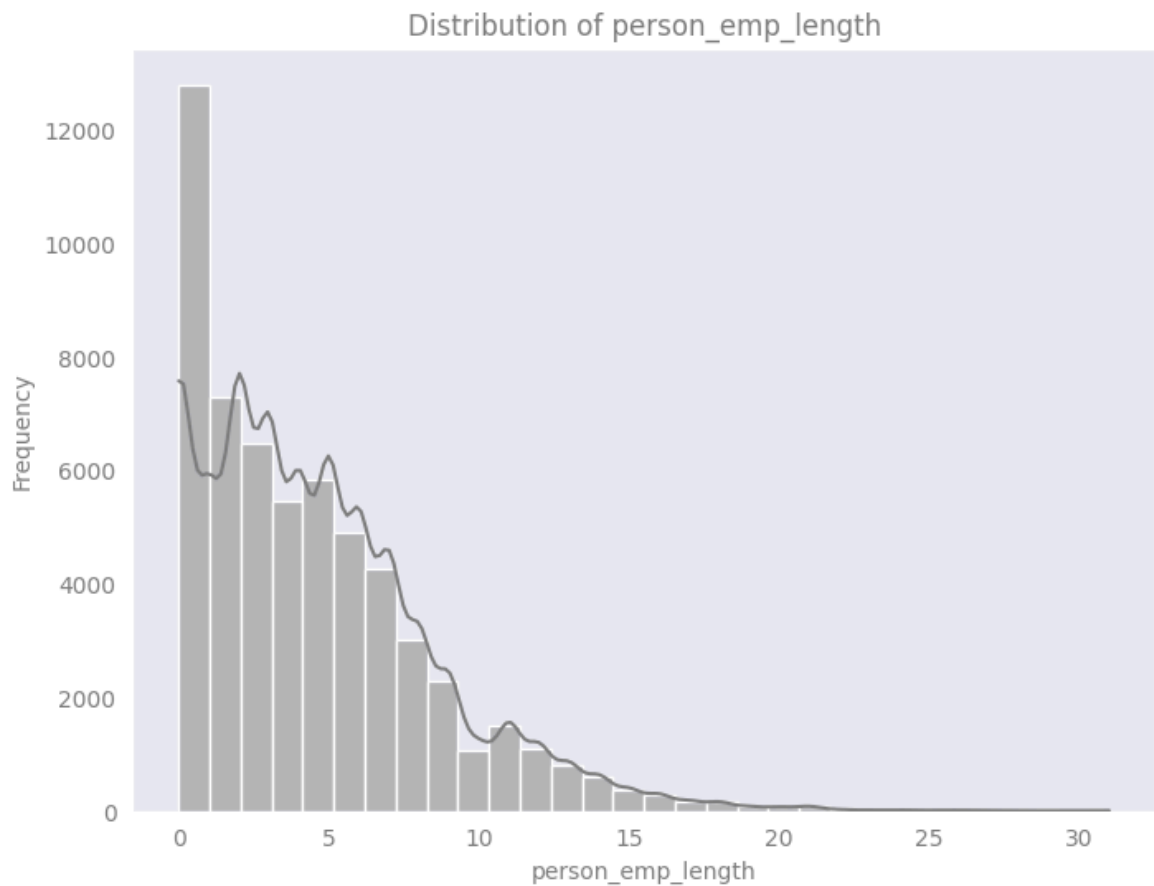
```
In [41]: df3 = outlier_removal(df1, 'person_emp_length', 0.25, 0.98)
```



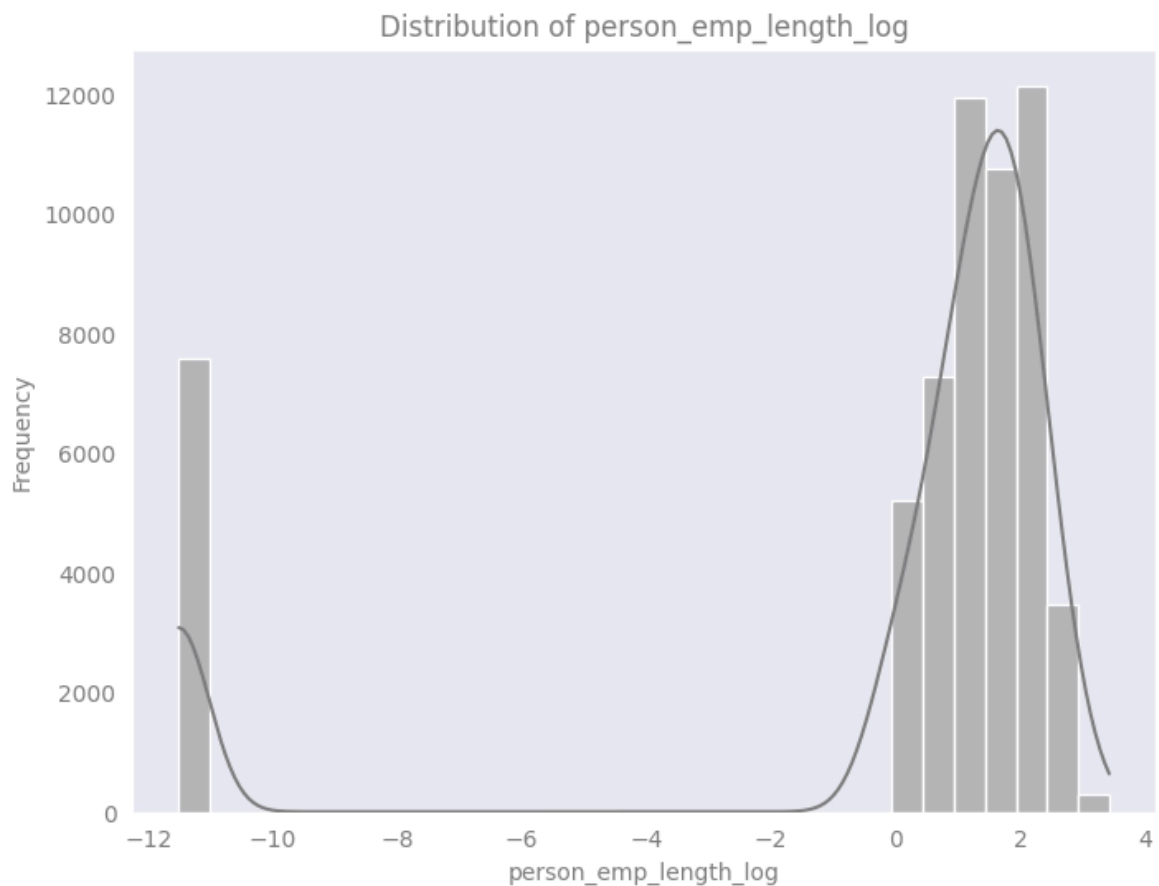
Q1 = 2.0, Q3 = 15.0, IQR = 13.0  
lower\_bound = -17.5, upper\_bound = 34.5  
Scatter after outlier removal



```
In [42]: bar_plotter(df3, 'person_emp_length')
```




```
In [43]: # Log transformation
df3['person_emp_length_log'] = df3['person_emp_length'].apply(lambda x: np.log(x+1))
bar_plotter(df3, 'person_emp_length_log')
```




----> loan amount column

```
In [44]: outlier_remover(df3, 'loan_amnt', 0.25, 0.99)
```



Q1 = 5000.0, Q3 = 25000.0, IQR = 20000.0  
lower\_bound = -25000.0, upper\_bound = 55000.0  
Scatter after outlier removal



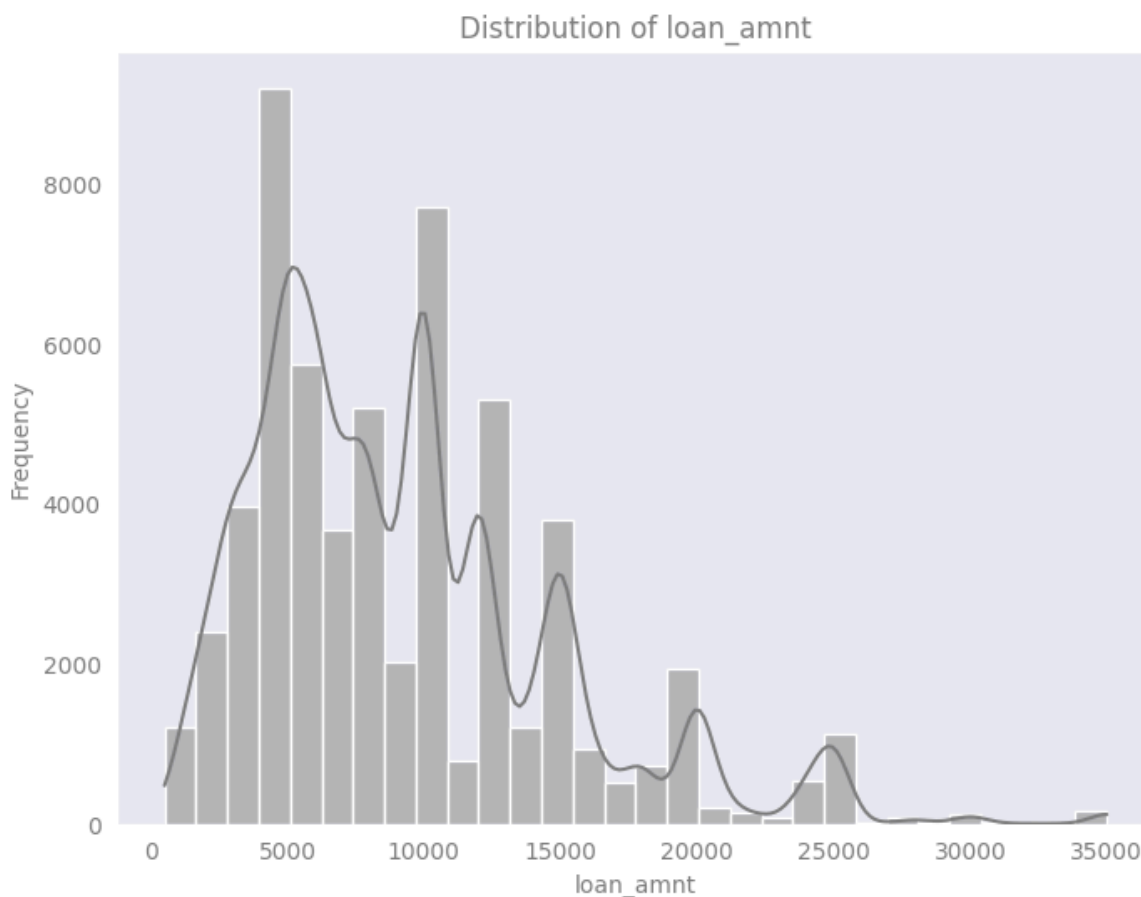
Out[44]:

	id	person_age	person_income	person_home_ownership	person_emp_l
0	0.00001	37	35000	RENT	0
1	1.00000	22	56000	OWN	6
2	2.00000	29	28800	OWN	8
3	3.00000	30	70000	RENT	14
4	4.00000	22	60000	RENT	2
...	...	...	...	...	...
58640	58640.00000	34	120000	MORTGAGE	5
58641	58641.00000	28	28800	RENT	0
58642	58642.00000	23	44000	RENT	7
58643	58643.00000	22	30000	RENT	2
58644	58644.00000	31	75000	MORTGAGE	2


58638 rows × 15 columns



In [45]: bar\_plotter(df3, 'loan\_amnt')




```
In [46]: # Log transformation
df3['loan_amnt_log'] = df3['loan_amnt'].apply(lambda x: np.log(x))
bar_plotter(df3, 'loan_amnt_log')
```



----> *loan\_int\_rate column*


```
In [47]: df4 = outlier_removal(df3, 'loan_int_rate', 0.25, 0.75)
```



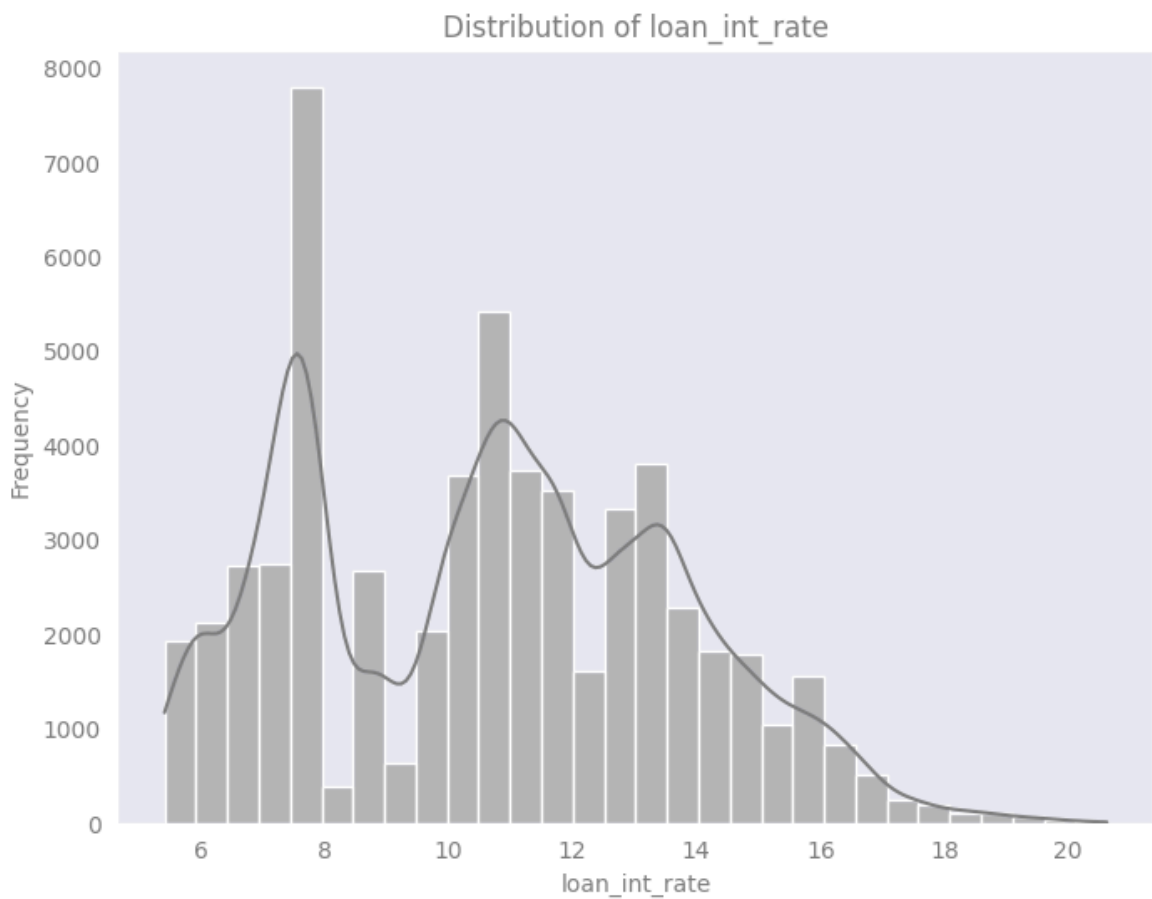
Q1 = 7.88, Q3 = 12.99, IQR = 5.11

lower\_bound = 0.21499999999999897, upper\_bound = 20.655

Scatter after outlier removal




```
In [48]: bar_plotter(df4, 'loan_int_rate')
```




----> *loan\_percent\_income* column

```
In [49]: df5 = outlier_removal(df4, 'loan_percent_income', 0.25, 0.85)
```

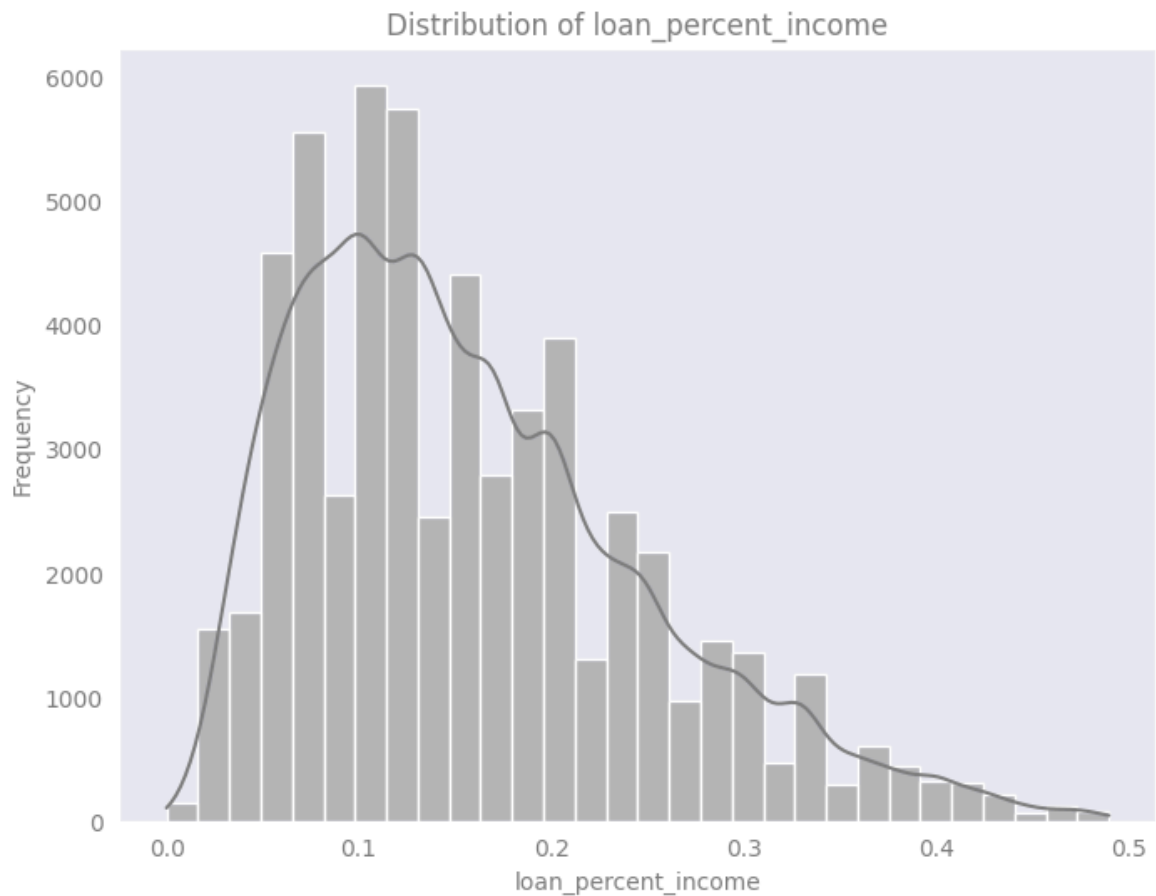


Q1 = 0.09, Q3 = 0.25, IQR = 0.16  
lower\_bound = -0.15, upper\_bound = 0.49  
Scatter after outlier removal





```
In [50]: bar_plotter(df5, 'loan_percent_income')
```




----> *cb\_person\_cred\_hist\_length* column


## Data Transformation Techniques

- **Outlier Detection:** Identifying extreme values in the dataset that may skew model performance.


```
In [51]: df6 = outlier_remover(df5, 'cb_person_cred_hist_length', 0.25, 0.85)
```



Q1 = 3.0, Q3 = 10.0, IQR = 7.0  
 lower\_bound = -7.5, upper\_bound = 20.5  
 Scatter after outlier removal



```
In [52]: bar_plotter(df6, 'cb_person_cred_hist_length')
```



In [53]: `original_df_train.shape[0] - df6.shape[0]`

Out[53]: 513



## Categorical Column [Label Encoding]

- 1. Label Encoding

```
In [54]: cat_columns = ['loan_intent', 'loan_grade', 'cb_person_default_on_file', 'person_home_ownership']

for col in cat_columns:
    print(f"{col} Unique values = ", df6[col].unique())
```

```
loan_intent Unique values = ['EDUCATION' 'MEDICAL' 'PERSONAL' 'VENTURE' 'DEBTCONSOLIDATION'
                             'HOMEIMPROVEMENT']
loan_grade Unique values = ['B' 'C' 'A' 'D' 'E' 'F' 'G']
cb_person_default_on_file Unique values = ['N' 'Y']
person_home_ownership Unique values = ['RENT' 'OWN' 'MORTGAGE' 'OTHER']
```

```
In [55]: from sklearn.preprocessing import LabelEncoder

encoder = LabelEncoder()

for col in cat_columns:
```

```
df6.loc[:,col+"_en"] = encoder.fit_transform(df6[col])
df6.head()
```

Out[55]:

	id	person_age	person_income	person_home_ownership	person_emp_length	loan_status
0	0.00001	37	35000	RENT	0.00001	EMERGENCY
1	1.00000	22	56000	OWN	6.00000	GOOD
2	2.00000	29	28800	OWN	8.00000	GOOD
3	3.00000	30	70000	RENT	14.00000	GOOD
4	4.00000	22	60000	RENT	2.00000	GOOD

In [56]:

```
updated_columns = ['person_age_log', 'person_emp_length_log', 'loan_amnt_log', 'loan_status_log']
original_columns = ['_'.join(val.split('_')[:-1]) for val in updated_columns]

# removing source columns
df7 = df6.drop(original_columns, axis='columns')
df7.head()
```

Out[56]:

	id	person_income	loan_int_rate	loan_percent_income	cb_person_cred_hist_length	loan_status
0	0.00001	35000	11.49	0.17	1	EMERGENCY
1	1.00000	56000	13.35	0.07	1	GOOD
2	2.00000	28800	8.90	0.21	1	GOOD
3	3.00000	70000	11.11	0.17	1	GOOD
4	4.00000	60000	6.92	0.10	1	GOOD

In [57]:

```
df7['loan_status_up'] = df7['loan_status'].apply(lambda x: int(x))
df8 = df7.copy()

df8['loan_status'] = df8['loan_status_up']

if 'loan_status_up' in df8.columns:
    df8.drop('loan_status_up', axis = 'columns', inplace=True)

df8.head()
```

Out[57]:

	id	person_income	loan_int_rate	loan_percent_income	cb_person_cred_hist_length	loan_status
0	0.00001	35000	11.49	0.17	1	EMERGENCY
1	1.00000	56000	13.35	0.07	1	GOOD
2	2.00000	28800	8.90	0.21	1	GOOD
3	3.00000	70000	11.11	0.17	1	GOOD
4	4.00000	60000	6.92	0.10	1	GOOD

In [58]: `df8.shape, original_df_train.shape`

Out[58]: `((58132, 13), (58645, 13))`

```
In [59]: def post_remover(val):
        if '_log' in val:
            val = val.replace('_log', '')
        elif "_en" in val:
            val = val.replace('_en', '')
        return val

input_columns = [post_remover(val) for val in df8.columns]
input_columns
```

Out[59]: `['id',
'person_income',
'loan_int_rate',
'loan_percent_income',
'cb_person_cred_hist_length',
'loan_status',
'person_age',
'person_emp_length',
'loan_amnt',
'loan_intent',
'loan_grade',
'cb_person_default_on_file',
'person_home_ownership']`

```
In [60]: for c1,c2 in zip(input_columns, df8.columns):
        if c1 not in c2:
            print(c1, c2)
            print("Not matched")
```

```
In [61]: df8.columns = input_columns
df8.head()
```

Out[61]:

	id	person_income	loan_int_rate	loan_percent_income	cb_person_cred_hist_leng
--	----	---------------	---------------	---------------------	--------------------------

0	0.00001	35000	11.49	0.17	
1	1.00000	56000	13.35	0.07	
2	2.00000	28800	8.90	0.21	
3	3.00000	70000	11.11	0.17	
4	4.00000	60000	6.92	0.10	

<  >

In [97]:



## Data Processing [Test Data]

- 1. Label Encoding

## 2. Log Transformation


```
In [98]: num_cols, cat_cols = original_columns[:3], original_columns[3:]
num_cols, cat_cols
```

```
Out[98]: (['person_age', 'person_emp_length', 'loan_amnt'],
 ['loan_intent',
  'loan_grade',
  'cb_person_default_on_file',
  'person_home_ownership'])
```

```
In [99]: test_df.head()
```

```
Out[99]:
```

	person_age	person_income	person_home_ownership	person_emp_length	loan_intent
0	23	69000	RENT	3.0	HOMEIMPROVEMENT
1	26	96000	MORTGAGE	6.0	DEBTCONSOLIDATION
2	26	30000	RENT	5.0	DEBTCONSOLIDATION
3	33	50000	RENT	4.0	DEBTCONSOLIDATION
4	26	102000	MORTGAGE	8.0	HOMEIMPROVEMENT

<  >

```
In [100... if 'id' in test_df.columns:
test_df = test_df.drop('id', axis='columns')
```

```
In [101... test_df.shape, df8.shape
```

```
Out[101... ((39098, 11), (58132, 13))
```

```
In [102... test_df_log = test_df[num_cols]
test_df_en = test_df[cat_cols]
test_remain = test_df.drop(num_cols+cat_cols, axis='columns')
```

```
In [103... test_df_log_up = test_df_log.applymap(lambda x: np.log(x))
test_df_log_up.head()
```

```
Out[103...

```

	person_age	person_emp_length	loan_amnt
0	3.135494	1.098612	10.126631
1	3.258097	1.791759	9.210340
2	3.258097	1.609438	8.294050
3	3.496508	1.386294	8.853665
4	3.258097	2.079442	9.615805

```
In [104... for col in cat_cols:
test_df_en.loc[:,col] = encoder.fit_transform(test_df_en[col])
test_df_en.head()
```

Out[104...

	loan_intent	loan_grade	cb_person_default_on_file	person_home_ownership
0	2	5	0	3
1	4	2	1	0
2	5	4	1	3
3	0	0	0	3
4	2	3	1	0

In [105...

```
final_test_df = pd.concat([test_df_log_up, test_df_en, test_remain], axis='column')
final_test_df.head()
```

Out[105...

	person_age	person_emp_length	loan_amnt	loan_intent	loan_grade	cb_person_defai
0	3.135494	1.098612	10.126631	2	5	
1	3.258097	1.791759	9.210340	4	2	
2	3.258097	1.609438	8.294050	5	4	
3	3.496508	1.386294	8.853665	0	0	
4	3.258097	2.079442	9.615805	2	3	



In [106...

```
final_test_df.shape
```

Out[106...

```
(39098, 11)
```

In [107...

```
final_test_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 39098 entries, 0 to 39097
Data columns (total 11 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   person_age                            39098 non-null  float64
1   person_emp_length                     39098 non-null  float64
2   loan_amnt                             39098 non-null  float64
3   loan_intent                           39098 non-null  object
4   loan_grade                            39098 non-null  object
5   cb_person_default_on_file             39098 non-null  object
6   person_home_ownership                 39098 non-null  object
7   person_income                         39098 non-null  int64
8   loan_int_rate                         39098 non-null  float64
9   loan_percent_income                   39098 non-null  float64
10  cb_person_cred_hist_length             39098 non-null  int64
dtypes: float64(5), int64(2), object(4)
memory usage: 3.3+ MB
```

In [110...

```
X = df8.drop(['id', 'loan_status'], axis='columns').columns
y = ['loan_status']
```

In [111...

```
last_one = final_test_df[X]

last_one.head()
```

Out[111...

	person_income	loan_int_rate	loan_percent_income	cb_person_cred_hist_length	perso
0	69000	15.76	0.36	2	3.1
1	96000	12.68	0.10	4	3.2
2	30000	17.19	0.13	2	3.2
3	50000	8.90	0.14	7	3.4
4	102000	16.32	0.15	4	3.2

In [112...

df8.head()

Out[112...

	id	person_income	loan_int_rate	loan_percent_income	cb_person_cred_hist_leng
0	0.00001	35000	11.49	0.17	
1	1.00000	56000	13.35	0.07	
2	2.00000	28800	8.90	0.21	
3	3.00000	70000	11.11	0.17	
4	4.00000	60000	6.92	0.10	



## Model Training

- 1. Hyperparameter Tuning [Grid Search CV]
- 2. Model Training

In [114...

```
X_train, X_test, y_train, y_test = train_test_split(df8.drop(['loan_status', 'id'], axis=1), df8['loan_status'])

model = RandomForestClassifier()
model.fit(X_train, y_train)
"Train score", model.score(X_train, y_train), "Test score", model.score(X_test,
```

Out[114...

```
('Train score', 0.9999770636941214, 'Test score', 0.9504575793022776)
```

In [115...

```
from sklearn.metrics import accuracy_score, classification_report

def train_and_predict(classifier, X_train, X_test, y_train, y_test):
    classifier.fit(X_train, y_train)

    predictions = classifier.predict(X_test)

    train_score = classifier.score(X_train, y_train)
    test_score = accuracy_score(y_test, predictions)

    print(f"Model: {classifier.__class__.__name__}")
```



```
print(f"Training Score: {train_score:.4f}")
print(f"Testing Score: {test_score:.4f}")
print("\nClassification Report:")
print(classification_report(y_test, predictions))

return classifier, predictions

classifiers = [
    LogisticRegression(),
    SVC(),
    DecisionTreeClassifier(),
    RandomForestClassifier(),
    GradientBoostingClassifier(),
    AdaBoostClassifier(),
    KNeighborsClassifier(),
    GaussianNB()
]

for classifier in classifiers:
    print(f'Using classifier: {classifier.__class__.__name__}')
    trained_model, predictions = train_and_predict(classifier, X_train, X_test,
    print("\n" + "-" * 50 + "\n")
```

Using classifier: LogisticRegression  
Model: LogisticRegression  
Training Score: 0.8593  
Testing Score: 0.8594

## Classification Report:

	precision	recall	f1-score	support
0	0.86	1.00	0.92	12489
1	0.00	0.00	0.00	2044
accuracy			0.86	14533
macro avg	0.43	0.50	0.46	14533
weighted avg	0.74	0.86	0.79	14533

-----

Using classifier: SVC  
Model: SVC  
Training Score: 0.8593  
Testing Score: 0.8594

## Classification Report:

	precision	recall	f1-score	support
0	0.86	1.00	0.92	12489
1	0.00	0.00	0.00	2044
accuracy			0.86	14533
macro avg	0.43	0.50	0.46	14533
weighted avg	0.74	0.86	0.79	14533

-----

Using classifier: DecisionTreeClassifier  
Model: DecisionTreeClassifier  
Training Score: 1.0000  
Testing Score: 0.9131

## Classification Report:

	precision	recall	f1-score	support
0	0.96	0.94	0.95	12489
1	0.68	0.73	0.70	2044
accuracy			0.91	14533
macro avg	0.82	0.84	0.83	14533
weighted avg	0.92	0.91	0.91	14533

-----

Using classifier: RandomForestClassifier  
Model: RandomForestClassifier  
Training Score: 1.0000  
Testing Score: 0.9508

## Classification Report:

	precision	recall	f1-score	support
0	0.95	0.99	0.97	12489
1	0.92	0.71	0.80	2044
accuracy			0.95	14533
macro avg	0.94	0.85	0.89	14533
weighted avg	0.95	0.95	0.95	14533

-----

Using classifier: GradientBoostingClassifier  
 Model: GradientBoostingClassifier  
 Training Score: 0.9466  
 Testing Score: 0.9468

Classification Report:

	precision	recall	f1-score	support
0	0.95	0.99	0.97	12489
1	0.90	0.70	0.79	2044
accuracy			0.95	14533
macro avg	0.93	0.84	0.88	14533
weighted avg	0.95	0.95	0.94	14533

-----

Using classifier: AdaBoostClassifier  
 Model: AdaBoostClassifier  
 Training Score: 0.9278  
 Testing Score: 0.9269

Classification Report:

	precision	recall	f1-score	support
0	0.94	0.97	0.96	12489
1	0.80	0.65	0.71	2044
accuracy			0.93	14533
macro avg	0.87	0.81	0.84	14533
weighted avg	0.92	0.93	0.92	14533

-----

Using classifier: KNeighborsClassifier  
 Model: KNeighborsClassifier  
 Training Score: 0.9140  
 Testing Score: 0.8858

Classification Report:

	precision	recall	f1-score	support
0	0.91	0.96	0.94	12489
1	0.64	0.42	0.51	2044
accuracy			0.89	14533

macro avg	0.78	0.69	0.72	14533
weighted avg	0.87	0.89	0.88	14533

-----

Using classifier: GaussianNB  
 Model: GaussianNB  
 Training Score: 0.8717  
 Testing Score: 0.8686

Classification Report:

	precision	recall	f1-score	support
0	0.91	0.94	0.93	12489
1	0.54	0.41	0.47	2044
accuracy			0.87	14533
macro avg	0.73	0.68	0.70	14533
weighted avg	0.86	0.87	0.86	14533

-----

### ----> Grid Search CV

```
In [116... from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.model_selection import GridSearchCV

rf_param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [None, 120, 30],
    'min_samples_split': [5, 10],
    'min_samples_leaf': [2, 4]
}

gb_param_grid = {
    'n_estimators': [50, 100, 200],
    'learning_rate': [0.01, 0.1, 0.2],
    'max_depth': [3, 5, 7],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

rf_grid_search = GridSearchCV(estimator=RandomForestClassifier(), param_grid=rf_
                             scoring='accuracy', cv=5, n_jobs=-1, verbose=1)
rf_grid_search.fit(X_train, y_train)

print("Best parameters for RandomForestClassifier:")
print(rf_grid_search.best_params_)
print("Best training score:", rf_grid_search.best_score_)
```

Fitting 5 folds for each of 24 candidates, totalling 120 fits

Best parameters for RandomForestClassifier:

```
{'max_depth': 30, 'min_samples_leaf': 2, 'min_samples_split': 10, 'n_estimators':
200}
```

Best training score: 0.9488750182823338

**----> Model Training**

```
In [117... best_params = {
    'max_depth': 30,
    'min_samples_leaf': 2,
    'min_samples_split': 5,
    'n_estimators': 200
}
final_model = RandomForestClassifier(**best_params)
final_model.fit(X_train, y_train)
y_pred = final_model.predict(X_test)
test_accuracy = accuracy_score(y_test, y_pred)
print(f'Test Accuracy: {test_accuracy:.4f}')
```

Test Accuracy: 0.9511

```
In [125... sub_model = RandomForestClassifier(**best_params)
sub_model.fit(df8.drop(['loan_status', 'id'], axis='columns'), df8['loan_status'])
result = sub_model.predict(last_one)
print(len(result), test_df.shape)
```

39098 (39098, 11)

**----> Model Evaluation**

## Model Evaluation


- 1. Confusion Matrix
- 2. Classification Report

```
In [126... cm = confusion_matrix(y_test, y_pred)

plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Predicted: 0', 'Predicted: 1'],
            yticklabels=['Actual: 0', 'Actual: 1'])

plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')

plt.show()
```



```
In [127... print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.95	0.99	0.97	12489
1	0.93	0.71	0.80	2044
accuracy			0.95	14533
macro avg	0.94	0.85	0.89	14533
weighted avg	0.95	0.95	0.95	14533

```
In [128... import joblib
joblib.dump(final_model, "loan_status.pkl")
```

```
Out[128... ['loan_status.pkl']
```



## Submission

- 1. Confusion Matrix
- 2. Classification Report

```
In [129... final = pd.concat([original_df_test['id'] , pd.DataFrame(result)], axis='columns')
final.columns = ['id', 'loan_status']
final.head()
```

```
Out[129...      id  loan_status
0  58645           1
1  58646           0
2  58647           1
3  58648           0
4  58649           0
```

```
In [131... final.to_csv("loan_status_submission_2.csv", index=False)
print("Collected submissions")
```

Collected submissions

```
In [ ]:
```